


The XSB System
Version 3.4
Volume 2: Libraries, Interfaces and Packages



October 14, 2013

Credits

Interfaces have become an increasingly important part of XSB. The interface from C to Prolog was implemented by David Warren as was the DLL interface; the interface from Prolog to C (foreign language interface) was developed by Jiyang Xu, Kostis Sagonas and Steve Dawson. The Oracle interface was written by Hassan Davulcu and Ernie Johnson. The ODBC took as its starting point the Oracle interface, and was written by Lily Dong and Baoqiu Cui, and maintained by David Warren. The interface to POSIX regular expression and wildcard matching as well as the Libwww-based Web access package was written by Michael Kifer. The interface to Perl pattern matching routines was written by Michael Kifer and Jin Yu. The SModels interface was written by Luis F. Castro.

The SLX preprocessor was written by José Júlio Alferes and Luís Moniz Pereira. Unix-style scripting libraries were written by Terrance Swift, and the ordset library was written by Richard O’Keefe.

Contents

1 Library Utilities	1
1.1 List Processing	1
1.1.1 Processing Comma Lists	3
1.2 Attributed Variables	3
1.2.1 Lowlevel Interface	4
1.3 <code>constraintLib</code> : a library for CLP	7
1.4 Formatted Output	9
1.5 String Manipulation	11
1.6 Script Writing Utilities	12
1.6.1 Communication with Subprocesses	14
1.7 Socket I/O	21
1.8 Arrays	27
1.9 The Profiling Library	27
1.10 Gensym	29
1.11 Random Number Generator	30
1.12 Loading Separated Files	31
1.13 Scanning in Prolog	32
1.14 XSB Lint	33
1.15 Miscellaneous Predicates	35
1.16 Other Libraries	35
1.16.1 Justification	36
1.16.2 AVL Trees	36
1.16.3 Ordered Sets: <code>ordsets.P</code>	36
1.16.4 Unweighted Graphs: <code>ugraphs.P</code>	36

1.16.5	Heaps: <code>heaps.P</code>	37
2	XSB-ODBC Interface	38
2.1	Introduction	38
2.2	Using the Interface	39
2.2.1	Connecting to and Disconnecting from Data Sources	39
2.2.2	Accessing Tables in Data Sources Using SQL	40
2.2.3	Cursor Management	42
2.2.4	Accessing Tables in Data Sources through the Relation Level	42
2.2.5	Using the Relation Level Interface	43
2.2.6	Handling NULL values	44
2.2.7	The View Level Interface	46
2.2.8	Insertions and Deletions of Rows through the Relational Level	48
2.2.9	Access to Data Dictionaries	50
2.2.10	Other Database Operations	50
2.2.11	Transaction Management	51
2.2.12	Interface Flags	51
2.2.13	Datalog	52
2.3	Error messages	53
2.4	Notes on specific ODBC drivers	53
3	The New XSB-Database Interface	54
3.1	Introduction	54
3.2	Configuring the Interface	54
3.3	Using the Interface	57
3.3.1	Connecting to and Disconnecting from Databases	57
3.3.2	Querying Databases	58
3.4	Error Handling	59
3.5	Notes on specific drivers	61
4	Introduction to XSB Packages	63
5	Wildcard Matching	64

6	pcre: Pattern Matching and Substitution Using PCRE	66
6.1	Introduction	66
6.2	Pattern matching	66
6.3	String Substitution	67
6.4	Installation and configuration	68
6.4.1	Configuring for Linux, Mac, and other Unices	68
6.4.2	Configuring for Windows	68
7	curl: The XSB Internet Access Package	70
7.1	Introduction	70
7.2	Integration with File I/O	70
7.2.1	Opening a Web Document	71
7.2.2	Closing a Web Document	71
7.3	Low Level Predicates	72
7.3.1	Loading web documents	72
7.3.2	Retrieve the properties of a web document	72
7.3.3	Encode Url	72
7.3.4	Obtaining the Redirection URL	73
7.4	Installation and configuration	73
8	sgml and xpath: SGML/XML/HTML Parsers and XPath	74
8.1	Introduction	74
8.2	Overview of the SGML Parser	75
8.3	Predicate Reference	76
8.3.1	Loading Structured Documents	76
8.3.2	Handling of White Spaces	79
8.3.3	XML documents	79
8.3.4	DTD-Handling	80
8.3.5	Low-level Parsing Primitives	81
8.3.6	External Entities	84
8.3.7	Exceptions	84
8.3.8	Unsupported features	85
8.3.9	Summary of Predicates	86

8.4 XPath support	86
9 rdf: The XSB RDF Parser	89
9.1 Introduction	89
9.2 High-level API	89
9.2.1 RDF Object representation	90
9.2.2 Name spaces	91
9.2.3 Low-level access	91
9.3 Testing the RDF translator	92
10 Constraint Packages	93
10.1 clpr: The CPL(R) package	93
10.1.1 The CLP(R) API	94
10.2 The bounds Package	99
10.2.1 The bounds API	101
11 Constraint Handling Rules	104
11.1 Introduction	104
11.2 Syntax and Semantics	104
11.2.1 Syntax	104
11.2.2 Semantics	105
11.3 CHR in XSB Programs	107
11.3.1 Embedding in XSB Programs	107
11.3.2 Compilation	107
11.4 Useful Predicates	108
11.5 Examples	108
11.6 CHR and Tabling	109
11.6.1 General Issues and Principles	109
11.6.2 Call Abstraction	110
11.6.3 Answer Projection	110
11.6.4 Answer Combination	112
11.6.5 Overview of Tabling-related Predicates	114
11.7 Guidelines	114

11.8 CHRd	114
12 XASP: Answer Set Programming with XSB and Smodels	116
12.1 Installing the Interface	117
12.1.1 Installing the Interface under Unix	117
12.1.2 Installing XASP under Windows using Cygwin	118
12.2 The Smodels Interface	120
12.2.1 Using the Smodels Interface with Multiple Threads	123
12.3 The xnmr_int Interface	123
13 PITA: Probabilistic Inference with Tabling and Answer subsumption	126
13.0.1 Installation	126
13.0.2 Use	128
14 Other XSB Packages	130
14.1 Programming with FLORA-2	130
14.2 Summary of xmc: Model-checking with XSB	132
14.3 slx: Extended Logic Programs under the Well-Founded Semantics	133
14.4 gapza: Generalized Annotated Programs	135

Chapter 1

Library Utilities

In this chapter we introduce libraries of some useful predicates that are supplied with XSB. Interfaces and more elaborate packages are documented in later chapters. These predicates are available only when imported them from (or explicitly consult) the corresponding modules.

1.1 List Processing

The XSB library contains various list utilities, some of which are listed below. These predicates should be explicitly imported from the module specified after the skeletal specification of each predicate. There are a lot more useful list processing predicates in various modules of the XSB system, and the interested user can find them by looking at the sources.

<code>append(?List1, ?List2, ?List3)</code>	<code>module: basics</code>
Succeeds if list <code>List3</code> is the concatenation of lists <code>List1</code> and <code>List2</code> .	
<code>member(?Element, ?List)</code>	<code>module: basics</code>
Checks whether <code>Element</code> unifies with any element of list <code>List</code> , succeeding more than once if there are multiple such elements.	
<code>memberchk(?Element, ?List)</code>	<code>module: basics</code>
Similar to <code>member/2</code> , except that <code>memberchk/2</code> is deterministic, i.e. does not succeed more than once for any call.	
<code>ith(?Index, ?List, ?Element)</code>	<code>module: basics</code>
Succeeds if the <code>Index</code> th element of the list <code>List</code> unifies with <code>Element</code> . Fails if <code>Index</code> is not a positive integer or greater than the length of <code>List</code> . Either <code>Index</code> and <code>List</code> , or <code>List</code> and <code>Element</code> , should be instantiated (but not necessarily ground) at the time of the call.	
<code>delete_ith(+Index, +List, ?Element, ?RestList)</code>	<code>module: listutil</code>
Succeeds if the <code>Index</code> th element of the list <code>List</code> unifies with <code>Element</code> , and <code>RestList</code> is <code>List</code> with <code>Element</code> removed. Fails if <code>Index</code> is not a positive integer or greater than the length of <code>List</code> .	

`log_ith(?Index, ?Tree, ?Element)` module: basics

Succeeds if the `Index`th element of the Tree `Tree` unifies with `Element`. Fails if `Index` is not a positive integer or greater than the number of elements that can be in `Tree`. Either `Index` and `Tree`, or `Tree` and `Element`, should be instantiated (but not necessarily ground) at the time of the call. `Tree` is a list of full binary trees, the first being of depth 0, and each one being of depth one greater than its predecessor. So `log_ith/3` is very similar to `ith/3` except it uses a tree instead of a list to obtain log-time access to its elements.

`log_ith_bound(?Index, ?Tree, ?Element)` module: basics

is like `log_ith/3`, but only if the `Index`th element of `Tree` is non-variable and equal to `Element`. This predicate can be used in both directions, and is most useful with `Index` unbound, since it will then bind `Index` and `Element` for each non-variable element in `Tree` (in time proportional to $N * \log N$, for N the number of non-variable entries in `Tree`.)

`length(?List, ?Length)` module: basics

Succeeds if the length of the list `List` is `Length`. This predicate is deterministic if `List` is instantiated to a list of definite length, but is nondeterministic if `List` is a variable or has a variable tail. If `List` is uninstantiated, it is unified with a list of length `Length` that contains variables.

`same_length(?List1, ?List2)` module: basics

Succeeds if list `List1` and `List2` are both lists of the same number of elements. No relation between the types or values of their elements is implied. This predicate may be used to generate either list (containing variables as elements) given the other, or to generate two lists of the same length, in which case the arguments will be bound to lists of length 0, 1, 2, ...

`select(?Element, ?L1, ?L2)` module: basics

`List2` derives from `List1` by selecting (removing) an `Element` non-deterministically.

`reverse(+List, ?ReversedList)` module: basics

Succeeds if `ReversedList` is the reverse of list `List`. If `List` is not a proper list, `reverse/2` can succeed arbitrarily many times. It works only one way.

`perm(+List, ?Perm)` module: basics

Succeeds when `List` and `Perm` are permutations of each other. The main use of `perm/2` is to generate permutations of a given list. `List` must be a proper list. `Perm` may be partly instantiated.

`subseq(?Sequence, ?SubSequence, ?Complement)` module: basics

Succeeds when `SubSequence` and `Complement` are both subsequences of the list `Sequence` (the order of corresponding elements being preserved) and every element of `Sequence` which is not in `SubSequence` is in the `Complement` and vice versa. That is,

$$length(Sequence) = length(SubSequence) + length(Complement)$$

for example, `subseq([1,2,3,4], [1,3], [2,4])`. The main use of `subseq/3` is to generate subsets and their complements together, but can also be used to interleave two lists in all possible ways.

`merge(+List1, +List2, ?List3)` module: listutil

Succeeds if `List3` is the list resulting from “merging” lists `List1` and `List2`, i.e. the elements of `List1` together with any element of `List2` not occurring in `List1`. If `List1` or `List2` contain duplicates, `List3` may also contain duplicates.

`absmerge(+List1, +List2, ?List3)` module: listutil

Predicate `absmerge/3` is similar to `merge/3`, except that it uses predicate `absmember/2` described below rather than `member/2`.

`absmember(+Element, +List)` module: listutil

Similar to `member/2`, except that it checks for identity (through the use of predicate `'=='/2`) rather than unifiability (through `'='/2`) of `Element` with elements of `List`.

`member2(?Element, ?List)` module: listutil

Checks whether `Element` unifies with any of the actual elements of `List`. The only difference between this predicate and predicate `member/2` is on lists having a variable tail, e.g. `[a, b, c | _]`: while `member/2` would insert `Element` at the end of such a list if it did not find it, Predicate `member2/2` only checks for membership but does not insert the `Element` into the list if it is not there.

`closetail(?List)` module: listutil

Predicate `closetail/1` closes the tail of an open-ended list. It succeeds only once.

1.1.1 Processing Comma Lists

It is often useful to process comma lists when meta-interpreting or preprocessing. XSB libraries include the following simple utilities.

`comma_to_list(+CommaList, -List)` module: basics

Transforms `CommaList` to `List`.

`comma_append(?CL1, ?CL2, ?CL3)` basics

`comma_length(?CommaList, ?Length)` basics

`comma_member(?Element, ?CommaList)` basics

`comma_member(?Element, ?CommaList)` module: basics

Analogues for comma lists of `append/3`, `length/3`, `member/2` and `memberchk/2`, respectively.

1.2 Attributed Variables

Attributed variables are a special data type that associates variables with arbitrary attributes as well as supports extensible unification. Attributed variables have proven to be a flexible and powerful mechanism to extend a classic logic programming system with the ability of constraint solving. Our low-level API for constraints closely resembles that of hProlog [?] and SWI [?].

1.2.1 Lowlevel Interface

Attributes of variables are pairs of attribute module names and values. An attribute module name can be any atom. A value can be any XSB value (term, variable, atom, ...). Any variable has at most one attribute for a particular attribute module. Attribute modules are distinct from XSB modules: although it is most efficient to keep each handlers for each attribute module in their own XSB module. Attributes can be manipulated with the following three predicates (`get_attr/3`, `put_attr/3` and `del_attr/2`) defined in the module `machine`.

```
get_attr(-Var,+Mod, ?Val)                                module: machine
    Gets the value of the attribute of Var in attribute module Mod. Non-variable terms in Var
    cause a type error. Val will be unified with the value of the attribute, if it exists. Otherwise
    the predicate fails.
```

```
put_attr(-Var,+Mod, ?Val)                                module: machine
    Sets the value of the attribute of Var in attribute module Mod. Non-variable terms in Var
    cause a type error. The previous value of the attribute is overwritten, if it exists.
```

```
del_attr(-Var, +Mod)                                     module: machine
    Removes the attribute of Var in attribute module Mod. Non-variable terms in Var cause a
    type error. The previous value of the attribute is removed, if it exists.
```

One has to extend the default unification algorithm for used attributes by installing a handler in the following way:

```
:- install_verify_attribute_handler(+Mod,-AttrValue,-Target,+Handler,+WarningFlag)
:- install_verify_attribute_handler(+Mod,-AttrValue,-Target,+Handler)
```

The predicates `install_verify_attribute_handler/5` and `install_verify_attribute_handler/4` are defined in module `machine`. *Mod* is the attribute Module and *Handler* is a term with arguments *AttrValue* and *Target*. The *Handler* term has to correspond to a handler predicate that takes the value of the attribute (*AttrValue*) and the term that the attributed value is bound to (*Target*) as arguments. The argument *WarningFlag* in the 5-argument version of the predicate can be used to suppress the warning issued when replacing the `verify_attribute_handler` for a module. If the argument is `warning_on` then the warning is issued if a handler for the module already exists. Otherwise, the warning is suppressed. The 4-argument version of the predicate does *not* suppress the warning.

To get good efficiency, it is usually best to keep the handlers for each attribute module in separate XSB modules. The handler is called after the unification of an attributed variable with a term or other attributed variable, if the attributed variable has an attribute in the corresponding module. The two arguments of the unification are already bound at the time the handler is called, i.e. the handler is a post-unify handler.

Here, by giving the implementation of a simple finite domain constraint solver (see the file `fd.P` below), we show how these lowlevel predicates for attributed variables can be used. In this example, an attribute in the module `fd` is used and the value of this attribute is a list of terms.

```

%% File: fd.P
%%
%% A simple finite domain constraint solver implemented using the lowlevel
%% attributes variables interface.

:- import put_attr/3, get_attr/3, del_attr/2,
    install_verify_attribute_handler/4 from machine.
:- import member/2 from basics.

:- install_verify_attribute_handler(fd,AttrValue,Target,fd_handler(AttrValue,Target)).

fd_handler(Da, Target) :-
    (var(Target),                                     % Target is an attributed variable
     get_attr(Target, fd, Db) ->                     % has a domain
     intersection(Da, Db, [E|Es]),                   % intersection not empty
     (Es = [] ->                                     % exactly one element
      Target = E                                     % bind Var (and Value) to E
      ; put_attr(Target, fd, [E|Es])                 % update Var's (and Value's)
     )
    ; member(Target, Da)                             % is Target a member of Da?
    ).

intersection([], _, []).
intersection([H|T], L2, [H|L3]) :-
    member(H, L2), !,
    intersection(T, L2, L3).
intersection([_|T], L2, L3) :-
    intersection(T, L2, L3).

domain(X, Dom) :-
    var(Dom), !,
    get_attr(X, fd, Dom).
domain(X, List) :-
    List = [E1|Els],                                % at least one element
    (Els = []                                       % exactly one element
     -> X = E1                                     % implied binding
    ; put_attr(Fresh, fd, List),                  % create a new attributed variable
      X = Fresh                                    % may call verify_attributes/2
    ).

show_domain(X) :-
    var(X),                                         % print out the domain of X
    get_attr(X, fd, D),                           % X must be a variable
    write('Domain of '), write(X),
    write(' is '), writeln(D).

```

When writing or porting a constraint package, it is usually useful to adjust the way that correct answer substitutions are shown in the command line. This can be controlled using the following two predicates:

```
install_attribute_portray_hook(Module,Attribute,Handler)          module: machine
```

This hook is called by the command-line interpreter when printing out the value of each variable in a top-level query. When a printing out an attributed variable, any appropriate handlers are called to portray the constraints represented by the attribute. As an example, the `bounds` package (Section 10.2) uses a hook to print out the bounds of variables:

```
| ?- X in 1..10,Y in 1..10,X + 4 #< Y -3.
```

```
X = _h629 { bounds : 1 .. 2 }
```

```
Y = _h673 { bounds : 9 .. 10 }
```

Writing a handler can be as simple as possible or as elaborate as desired. In the case of `bounds` the handler is simple:

```
bounds_attr_portray_hook(bounds(L,U,_)) :- write(L..U).
```

The hook is installed when the constraint package is loaded by placing in the package loader directive such as:

```
:- install_attribute_portray_hook(bounds,Attr,bounds_attr_portray_hook(Attr)).
```

Note that the hook will be indexed on the module associated with the attribute (in this case `bounds`). XSB's command-line interpreter will unify the second argument of the portray hook with the attribute, and then call `Handler`.

```
install_attribute_constraint_hook(Module,Vars,Names,Handler)    module: machine
```

For some constraint packages, it may not be particularly useful to associate constraints with variables: instead, the projection of global constraints onto the variables of the top-level query may be more useful. This is the case in the `CLP(R)` package (Section 10.1), where the command-line interaction may look as follows:

```
| ?- {X = 2*Y,Y >= 7},inf(X,F).
```

```
{ X >= 14.0000 }
```

```
{ Y = 0.5000 * X }
```

```
X = _h8841
```

```
Y = _h9506
```

```
F = 14.0000
```

In XSB, the (projection of the) global constraints in `CLP(R)` are displayed by the following routines:

```
clpr_portray_varlist(Vars,Names):-
filter_varlist(Vars,Names,V1,N1),
dump(V1,N1,Constraints),
member(C,Constraints),
console_write(' { '), console_write(C),console_writeln(' } '),
fail.
clpr_portray_varlist(_V,_N).
```

```
filter_varlist([],[],[],[]).
```

```
filter_varlist([V1|R1],[N1|R2],[V1|R3],[N1|R4]):-
```

```

var(V1),!,
filter_varlist(R1,R2,R3,R4).
filter_varlist([_V1|R1],[_N1|R2],R3,R4):-
filter_varlist(R1,R2,R3,R4).

```

This predicate sets up a call to the CLP(R) library predicate `dump/3`, whose constraints it then writes out to the console. Analogous to the `portray` hook, the console hook is installed using the directive:

```
:- install_constraint_portray_hook(clpr,Vars,Names,clpr_portray_varlist(Vars,Names)).
```

If the `clpr` module is loaded, the command line interpreter checks any constraint `portray` hooks upon the first success of a top-level goal. It then unifies the second argument `Vars` with the variables of the goal, and `Names` with the names of the variables of the goal which are then passed on to `Handler`

1.3 constraintLib: a library for CLP

XSB supports constraint logic programming through its engine-level support of attributed variables (Section 1.2), and its support for constraint handling rules (CHR) (Chapter 11). The `constraintLib` library includes routines for delaying and examining bindings that are commonly used to implement CHR and other constraint libraries.

When processing constraints, it is often useful to delay a goal based on the instantiation level of a term or set of terms. For instance a $3 > X + Y$ should be delayed until both `X` and `Y` are instantiated. However the goal should be reinvoked as soon as possible after both are instantiated in order to prune search paths that may not be useful to pursue. The predicate `when/2` provides a useful mechanism to delay goals based on instantiation patterns ¹.

```
when(+Condition,Goal)                                module: constraintLib
    Delays the execution of Goal until Condition is satisfied, whereupon Goal will be executed.
    Condition can have the form
```

- `?=(Term1,Term2)`
- `nonvar(Term)`
- `ground(Term)` ²
- `(Condition,Condition)`
- `(Condition ; Condition)`

Example: The following session illustrates the use of `when/2` to delay a goal.

¹Despite the similar name, this method of delaying is conceptually different from SLG DELAYING discussed in Volume 1 of this manual, which is used for resolving cycles of dependencies in computing the well-founded semantics, and is not based on the state of instantiation of a term.

²To use `ground/1` in the condition, it must be imported into the file where it is used.

```
|?- when(nonvar(X),writeln(test(1-2,nonvar))),writeln(test(1,nonvar)),X = f(_Y).
```

```
test(1,nonvar)
test(1 - 2,nonvar)
```

```
X = f(_h245)
```

unifiable(X, Y, -Unifier) module: constraintLib

If **X** and **Y** can unify, succeeds unifying **Unifier** with a list of terms of the form **Var = Value** representing a most general unifier of **X** and **Y**. **unifiable/3** can handle cyclic terms. Attributed variables are handled as normal variables. Associated hooks are not executed ³.

setarg(+Index,+Term,+Value) module: constraintLib

The predicate **setarg/3** provides an efficient but non-logical way to update argument **Index** of a Prolog term **Term** to **Value** via destructive assignment and without the necessity of copying **Term**. **setarg/3** should be used sparingly, to ensure both clarity and portability of code.

Example

```
|?- X = p(f(1),g(2),r([a])),
    writeln(zero(X)),
    ( set_arg(X,2,g([b])),
      writeln(one(X)),
      fail
    ; writeln(two(X))) .
zero(p(f(1),g(2),r([a])))
one(p(f(1),g([b]),r([a])))
two(p(f(1),g(2),r([a])))
```

```
X = p(f(1),g(2),r([a]))
```

Error Cases

- Index is a variable
 - instantiation_error
- Index neither a variable nor an integer
 - type_error(integer,Index)
- Index is less than 0
 - domain_error(not_less_than_zero,Index)
- Term is a variable
 - instantiation_error
- Term neither a variable nor a compound term
 - type_error(compound,Term)

term_variables(+Term,-Variables) module: constraintLib

Given any Prolog term **Term** as input, returns a sorted list of variables in the term.

³In Version 3.4, **unifiable/3** is implemented as a Prolog predicate and so is slower than many of the predicates in this section.

1.4 Formatted Output

```
format(+String,+Control)                                module:  format
```

```
format(+Stream,+String,+Control)                       module:  format
```

`format/2` and `format/3` act as a Prolog analog to the C `stdio` function `printf()`, allowing formatted output ⁴.

Output is formatted according to `String` which can contain either a format control sequence, or any other character which will appear verbatim in the output. Control sequences act as place-holders for the actual terms that will be output. Thus

```
?- format("Hello ~q!",world).
```

will print `Hello world!`.

If there is only one control sequence, the corresponding element may be supplied alone in `Control`. If there are more, `Control` must be a list of these elements. If there are none then `Control` must be an empty list. There have to be as many elements in `Control` as control sequences in `String`.

The character `~` introduces a control sequence. To print a `~` just repeat it:

```
?- format("Hello ~~world!", []).
```

will output `Hello ~world!`.

The general format of a control sequence is `~NC`. The character `C` determines the type of the control sequence. `N` is an optional numeric argument. An alternative form of `N` is `*`. `*` implies that the next argument in `Arguments` should be used as a numeric argument in the control sequence. For example:

```
?- format("Hello~4cworld!", [0'x]).
```

and

```
?- format("Hello~*cworld!", [4,0'x]).
```

both produce

```
Helloxxxxworld!
```

The following control sequences are available in XSB.

⁴The `format` family of predicates is due to Quintus Prolog, by way of Ciao.

- `~a` The argument is an atom. The atom is printed without quoting.
- `~Nc` (Print character.) The argument is a number that will be interpreted as an ASCII code. `N` defaults to one and is interpreted as the number of times to print the character.
- `~f` (Print float). The argument is a float. The float will be printed out by XSB.
- `~d` (Print integer). The argument is an integer, and will be printed out by XSB.
- `~Ns` (Print string.) The argument is a list of ASCII codes. Exactly `N` characters will be printed. `N` defaults to the length of the string. Example:

```
?- format("Hello ~4s ~4s!", ["new","world"]).
?- format("Hello ~s world!", ["new"]).
```

will print as

```
Hello new worl!
Hello new world!
```

respectively.

- `~i` (Ignore argument.) The argument may be of any type. The argument will be ignored. Example:

```
?- format("Hello ~i~s world!", ["old","new"]).
```

will print as

```
Hello new world!
```

- `~k` (Print canonical.) The argument may be of any type. The argument will be passed to `write_canonical/2`). Example:

```
?- format("Hello ~k world!", a+b+c).
```

will print as

```
Hello ++(a,b),c) world!
```

- `~q` (Print quoted.) The argument may be of any type. The argument will be passed to `writeq/2`. Example:

```
?- format("Hello ~q world!", [['A','B']]).
```

will print as

```
Hello ['A','B'] world!
```

- `~w` (write.) The argument may be of any type. The argument will be passed to `write/2`. Example:

```
?- format("Hello ~w world!", [['A','B']]).
```

will print as

```
Hello [A,B] world!
```

- `~Nn` (Print newline.) Print `N` newlines. `N` defaults to 1. Example:

```
?- format("Hello ~n world!", []).
```

will print as

```
Hello
world!
```

1.5 String Manipulation

XSB has a number of powerful predicates that simplify the job of string manipulation. These predicates are especially powerful when they are combined with pattern-matching facilities provided by the `pcre` package described in Chapter 6⁵.

```
str_sub(+Sub, +Str, ?Pos) module: string
```

Succeeds if `Sub` is a substring of `Str`. In that case, `Pos` unifies with the position where the match occurred. Positions start from 0. `str_sub/2` is also available, which is equivalent to having `_` in the third argument of `str_sub/3`.

```
str_match(+Sub, +Str, +Direction, ?Beg, ?End) module: string
```

This is an enhanced version of the previous predicate. `Direction` can be `forward` or `reverse` (or any abbreviation of these). If `forward`, the predicate finds the first match of `Sub` from the beginning of `Str`. If `reverse`, it finds the first match from the end of the string (*i.e.*, the last match of `Sub` from the beginning of `Str`). `Beg` and `End` must be integers or unbound variables. (It is possible that one is bound and another is not.) `Beg` unifies with the offset of the first character where `Sub` matched, and `End` unifies with the offset of the next character to the right of `Sub` (such a character might not exist, but the offset is still defined). Offsets start from 0.

Both `Beg` and `End` can be bound to negative integers. In this case, the value represents the offset from the *second* character past the end of `Str`. Thus `-1` represents the character next to the end of `Str` and can be used to check where the end of `Sub` matches in `Str`. In the following examples

```
?- string_match(Sub,Str,forw,X,-1).
?- string_match(Sub,Str,rev,X,-1).
?- string_match(Sub,Str,forw,0,X).
```

⁵Not all string manipulation predicates have been made thread-safe in Version 3.4.

the first checks if the *first* match of **Sub** from the beginning of **Str** is a suffix of **Str** (because **End** represents the character next to the last character in **Sub**, so **End=-1** means that the last characters of **Sub** and of **Str** occupy the same position). If so, **X** is bound to the offset (from the end of **Str**) of the first character of **Sub**. The second example checks if the *last* match of **Sub** in **Str** is a suffix of **Str** and binds **X** to the offset of the beginning of that match (counted from the beginning of **Str**). The last example checks if the first match of **Sub** is a prefix of **Str**. If so, **X** is bound to the offset (from the beginning of **Str**) of the last character of **Sub**.

```
substring(+String, +BeginOffset, +EndOffset, -Result)           module: string
```

String can be an atom or a list of characters, and the offsets must be integers. If **EndOffset** is negative, `endof(String)+EndOffset+1` is assumed. Thus, **-1** means end of string. If **BeginOffset** is less than 0, then 0 is assumed; if it is greater than the length of the string, then string end is assumed. If **EndOffset** is non-negative, but is less than **BeginOffset**, then empty string is returned.

Offsets start from 0.

The result returned in the fourth argument is a string, if **String** is an atom, or a list of characters, if so is **String**.

The `substring/4` predicate always succeeds (unless there is an error, such as wrong argument type).

Here are some examples:

```
| ?- substring('abcdefg', 3, 5, L).
```

```
L = de
```

```
| ?- substring("abcdefg", 4, -1, L).
```

```
L = [101,102]
```

(*i.e.*, **L** = **ef** represented using ASCII codes).

1.6 Script Writing Utilities

Prolog, (in particular XSB!) can be useful for writing scripts. Prolog's simple syntax and declarative semantics make it especially suitable for scripts that involve text processing. There are several ways to access script-writing commands from XSB. The first is to execute the command via the predicates `shell/1` or `shell/2`. These predicates can execute any command but they do not provide streamability across UNIX and Windows commands, and they do not return any output of commands to Prolog. Special predicates are provided to handle cross-platform compatibility and to bring output into XSB.

Effort has been made to make these thread-safe; however in Version 3.4, calls to the XSB script writing utilities go through a single mutex, and may cause contention if many threads seek to concurrently use sockets.

`expand_filename(+FileName,-ExpandedName)` module: machine

Expands the file name passed as the first argument and binds the variable in the second argument to the expanded name. This includes (1) expanding Unix tildes, (2) prepending `FileName` to the current directory, and (3) “rectifying” the expanded file name. In rectification, the expanded file name is “rectified” so that multiple repeated slashes are replaced with a single slash, the intervening “./” are removed, and “./” are applied so that the preceding item in the path name is deleted. For instance, if the current directory is `/home`, then `abc//cde/..///ff/./b` will be converted into `/home/abc/ff/b`.

Under Windows, this predicate does rectification as described above, (using backslashes when appropriate), but it does not expand the tildes.

`expand_filename_no_prepend(+FileName,-ExpandedName)` module: shell

This predicate behaves as `expand_filename/2`, but only expands tildes and does rectification. It does not prepend the current working directory to relative file names.

`parse_filename(+FileName,-Dir,-Base,-Extension)` module: machine

This predicate parses file names by separating the directory part, the base name part, and file extension. If file extension is found, it is removed from the base name. Also, directory names are rectified and if a directory name starts with a tilde (in Unix), then it is expanded. Directory names always end with a slash or a backslash, as appropriate for the OS at hand.

For instance, `~john///doe/dir1//../foo.bar` will be parsed into: `/home/john/doe/`, `foo`, and `bar` (where we assume that `/home/john` is what `~john` expands into).

`sleep(+Seconds)` module: shell

Put XSB to sleep for a given number of seconds.

Error Cases

- `Seconds` is a variable
 - `instantiation_error`.
- `Seconds` is not an integer
 - `type_error(integer, Seconds)`.

`sys_pid(-Pid)` module: shell

Get Id of the current process.

`getenv(+VarName,-VarVal)` module: machine

Unifies `VarVal` with the value of `VarName` in the current shell. If `VarName` is not an environment variable, the predicate fails.

Example:

```
:- import getenv/2 from machine.
```

```
yes
| ?- getenv('HOSTTYPE',F).
```

```
F = intel-pc
```

putenv(+String) module: machine

If **String** is of the form **VarName=Value**, inserts or resets the environment variable **VarName**. If **VarName** does not exist, it is inserted with **VarVal**. If the **VarName** does exist, it is reset to **VarVal**. **putenv/2** always succeeds.

Exceptions:

instantiation_error **String** is not instantiated at the time of call.

type_error **VarName** or **VarVal** is not an atom or a list of atoms.

epoch_seconds(-Seconds) module: machine

Returns the number of seconds since the beginning of the UNIX/POSIX epoch (January 1, 1970) ⁶. May cause overflow on 32-bit platforms.

1.6.1 Communication with Subprocesses

In the previous section, we have seen several predicates that allow XSB to create other processes. However, these predicates offer only a very limited way to communicate with these processes. The predicate **spawn_process/5** and friends come to the rescue. It allows a user to spawn any process (including multiple copies of XSB) and redirect its standard input and output to XSB streams. XSB can then write to the process and read from it. The section of socket I/O describes yet another mode of interprocess communication.

In addition, the predicate **pipe_open/2** described in this section lets one create any number of pipes (that do not need to be connected to the standard I/O stream) and talk to child processes through these pipes. All predicates in this section, except **pipe_open/2** and **fd2stream/2**, must be imported from module **shell**. The predicates **pipe_open/2** and **fd2stream/2** must be imported from **file_io**.

spawn_process(+CmdSpec,-StreamToProc,-StreamFromProc,-ProcStderrStream,-ProcId)

Spawn a new process specified by **CmdSpec**. **CmdSpec** must be either a single atom or a *list* of atoms. If it is an atom, then it must represent a shell command. If it is a list, the first member of the list must be the name of the program to run and the other elements must be arguments to the program. Program name must be specified in such a way as to make sure the OS can find it using the contents of the environment variable **PATH**. Also note that pipes, I/O redirection and such are not allowed in command specification. That is, **CmdSpec** must represent a single command. (But read about process plumbing below and about the related predicate **shell/5**.)

The next three parameters of **spawn_process** are XSB I/O stream identifiers for the process (leading to the subprocess standard input), from the process (from its standard output), and a stream capturing the subprocess standard error output. The last parameter is the system process id.

Here is a simple example of how it works.

⁶Uses the Posix call **time(0)**, so the number of seconds will be returned on non-Unix platforms, such as Microsoft.

```
| ?- import file_flush/2, file_read_line_atom/2 from file_io.
| ?- import file_nl/1 , file_write/2 from xsb_writ.

| ?- spawn_process([cat, '-'], To, From, Stderr, Pid),
      writeln(To,'Hello cat!'), flush_output(To,_), file_read_line_atom(From,Y).

To = 3
From = 4
Stderr = 5
Pid = 14328
Y = Hello cat!

yes
```

Here we created a new process, which runs the “cat” program with argument “-”. This forces `cat` to read from standard input and write to standard output. The next line writes an atom and newline to the XSB stream `To`, which is bound to the standard input of the `cat` process (proc id 14328). The `cat` process then copies the input to its standard output. Since standard output of the `cat` process is redirected to the XSB stream `From` in the parent process, the last line in our program is able to read it and return in the variable `Y`. Note that in the second line we used `flush_output/2`. Flushing the output is extremely important here, because XSB I/O pipe (file) streams are buffered. Thus, `cat` might not see its input until the buffer is filled up, so the above clause might hang. `flush_output/2` makes sure that the input is immediately available to the subprocess.

In addition to the above general schema, the user can tell `spawn_process/5` not to open one of the communication streams or to use one of the existing communication streams. This is useful when you do not expect to write or read to/from the subprocess or when one process wants to write to another (see the process plumbing example below). To tell that a certain stream is not needed, it suffices to bind that stream to an atom. For instance,

```
| ?- spawn_process([cat, '-'], To, none, none, _),
      nl(To), writeln(To,'Hello cat!'), flush_output(To).

To = 3,
Hello cat!

reads from XSB and copies the result to standard output. Likewise,

| ?- spawn_process('cat library.tex', none, From, none, _),
      file_read_line_atom(From, S).

From = 4
S = \chapter{Library Utilities} \label{library_utilities}
```

In each case, only one of the streams is open. (Note that the shell command is specified as an atom rather than a list.) Finally, if both streams are suppressed, then `spawn_process` reduces to the usual `shell/1` call (in fact, this is how `shell/1` is implemented):

```
| ?- spawn_process([pwd], none, none).
```

```
/usr/local/foo/bar
```

On the other hand, if any one of the three stream variables in `spawn_process` is bound to an already existing file stream, then the subprocess will use that stream (see the process plumbing example below).

One of the uses of XSB subprocesses is to create XSB servers that spawn subprocesses and control them. A spawned subprocess can be another XSB process. The following example shows one XSB process spawning another, sending it a goal to evaluate and obtaining the result:

```
| ?- spawn_process([xsb], To, From, Err, _),
    write(To, 'assert(p(1)).'), flush_output(To, _),
    write(To, 'p(X), writeln(X).'), flush_output(To, _),
    file_read_line_atom(From, XX).
```

```
XX = 1
```

```
yes
| ?-
```

Here the parent XSB process sends “`assert(p(1)).`” and then “`p(X), writeln(X).`” to the spawned XSB subprocess. The latter evaluates the goal and prints (via “`writeln(X)`”) to its standard output. The main process reads it through the `From` stream and binds the variable `XX` to that output.

Finally, we should note that the stream variables in the `spawn_process` predicate can be used to do process plumbing, *i.e.*, redirect output of one subprocess into the input of another. Here is an example:

```
| ?- open(test, write, Stream),
    spawn_process([cat, 'data'], none, FromCat1, none, _),
    spawn_process([sort], FromCat1, Stream, none, _).
```

Here, we first open file `test`. Then `cat data` is spawned. This process has the input and standard error stream blocked (as indicated by the atom `none`), and its output goes into stream `FromCat1`. Then we spawn another process, `sort`, which picks the output from the first process (since it uses the stream `FromCat1` as its input) and sends its own output (the sorted version of `data`) to its output stream `Stream`. However, `Stream` has already been open for output into the file `test`. Thus, the overall result of the above clause is tantamount to the following shell command:

```
cat data | sort > test
```

Important notes about spawned processes:

1. Asynchronous processes spawned by XSB do not disappear (at least on Unix) when they terminate, *unless* the XSB program executes a *wait* on them (see `process_control` below). Instead, such processes become defunct *zombies* (in Unix terminology); they do not do anything, but consume resources (such as file descriptors). So, when a subprocess is known to terminate, it must be waited on.
2. The XSB parent process must know how to terminate the asynchronous subprocesses it spawns. The drastic way is to kill it (see `process_control` below). Sometimes a subprocess might terminate by itself (*e.g.*, having finished reading a file). In other cases, the parent and the child programs must agree on a protocol by which the parent can tell the child to exit. The programs in the XSB subdirectory `examples/subprocess` illustrate this idea. If the child subprocess is another XSB process, then it can be terminated by sending the atom `end_of_file` or `halt` to the standard input of the child. (For this to work, the child XSB must be waiting at the prompt).
3. It is very important to not forget to close the streams that the parent uses to communicate with the child. These are the streams that are provided in arguments 2,3,4 of `spawn_process`. The reason is that the child might terminate, but these streams to the standard input of the child will remain open, since they belong to the parent process. As a result, the parent will own defunct I/O streams and might eventually run out of file descriptors or streams.

`process_status(+Pid,-Status)`

This predicate always succeeds. Given a process id, it binds the second argument (which must be an unbound variable) to one of the following atoms: `running`, `stopped`, `exited_normally`, `exited_abnormally`, `aborted`, `invalid`, and `unknown`. The `invalid` status is given to processes that never existed or that are not children of the parent XSB process. The `unknown` status is assigned when none of the other statuses can be assigned.

Note: process status (other than `running`) is system dependent. Windows does not seem to support `stopped` and `aborted`. Also, processes killed using the `process_control` predicate (described next) are often marked as `invalid` rather than `exited`, because Windows seems to lose all information about such processes. Process status might be inaccurate in some Unix systems as well, if the process has terminated and `wait()` has been executed on that process.

`process_control(+Pid,+Operation)`

Perform a process control `operation` on the process with the given `Pid`. Currently, the only supported operations are `kill` (an atom) and `wait(Code)` (a term). The former causes the process to exit unconditionally, and the latter waits for process completion. When the process exits, `Code` is bound to the process exit code. The code for normal termination is 0.

This predicate succeeds, if the operation was performed successfully. Otherwise, it fails. The `wait` operation fails if the process specified in `Pid` does not exist or is not a child of the parent XSB process.

The `kill` operation might fail, if the process to be killed does not exist or if the parent XSB process does not have the permission to terminate that process. Unix and Windows have

different ideas as to what these permissions are. See *kill(2)* for Unix and *TerminateProcess* for Windows.

Note: under Windows, the programmer’s manual warns of dire consequences if one kills a process that has DLLs attached to it.

`get_process_table(-ProcessList)` module: `shell`

This predicate is imported from module `shell`. It binds `ProcessList` to the list of terms, each describing one of the active XSB subprocesses (created via `spawn_process/5`). Each term has the form:

`process(Pid,ToStream,FromStream,StderrStream,CommandLine).`

The first argument in the term is the process id of the corresponding process, the next three arguments describe the three standard streams of the process, and the last is an atom that shows the command line used to invoke the process. This predicate always succeeds.

`shell(+CmdSpec,-StreamToProc, -StreamFromProc, -ProcStderr, -ErrorCode)`

The arguments of this predicate are similar to those of `spawn_process`, except for the following: (1) The first argument is an atom or a list of atoms, like in `spawn_process`. However, if it is a list of atoms, then the resulting shell command is obtained by string concatenation. This is different from `spawn_process` where each member of the list must represent an argument to the program being invoked (and which must be the first member of that list). (2) The last argument is the error code returned by the shell command and not a process id. The code -1 and 127 mean that the shell command failed.

The `shell/5` predicate is similar to `spawn_process` in that it spawns another process and can capture that process’ input and output streams. The important difference, however, is that XSB will wait until the process spawned by `shell/5` terminates. In contrast, the process spawned by `spawn_process` will run concurrently with XSB. In this latter case, XSB must explicitly synchronize with the spawned subprocess using the predicate `process_control/2` (using the `wait` operation), as described earlier.

The fact that XSB must wait until `shell/5` finishes has a very important implication: the amount of data that can be sent to and from the shell command is limited (1K is probably safe). This is because the shell command communicates with XSB via pipes, which have limited capacity. So, if the pipe is filled, XSB will hang waiting for `shell/5` to finish and `shell/5` will wait for XSB to consume data from the pipe. Thus, use `spawn_process/5` for any kind of significant data exchange between external processes and XSB.

Another difference between these two forms of spawning subprocesses is that `CmdSpec` in `shell/5` can represent *any* shell statement, including those that have pipes and I/O redirection. In contrast, `spawn_process` only allows command of the form “program args”. For instance,

```
| ?- open(test,write,Stream),
    shell('cat | sort > data', Stream, none, none, ErrCode)
```

As seen from this example, the same rules for blocking I/O streams apply to `shell/5`. Finally, we should note that the already familiar standard predicates `shell/1` and `shell/2` (documented in Volume 1) are implemented using `shell/5`, and `shell/5` shares their error cases.

Notes:

1. With `shell/5`, you do not have to worry about terminating child processes: XSB waits until the child exits automatically. However, since communication pipes have limited capacity, this method can be used only for exchanging small amounts of information between parent and child.
2. The earlier remark about the need to close I/O streams to the child *does* apply.

`pipe_open(-ReadPipe, -WritePipe)`

Open a new pipe and return the read end and the write end of that pipe. If the operation fails, both `ReadPipe` and `WritePipe` are bound to negative numbers. The pipes returned by the `pipe_open/2` predicate are small integers that represent file descriptors used by the underlying OS. They are **not XSB I/O streams**, and they cannot be used for I/O directly. To use them, one must convert them to streams using `open/3` or `open/4`.⁷

The best way to illustrate how one can create a new pipe to a child (even if the child has been created earlier) is to show an example. Consider two programs, `parent.P` and `child.P`. The parent copy of XSB consults `parent.P`, which does the following: First, it creates a pipe and spawns a copy of XSB. Then it tells the child copy of XSB to assert the fact `pipe(RP)`, where `RP` is a number representing the read part of the pipe. Next, the parent XSB tells the child XSB to consult the program `child.P`. Finally, it sends the message `Hello!`.

The `child.P` program gets the pipe from predicate `pipe/1` (note that the parent tells the child XSB to first assert `pipe(RP)` and only then to consult the `child.P` file). After that, the child reads a message from the pipe and prints it to its standard output. Both programs are shown below:

```
%% parent.P
:- import pipe_open/2 from file_io.
%% Create the pipe and pass it to the child process
?- pipe_open(RP,WP),
   %% WF is now the XSB I/O stream bound to the write part of the pipe
   open(pipe(WP),write,WF),
   %% ProcInput becomes the XSB stream leading directly to the child's stdin
   spawn_process(nxsb1, ProcInput, block, block, Process),
   %% Tell the child where the reading part of the pipe is
   fmt_write(ProcInput, "assert(pipe(%d)).\n", arg(RP)),
   fmt_write(ProcInput, "[child].\n", _),
```

⁷ XSB does not convert pipe file descriptors into I/O streams automatically. Because of the way XSB I/O streams are represented, they are not inherited by the child process and they do not make sense to the child process (especially if the child is not another XSB process). Therefore, we must pass the child processes an OS file descriptor instead. The child then converts these descriptor into XSB I/O streams.

```

flush_output(ProcInput, _),
%% Pass a message through the pipe
fmt_write(WF, "Hello!\n", _),
flush_output(WF, _),
fmt_write(ProcInput, "end_of_file.\n",_), % send end_of_file atom to child
flush_output(ProcInput, _),
%% wait for child (so as to not leave zombies around;
%% zombies quit when the parent finishes, but they consume resources)
process_control(Process, wait),
%% Close the ports used to communicate with the process
%% Otherwise, the parent might run out of file descriptors
%% (if many processes were spawned)
close(ProcInput), close(WF).

%% child.P
:- import file_read_line_atom/2 from file_io.
:- dynamic pipe/1.
?- pipe(P), open(pipe(P),read,F),
    %% Acknowledge receipt of the pipe
    fmt_write("\nPipe %d received\n", arg(P)),
    %% Get a message from the parent and print it to stdout
    file_read_line_atom(F, Line), write('Message was: '), writeln(Line).

```

This produces the following output:

```

| ?- [parent].                <- parent XSB consults parent.P
[parent loaded]
yes
| ?- [xsb_configuration loaded] <- parent.P spawns a child copy of XSB
[sysinitrc loaded]             Here we see the startup messages of
[packaging loaded]             the child copy
XSB Version 2.0 (Gouden Carolus) of June 27, 1999
[i686-pc-linux-gnu; mode: optimal; engine: slg-wam; scheduling: batched]
| ?-
yes
| ?- [Compiling ./child]       <- The child copy of received the pipe from
[child compiled, cpu time used: 0.1300 seconds] the parent and then the
[child loaded]                 request to consult child.P
Pipe 15 received               <- child.P acknowledges receipt of the pipe
Message was: Hello!            <- child.P gets the message and prints it
yes

```

Observe that the parent process is very careful about making sure that the child terminates and also about closing the I/O streams after they are no longer needed.

Finally, we should note that this mechanism can be used to communicate through pipes with non-XSB processes as well. Indeed, an XSB process can create a pipe using `pipe_open` (*before* spawning a child process), pass one end of the pipe to a child process (which can be a C program), and use `open/3` to convert the other end of the pipe to an XSB stream. The C program, of course, does not need `open/3`, since it can use the pipe file handle directly. Likewise, a C program can spawn off an XSB process and pass it one end of a pipe. The XSB child-process can then convert this pipe fd to a file using `fd2iostream` and then talk to the parent C program.

`fd2iostream(+Pipe, -IOstream)`

Take a file descriptor and convert it to an XSB I/O stream. This predicate should be used only for user-defined I/O. Otherwise, use `open/{3,4}` when possible.

1.7 Socket I/O

The XSB socket library defines a number of predicates for communication over BSD-style sockets. Most are modeled after and are interfaces to the socket functions with the same name. For detailed information on sockets, the reader is referred to the Unix man pages (another good source is *Unix Network Programming*, by W. Richard Stevens). Several examples of the use of the XSB sockets interface can be found in the `XSB/examples/` directory in the XSB distribution.

XSB supports two modes of communication via sockets: *stream-oriented* and *message-oriented*. In turn, stream-oriented communication can be *buffered* or *character-at-a-time*.

To use *buffered* stream-oriented communication, system socket handles must be converted to XSB I/O streams using `fd2iostream/2`. In these stream-oriented communication, messages have no boundaries, and communication appears to the processes as reading and writing to a file. At present, buffered stream-oriented communication works under Unix only.

Character-at-a-time stream communication is accomplished using the primitives `socket_put/3` and `socket_get0/3`. These correspond to the usual Prolog `put/1` and `get0/1` I/O primitives.

In message-oriented communication, processes exchange messages that have well-defined boundaries. The communicating processes use `socket_send/3` and `socket_recv/3` to talk to each other. XSB messages are represented as strings where the first four bytes (`sizeof(int)`) is an integer (represented in the binary network format — see the functions `htonl` and `ntohl` in socket documentation) and the rest is the body of the message. The integer in the header represents the length of the message body.

Effort has been made to make the socket interface thread-safe; however in Version 3.4, calls to the XSB socket interface go through a single mutex, and may cause contention if many threads seek to concurrently use sockets.

We now describe the XSB socket interface. All predicates below must be imported from the module `socket`. Note that almost all predicates have the last argument that unifies with the error

code returned from the corresponding socket operation. This argument is explained separately.

General socket calls. These are used to open/close sockets, to establish connections, and set special socket options.

`socket(-Sockfd, ?ErrorCode)`

A socket `Sockfd` in the `AF_INET` domain is created. (The `AF_UNIX` domain is not yet implemented). `Sockfd` is bound to a small integer, called socket descriptor or socket handle.

`socket_set_option(+Sockfd,+OptionName,+Value)`

Set socket option. At present, only the `linger` option is supported. “Lingering” is a situation when a socket continues to live after it was shut down by the owner. This is used in order to let the client program that uses the socket to finish reading or writing from/to the socket. `Value` represents the number of seconds to linger. The value -1 means do not linger at all.

`socket_close(+Sockfd, ?ErrorCode)`

`Sockfd` is closed. Sockets used in `socket_connect/2` should not be closed by `socket_close/1` as they will be closed when the corresponding stream is closed.

`socket_bind(+Sockfd,+Port, ?ErrorCode)`

The socket `Sockfd` is bound to the specified local port number.

`socket_connect(+Sockfd,+Port,+Hostname,?ErrorCode)`

The socket `Sockfd` is connected to the address (`Hostname` and `Port`). If `socket_connect/4` terminates abnormally for any reason (connection refused, timeout, etc.), then XSB closes the socket `Sockfd` automatically, because such a socket cannot be used according to the BSD semantics. Therefore, it is always a good idea to check to the return code and reopen the socket, if the error code is not `SOCK_OK`.

`socket_listen(+Socket, +Length, ?ErrorCode)`

The socket `Sockfd` is defined to have a maximum backlog queue of `Length` pending connections.

`socket_accept(+Sockfd,-SockOut, ?ErrorCode)`

Block the caller until a connection attempt arrives. If the incoming queue is not empty, the first connection request is accepted, the call succeeds and returns a new socket, `SockOut`, which can be used for this new connection.

Buffered, message-based communication. These calls are similar to the `recv` and `send` calls in C, except that XSB wraps a higher-level message protocol around these low-level functions. More precisely, `socket_send/3` prepends a 4-byte field to each message, which indicates the length of the message body. When `socket_recv/3` reads a message, it first reads the 4-byte field to determine the length of the message and then reads the remainder of the message.

All this is transparent to the XSB user, but you should know these details if you want to use these details to communicate with external processes written in C and such. All this means that these external programs must implement the same protocol. The subtle point here is that different

machines represent integers differently, so an integer must first be converted into the machine-independent network format using the functions `htonl` and `ntohl` provided by the socket library. For instance, to send a message to XSB, one must do something like this:

```
char *message, *msg_body;
unsigned int msg_body_len, network_encoded_len;

msg_body_len = strlen(msg_body);
network_encoded_len = (unsigned int) htonl((unsigned long int) msg_body_len);
memcpy((void *) message, (void *) &network_encoded_len, 4);
strcpy(message+4, msg_body);
```

To read a message sent by XSB, one can do as follows:

```
int actual_len;
char lenbuf[4], msg_buff;
unsigned int msglen, net_encoded_len;

actual_len = (long)recvfrom(sock_handle, lenbuf, 4, 0, NULL, 0);
memcpy((void *) &net_encoded_len, (void *) lenbuf, 4);
msglen = ntohl(net_encoded_len);

msg_buff = calloc(msglen+1, sizeof(char)); // check if this succeeded!!!
recvfrom(sock_handle, msg_buff, msglen, 0, NULL, 0);
```

If making the external processes follow the XSB protocol is not practical (because you did not write these programs), then you should use the character-at-a-time interface or, better, the buffered stream-based interface both of which are described in this section. At present, however, the buffered stream-based interface does not work on Windows.

socket_recv(+Sockfd,-Message, ?ErrorCode)

Receives a message from the connection identified by the socket descriptor `Sockfd`. Binds `Message` to the message. `socket_recv/3` provides a message-oriented interface. It understands message boundaries set by `socket_send/3`.

socket_send(+Sockfd,+Message, ?ErrorCode)

Takes a message (which must be an atom) and sends it through the connection specified by `Sockfd`. `socket_send/3` provides message-oriented communication. It prepends a 4-byte header to the message, which tells `socket_recv/3` the length of the message body.

Stream-oriented, character-at-a-time interface. Internally, this interface uses the same `sendto` and `recvfrom` socket calls, but they are executed for each character separately. This interface is appropriate when the message format is not known or when message boundaries are determined using special delimiters.

`socket_get0/3` creates the end-of-file condition when it receives the end-of-file character `CH_EOF_P` (a.k.a. 255) defined in `char_defs.h` (which must be included in the XSB program). C programs that need to send an end-of-file character should send `(char)-1`.

`socket_get0(+Sockfd, -Char, ?ErrorCode)`

The equivalent of `get0` for sockets.

`socket_put(+Sockfd, +Char, ?ErrorCode)`

Similar to `put/1`, but works on sockets.

Socket-probing. With the help of the predicate `socket_select/6` one can establish a group of asynchronous or synchronous socket connections. In the synchronous mode, this call is blocked until one of the sockets in the group becomes available for reading or writing, as described below. In the asynchronous mode, this call is used to probe the sockets periodically, to find out which sockets have data available for reading or which sockets have room in the buffer to write to.

The directory `XSB/examples/socket/select/` has a number of examples of the use of the socket-probing calls.

`socket_select(+SymConName,+Timeout,-ReadSockL,-WriteSockL,-ErrSockL,?ErrorCode)`

`SymConName` must be an atom that denotes an existing connection group, which must be previously created with `socket_set_select/4` (described below). `ReadSockL`, `WriteSockL`, `ErrSockL` are lists of socket handles (as returned by `socket/2`) that specify the available sockets that are available for reading, writing, or on which exception conditions occurred. `Timeout` must be an integer that specifies the timeout in seconds (0 means probe and exit immediately). If `Timeout` is a variable, then wait indefinitely until one of the sockets becomes available.

`socket_set_select(+SymConName,+ReadSockFdLst,+WriteSockFdLst,+ErrorSockFdLst)`

Creates a connection group with the symbolic name `SymConName` (an atom) for subsequent use by `socket_select/6`. `ReadSockFdLst`, `WriteSockFdLst`, and `ErrorSockFdLst` are lists of sockets for which `socket_select/6` will be used to monitor read, write, or exception conditions.

`socket_select_destroy(+SymConName)`

Destroys the specified connection group.

Error codes. The error code argument unifies with the error code returned by the corresponding socket commands. The error code -2 signifies *timeout* for timeout-enabled primitives (see below). The error code of zero signifies normal termination. Positive error codes denote specific failures, as defined in BSD sockets. When such a failure occurs, an error message is printed, but the predicate succeeds anyway. The specific error codes are part of the socket documentation. Unfortunately, the symbolic names and error numbers of these failures are different between Unix compilers and Visual C++. Thus, there is no portable, reliable way to refer to these error codes. The only reliably portable error codes that can be used in XSB programs defined through these symbolic constants:

```

#include "socket_defs_xsb.h"

#define SOCK_OK      0      /* indicates successful return from socket */
#define SOCK_EOF     -1     /* end of file in socket_recv, socket_get0 */

#include "timer_defs_xsb.h"

#define TIMEOUT_ERR -2      /* Timeout error code */

```

Timeouts. XSB socket interface allows the programmer to specify timeouts for certain operations. If the operation does not finish within the specified period of time, the operation is aborted and the corresponding predicate succeeds with the `TIMEOUT_ERR` error code. The following primitives are timeout-enabled: `socket_connect/4`, `socket_accept/3`, `socket_recv/3`, `socket_send/3`, `socket_get0/3`, and `socket_put/3`. To set a timeout value for any of the above primitives, the user should execute `set_timer/1` right before the subgoal to be timed. Note that timeouts are disabled after the corresponding timeout-enabled call completes or times out. Therefore, one must use `set_timer/1` before each call that needs to be controlled by a timeout mechanism.

The most common use of timeouts is to either abort or retry the operation that times out. For the latter, XSB provides the `sleep/1` primitive, which allows the program to wait for a few seconds before retrying.

The `set_timer/1` and `sleep/1` primitives are described below. They are standard predicates and do not need to be explicitly imported.

`set_timer(+Seconds)`

Set timeout value. If a timer-enabled goal executes after this value is set, the clock begins ticking. If the goal does not finish in time, it succeeds with the error code set to `TIMEOUT_ERR`. The timer is turned off after the goal executes (whether timed out or not and whether it succeeds or fails). This goal always succeeds.

Note that if the timer is not set, the timer-enabled goals execute “normally,” without timeouts. In particular, they might block (say, on `socket_recv`, if data is not available).

`sleep(+Seconds)`

Put XSB to sleep for the specified number of seconds. Execution resumes after the `Seconds` number of seconds. This goal always succeeds.

Here is an example of the use of the timer:

```
:- compiler_options([xpp_on]).
#include "timer_defs_xsb.h"

?- set_timer(3), % wait for 3 secs
   socket_recv(Sockfd, Msg, ErrorCode),
   (ErrorCode == TIMEOUT_ERR
   -> writeln('Socket read timed out, retrying'),
       try_again(Sockfd)
   ; write('Data received: '), writeln(Msg)
   ).
```

Apart from the above timer-enabled primitives, a timeout value can be given to `socket_select/6` directly, as an argument.

Buffered, stream-oriented communication. In Unix, socket descriptors can be “promoted” to file streams and the regular read/write commands can be used with such streams. In XSB, such promotion can be done using the following predicate:

```
fd2ioport(+Pipe, -IOport)                                module: shell
    Take a socket descriptor and convert it to an XSB I/O port that can be used for regular file
    I/O.
```

Once `IOport` is obtained, all normal I/O primitives can be used by specifying the `IOport` as their first argument. This is, perhaps, the easiest and the most convenient way to use sockets in XSB. (This feature has not been implemented for Windows.)

Here is an example of the use of this feature:

```
:- compiler_options([xpp_on]).
#include "socket_defs_xsb.h"

?- (socket(Sockfd, SOCK_OK)
   -> socket_connect(Sockfd1, 6020, localhost, Ecode),
       (Ecode == SOCK_OK
       -> fd2ioport(Sockfd, SockIOport),
           file_write(SockIOport, 'Hello Server!')
       ; writeln('Can''t connect to server')
       ),
   ; writeln('Can''t open socket'), fail
   ).
```

The module `array1` provides a simple backtrackable array implementation that requires no copying. In Version 3.2, this package was changed to make use of the backtrackable destructive assignment made possible by `setarg/3`. We note that as of Version 3.2 this library provides simple syntactic sugar for `functor/3`, `arg/3` and `setarg/3` and relies on error messages for these predicates.

<code>array_update(+Array, +Index, +Elem)</code>	module: <code>array</code>
<p>Updates the array <code>Array</code> such that the <code>Index</code>-th element of the new array is <code>Elem</code> using destructive assignement. The implementation is quite efficient in that it avoids the copying of the entire array.</p>	

no

XSB can provide Prolog-level profiling for Prolog programs, which allows the Prolog programmer to estimate what proportion of time is spent executing code for each predicate, and also what modes have been used to call a given predicate. It also helps to find unindexed accesses to dynamic predicates which may be the cause of poor performance. To enable profiling, XSB must be started with the command line parameter of `-p`. The module `xsb_profiling` contains the predicate `profile_call/1` that invokes profiling. The profiling library should only be used with the single-threaded engine in Version 3.4.

```
profile_call(+Goal)                                module:  xsb_profiling
```

Calls `Goal`, and when it first succeeds, prints to `userout` a table of predicate names indicating for each, the percentage of time spent executing that predicate’s code. Within the table, the sum of the predicate times for each module is also given. `Goal` may backtrack, but profiling is done only for the time to the first success, so it is most appropriate to profile succeeding deterministic goals ⁸.

Profiling works by starting another thread that interrupts every 100th of a second and sets a flag so that the XSB emulator will determine the predicate of the currently executing code. The printout also includes the total number of interrupts and for each predicate, the raw number of times its code was determined to be executing. A predicate is printed only if its code was interrupted at least once. The numbers will be meaningful only for relatively long-running predicates, taking more than a couple of seconds.

When an interrupt occurs, the *next interrupt instruction* to be executed – a WAM call, `execute`, `proceed` or `trust` instruction – will charge its associated predicate by logging that predicate to a table. The system does not keep track of code addresses for tries (used to represent the results of completed tables, and trie-indexed asserted code), so for some interrupts the associated executing predicate cannot be determined. In these cases the interrupt is charged against an “unknown/?” pseudo-predicate, and this count is included in the output.

Profiling does not give the context from which the predicate is called, so you may want to make renamed copies of basic predicates to use in particular circumstances to determine their times.

Predicates compiled with the “optimize” option may provide misleading results under profiling. Note that all system predicates (including those in `basics`) are compiled with the “optimize” option, by default. That option causes tail-recursive predicates to use a “jump” instruction rather than an “execute” instruction to make the recursive call, and so an interrupt in such a loop will not be charged until the next interrupt instruction is executed. If much time is spent in the recursion, this might not be for a long time, and the interrupt might ultimately be charged to another predicate. (If an interrupt has not been charged by the time of the next interrupt, it is lost.)

Profiling is currently available under Windows, Mac OS X, and Linux. However, for the profiling algorithm to provide a good estimation, the thread that wakes and sets the interrupt flag must be of high priority and given the CPU when it wants it. Accordingly, the estimates may be better or worse depending on the scheduling strategy of a given platform ⁹.

The profiling module also provides support for determining when a dynamic predicate is invoked in a mode that isn’t supported by any index. The XSB programmer can set a flag that will cause a message to be printed when a dynamic predicate is invoked, no index is applicable, and there are more than 20 potentially matching clauses. See `profile_unindexed_calls/1` below for details.

```
profile_mode_call(+Goal)                                module:  xsb_profiling
```

⁸This includes tabled subgoals under Local Evaluation, as such as goal will only succeed after deriving all of its answers.

⁹Windows and Mac OS X 10.6 provide good estimates. Some Linuxes however, do not charge about 20% of their interrupts due to thread scheduling issues. This loss of interrupts makes the profile estimate inefficient, but does not bias the estimate. We haven’t figured out how to get priority scheduling for interrupts on all machines, so if you want profiling to work more efficiently, maybe you can help figure out how to get appropriate scheduling.

Calls the goal `Goal` and constructs a table of the modes in which the predicate is called and the number of times it is called in that mode. Modes are simply “b” for ground and “f” for variable. Counts are kept in a table with entries of the form `Pred(Md1,Md2,...,Mdn)` where `Pred` is the name of the called predicate and the `Mdi` are either ‘f’ or ‘b’, indicating free or bound for the corresponding argument. The table can be printed using `profile_mode_dump/0` and can be cleared using `profile_mode_init/0`.

`profile_mode_dump` module: `xs_b_profiling`

Prints out the counts of calls in particular modes as accumulated using `profile_mode_call(+Goal)`.

`profile_mode_init` module: `xs_b_profiling`

Clears the table that accumulates counts of calls in particular modes (done by `profile_mode_call(+Goal)`).

`profile_unindexed_calls(+Par)` module: `xs_b_profiling`

Sets the kind of unindexed profiling to perform. If `Par` is 0 no unindexed logging will be done. (This is the default.) If `Par` is 1 each call to a dynamic predicate that cannot use any index (and would backtrack through more than 20 clauses) will generate a log message to `userout`. Each unindexed call to a predicate will be logged only once; after logging is done, the log instruction is changed to a branch, so it will never produce another log message for that dynamic code. If `Par` is 2, logging is done as for 1, except every unindexed call to any dynamic predicate will be logged; i.e. the logging instruction is not changed after logging. If `Par` is a predicate specification (of the form `Pred/Arity`, `Module:Pred/Arity`, `Term`, or `Module:Term`), only unindexed calls to the indicated goal will be logged, and when each is logged a back-trace will be printed. This allows the programmer to find the location of an unindexed call.

1.10 Gensym

The Gensym library provides a convenient way to generate unique integers or constants.

`prepare(+Index)` module: `gensym`

Sets the initial integer to be used for generation to `Index`. Thus, the command `?- prepare(0)` would cause the first call to `gennum/1` to return 1. `Index` must be a non-negative integer.

`gennum(-Var)` module: `gensym`

Unifies `Var` with a new integer.

`gensym(+Atom,-Var)` module: `gensym`

Generates a new integer, and concatenates this integer with `Atom`, unifying the result with `Var`. For instance a call `?- gensym(foo,Var)` might unify `Var` with `foo32`.

1.11 Random Number Generator

The following predicates are provided in module `random` to generate random numbers (both integers and floating numbers), based on the Wichmann-Hill Algorithm [?, ?]. The random number generator is entirely portable, and does not require any calls to the operating system. As noted below, it does require 3 seeds, each of which must be an integer in a given range. These seeds are thread-specific: thus different threads may generate independent sequences of random numbers.

`random(-Number)` module: `random`
 Binds `Number` to a random float in the interval [0.0, 1.0). Note that 1.0 will never be generated.

`random(+Lower,+Upper,-Number)` module: `random`
 Binds `Number` to a random integer in the interval [`Lower`,`Upper`) if `Lower` and `Upper` are integers. Otherwise `Number` is bound to a random float between `Lower` and `Upper`. `Upper` will never be generated.

`getrand(?State)` module: `random`
 Tries to unify `State` with the term `rand(X,Y,Z)` where `X`,`Y`, and `Z` are integers describing the state of the random generator.

`setrand(rand(+X,+Y,+Z))` module: `random`
 Sets the state of the random generator. `X`,`Y`, and `Z` must be integers in the ranges [1,30269), [1,30307), [1,30323), respectively.

`datetime_setrand` module: `random`
 This simple initialization utility sets the random seed triple based on a function of the current day, hour, minute and second.

`randseq(+K, +N, -RandomSeq)` module: `random`
 Generates a sequence of `K` unique integers chosen randomly in the range from 1 to `N`. `RandomSeq` is not returned in any particular order.

`randset(+K, +N, -RandomSet)` module: `random`
 Generates an ordered set of `K` unique integers chosen randomly in the range from 1 to `N`. The set is returned in reversed order, with the largest element first and the smallest last.

`gauss(-G1,-G2)` module: `random`
 Generates two random numbers that are normally distributed with mean 0 and standard deviation 1. It uses the polar form of the Box-Muller transformation [?] of uniform random variables as generated by `random/1`.

`weibull(K,Lambda,X)` module: `random`
 Generates a random number for the Weibull distribution:

$$f(x; k, \lambda) = \frac{k}{\lambda} \left(\frac{x}{\lambda}\right)^{k-1} e^{-(x/\lambda)^k}$$

based on the transformation

$$x = \lambda(-\ln(U))^{1/k}$$

of a uniformly distributed random variable produced by `random/1`

`exponential(K,X)`

module: `random`

Generates a random number for the exponential distribution:

$$f(x; k, \lambda) = \frac{e^{-(x/\lambda)^k}}{\lambda}$$

based on the transformation

$$x = \lambda(-\ln(U))$$

of a uniformly distributed random variable produced by `random/1`. This is the same as the Weibull distribution with $k = 1$.

1.12 Loading Separated Files

A common file format uses comma separated values, the so-called csv files. The XSB module, `proc_files`, supports the loading of files in this, and similar, formats to define Prolog predicates.

`load_csv(+FileName,+PredSpec)`

module: `proc_files`

`load_csv/2` takes a file name and a predicate specification, and reads a csv-formatted file into memory, defining the indicated dynamic predicate. The simplest form of `PredSpec` is `PredName/Arity`. In this case the arity must equal the number of fields in the csv file, and the predicate must be dynamic. Each line in the file will define one fact of the predicate `PredName/Arity`. Fields in the file enclosed in double quotes will be treated as single fields (and thus can contain commas and new-lines.) The dynamic predicate will be emptied before the facts from the file are added. Each field will be loaded as an atom (including fields that contain just integers.)

Alternatively, `PredSpec` may be of the form `predName(TypeSpec1,...,TypeSpecN)`, where `predName` is the name of the dynamic predicate to be defined by the file contents, and each `TypeSpecI` indicates the type of the corresponding field in the file. The permitted values of `TypeSpec` are:

- `atom` The corresponding field value will become an atom in the loaded fact.
- `integer` The corresponding field value will be converted to an integer in the loaded fact.
- `float` The corresponding field value will be converted to a float in the loaded fact.
- `term` The corresponding field must contain a Prolog term in canonical form, and it will be converted to that term in the loaded fact.
- `___` (A variable) Treated as `atom`.

```
load_dsv(+FileName,+PredSpec,+Options)                                module:  proc_files
```

This predicate supports the loading of more general forms of files with value-separated fields. The `FileName` and `PredSpec` parameters are exactly as in `load_csv/2`, as described just above. `Options` is a list of options. (With an empty list, `load_dsv` acts as `load_csv/2`.) The options are:

`separator="Sep"` which indicates that the character(s) `Sep` will be used as the field separator. There may be one or more characters.

`delimiter="C"` which indicates that the single character `C` will be used as the field delimiter (the default being `"'"`, and I've yet to find a situation in which I want to change it.)

`titles` which indicates that the first line of the file should be ignored and not contribute a fact to the dynamic predicate.

1.13 Scanning in Prolog

Scanners, (sometimes called tokenizers) take an input string, usually in ASCII or similar format, and produce a scanned sequence of tokens. The requirements that various applications have for scanning differ in small but important ways – a character that is special to one application may be part of the token of another; or some applications may want lower case text converted to upper-case text. The `stdscan.P` library provides a simple scanner written in XSB that can be configured in several ways. While useful, this scanner is not intended to be as powerful as general-purpose scanners such as *lex* or *flex*.

```
scan(+List,-Tokens)                                                  module:  stdscan
```

Given as input a `List` of character codes, `scan/2` scans this list producing a list of atoms constituting the lexical tokens. Its parameters are set via `set_scan_pars/1`.

Tokens produced are either a sequence of *letters* and/or *numbers* or consist of a single *special character* (e.g. `(` or `)`). Whitespaces may occur between tokens.

```
scan(+List,+FieldSeparator,-Tokens)                                module:  stdscan
```

Given as input a `List` of character codes, along with a character code for a field separator, `scan/3` scans this list producing a list of list of atoms constituting the lexical tokens in each field. `scan/3` thus can be used to scan tabular information. Its parameters are set via `set_scan_pars/1`.

```
set_scan_pars(+List)                                                module:  stdscan
```

`set_scan_pars(+List)` is used to configure the tokenizer to a particular need. `List` is a list of parameters including the following:

- **whitespace.** The default action of the scanner is to return a list of tokens, with any whitespace removed. If **whitespace** is a parameter, then the scanner returns the token " when it finds whitespace separating two tokens (unless the two tokens are letter sequences; since two letter sequences can be two tokens ONLY if they are separated by whitespace, such an indication of whitespace would be redundant.) Including the parameter **no_whitespace** undoes the effect of previously including **whitespace**.
- **upper_case** The default action of the parser is to treat lowercase letter differently from uppercase letters. This parameter should be set if conversion to uppercase should be done when producing a token that does *not* consist entirely of letters (e.g. one with mixed letters and digits). Including the parameter **no_case** undoes the effect of previously including **upper_case**.
- **upper_case_in_lit** The default action of the parser is to treat lowercase letter differently from uppercase letters. This parameter should be set if conversion to uppercase should be done when producing a token that consists entirely of letters. Including the parameter **no_case_in_lit** undoes the effect of previously including **upper_case**.
- **whitespace(Code)** adds **Code** as a whitespace code. By default, all ASCII codes less than or equal to 32 are regarded as whitespace.
- **letter(Code)** adds **Code** as a letter constituting a token. By default, ASCII codes for characters a-z and A-Z are regarded as letters.
- **special_char(Code)** adds **Code** as a special character. By default, ASCII codes for the following characters are regarded as special characters:

| { } [] " \$ % & ' () * + , - . / : ; < = > ? @ \ ^ _ ~ `

`get_scan_pars(-List)` module: stdscan
`get_scan_pars/1` returns a list of the currently active parameters.

1.14 XSB Lint

The `xsb_lint_impexp.P` file contains a simple tool to analyze import/exports and definitions and uses of predicates. It tries to find possible inconsistencies, producing warnings when it finds them and generating `document_import/document_export` declarations that might be useful. It can be used after a large multi-file, multi-module XSB program has been written to find possible inconsistencies in (or interesting aspects of) how predicates are defined and used.

XSB source files that contain an `export` compiler directive are considered as modules. Predicates defined in modules, but not exported, are local to that module. When compiling a module, the XSB compiler generates useful warnings when predicates are used but not defined or defined but not used. All predicates that are defined in source files that do not contain an `export` directive are compiled to be defined in a global module, called `usermod`, and no warning messages are generated. The user may add `document_export` and `document_import` compiler directives (exactly analogous to the `export` and `import` directives) to non-module source files. These directives are ignored by the compiler for its compilation, but cause the define-use analysis to be done and any

warning messages to be issued, if appropriate. This allows a user to get the benefit of the define-use analysis without using modules. (See Volume 1, Chapter 3 for more details.)

The `xsb_lint_impexp` utility processes both modules and regular XSB source files that contain `document_export` statements. `xsb_lint_impexp` is not itself a module. To use it, `[xsb_lint_impexp]` must be consulted, which will define the `checkImpExps/{1,2}` and `add_libraries/1` predicates in `usermod`.

`add_libraries(+DirectoryNameList)`

`add_libraries/1` takes a list of directory names and adds them to the `library_directory/1` predicate. This causes the XSB system to look for XSB source code files in these directories. To use `checkImpExps/{1,2}`, all the directories that contain files (or modules) referenced (recursively) in the files to be processed must be in the `library_directory/1` predicate. This predicate can be used to add a number of directories at once.

`checkImpExps(+Options,+FileNameList)`

`checkImpExps/1` reads all the XSB source files named in the list `FileNameList`, and all files they reference (recursively), and produces a listing that describes properties of how they reference predicates.

`Options` is a list of atoms (from the following list) indicating details of how `checkImpExps` should work.

1. `used_elsewhere`: Print a warning message in the case of a predicate defined in a file, not used there, but used elsewhere (in a file in `FileNameList`). This can be useful to see whether it might be better to move the predicate definition to another file, but it produces many warnings for predicates in multi-use libraries.
2. `unused`: Print a warning message in the case of a predicate that is exported but never used. This can be useful to see if predicate is not used anywhere, and thus could be deleted. Again this produces many warnings for predicates in multi-use libraries.
3. `all_files`: By default, only predicates in files that contain a `:- document_export` or `:- export` declaration are processed. This option causes predicates of *all* files (and modules) to be processed.
4. `all_symbol_uses`: Treat *all* non-predicate uses of symbols (even constants) as predicate uses for the purpose of generating imports.
5. `no_symbol_uses`: Don't treat any non-predicate uses of symbols as predicate uses for the purpose of generating imports.

We further explain the final two options, which allow the user to determine more precisely what uses of a symbol are considered as uses of it as the predicate symbol. All uses of symbols that appear in a “predicate context”, i.e., in the body of a rule or in a meta-predicate argument position of a use of a meta-predicate, are considered uses of that predicate symbol. The default is also to allow nonconstant symbols appearing in any other context to also count as uses of that symbol as that predicate symbol. This is useful for programs that define their own meta-predicates.

`checkImpExps(+FileNameList)`

`checkImpExps/1` is currently equivalent to `checkImpExps([],FileNameList)`.

1.15 Miscellaneous Predicates

`term_hash(+Term,+HashSize,-HashVal)`

module: `machine`

Given an arbitrary Prolog term, `Term`, that is to be hashed into a table of `HashSize` buckets, this predicate returns a hash value for `Term` that is between 0 and `HashSize - 1`.

`pretty_print(+ClausePairs)`

module: `pretty_print`

`pretty_print(+Stream,+ClausePairs)`

module: `pretty_print`

The input to `pretty_print/1`, `ClausePairs`, can be either a list of clause pairs or a single clause pair. A clause pair is either a Prolog clause (or declaration) or a pair:

(`Clause`,`Dict`)

Where `Dict` is a list of the form `A = V` where `V` is a variable in `Clause` and `A` is the string to be used to denote the variable ¹⁰.

By default, `pretty_print/1` outputs atomic terms using `writeln/1`, but specialized output can be configured via asserting in `usermod` a term of the form

`user_replacement_hook(Term,Call)`

which will use `Call` to output an atomic literal `A` whenever `A` unifies with `Term`. For example, pretty printing weight constraints in XSB's XASP package is done via the hook

`user_replacement_hook(weight_constr(Term),output_weight_constr(Term))`

which outputs a weight constraint in a (non-Prolog) syntax that is used by several ASP systems.

1.16 Other Libraries

Not all XSB libraries are fully documented. We provide brief summaries of some of these other libraries.

¹⁰Thus the list of variable names returned by `read_term/{2,3}` can be used directly in `Dict`.

1.16.1 Justification

By Hai-Feng Guo

Most Prolog debuggers, including XSB's, are based on a mechanism that allows a user to trace the evaluation of a goal by interrupting the evaluation at call, success, retry, or failure of various subgoals. While this has proved an excellent mechanism for evaluating SLD(NF) executions, it is difficult at best to use such a mechanism during a tabled evaluation. This is because, unlike with SLD(NF), SLG requires answers to be returned to tabled subgoals at various times (depending on whether batched or local evaluation is used), negative subgoals to be sometimes be delayed and/or simplified, etc.

One approach to understanding tabled evaluation better is to abstract away the procedural aspects of debugging and to use the tables produced by an evaluation to construct a *justification* after the evaluation has finished. The justification library does just this using algorithms described in [?].

1.16.2 AVL Trees

By Mats Carlsson

AVL trees provide a mechanism to maintain key value pairs so that loop up, insertion, and deletion all have complexity $\mathcal{O}(\log n)$. This library contains predicates to transform a sorted list to an AVL tree and back, along with predicates to manipulate the AVL trees.

1.16.3 Ordered Sets: `ordsets.P`

By Richard O'Keefe

(Summary from code documentation) `ordset.P` provides an XSB port of the widely used ordset library, whose summary we paraphrase here. In the ordset library, sets are represented by ordered lists with no duplicates. Thus $\{c, r, a, f, t\}$ is represented as `[a, c, f, r, t]`. The ordering is defined by the `@<` family of term comparison predicates, which is the ordering used by `sort/2` and `setof/3`. The benefit of the ordered representation is that the elementary set operations can be done in time proportional to the sum of the argument sizes rather than their product. Some of the unordered set routines, such as `member/2`, `length/2`, or `select/3` can be used unchanged.

1.16.4 Unweighted Graphs: `ugraphs.P`

By Mats Carlsson


XSB also includes a library for unweighted graphs. This library allows for the representation and manipulation of directed and non-directed unlabelled graphs, including predicates to find the transitive closure of a graph, maximal paths, minimal paths, and other features. This library represents graphs as an ordered set of their edges and does not use tabling. As a result, it may be slower for large graphs than similar predicates based on a datalog representatoin of edges.

1.16.5 Heaps: heaps.P

By Richard O’Keefe

(Summary from code documentation). A heap is a labelled binary tree where the key of each node is less than or equal to the keys of its sons. The point of a heap is that we can keep on adding new elements to the heap and we can keep on taking out the minimum element. If there are N elements total, the total time is $\mathcal{O}(N \lg(N))$. If you know all the elements in advance, you are better off doing a merge-sort, but this file is for when you want to do say a best-first search, and have no idea when you start how many elements there will be, let alone what they are.

A heap is represented as a triple $\mathbf{t}(\mathbf{N}, \mathbf{Free}, \mathbf{Tree})$ where \mathbf{N} is the number of elements in the tree, \mathbf{Free} is a list of integers which specifies unused positions in the tree, and \mathbf{Tree} is a tree made of \mathbf{t} terms for empty subtrees and $\mathbf{t}(\mathbf{Key}, \mathbf{Datum}, \mathbf{Lson}, \mathbf{Rson})$ terms for the rest. The nodes of the tree are notionally numbered like this:



The idea is that if the maximum number of elements that have been in the heap so far is M , and the tree currently has K elements, the tree is some subtree of the tree of this form having exactly M elements, and the \mathbf{Free} list is a list of $K - M$ integers saying which of the positions in the M -element tree are currently unoccupied. This free list is needed to ensure that the cost of passing N elements through the heap is $\mathcal{O}(N \lg(M))$ instead of $\mathcal{O}(N \lg N)$. For M say 100 and N say 10^4 this means a factor of two.

Chapter 2

XSB-ODBC Interface

By Baoqiu Cui, Lily Dong, and David S. Warren ¹.

2.1 Introduction

The XSB-ODBC interface is subsystem that allows XSB users to access databases through ODBC connections. This is mostly of interest to Microsoft Windows users. The interface allows XSB users to access data in any ODBC compliant database management system (DBMS). Using this uniform interface, information in different DBMS's can be accessed as though it existed as Prolog facts. The XSB-ODBC interface provides users with three levels of interaction: an *SQL level*, a *relation level* and a *view level*. The *SQL level* allows users to write explicit SQL statements to be passed to the interface to retrieve data from a connected database. The *relation level* allows users to declare XSB predicates that connect to individual tables in a connected database, and which when executed support tuple-at-a-time retrieval from the base table. The *view level* allows users to use a complex XSB query, including conjunction, negation and aggregates, to specify a database query. A listing of the features that the XSB-ODBC interface provides is as follows:

- Concurrent access from multiple XSB processes to a single DBMS
- Access from a single XSB process to multiple ODBC DBMS's
- Full data access and cursor transparency including support for
 - Full data recursion through XSB's tabling mechanism (depending on the capabilities of the underlying ODBC driver.
 - Runtime type checking
 - Automatic handling of NULL values for insertion, deletion and querying
- Full access to data source including

¹This interface was partly based on the XSB-Oracle Interface by Hassan Davulcu, Ernie Johnson and Terrance Swift.

- Transaction support
 - Cursor reuse for cached SQL statements with bind variables (thereby avoiding re-parsing and re-optimizing).
 - Caching compiler generated SQL statements with bind variables and efficient cursor management for cached statements
- A powerful Prolog / SQL compiler based on [?].
 - Full source code availability
 - Independence from database schema by the *relation level* interface
 - Performance as SQL by employing a *view level*
 - No mode specification is required for optimized view compilation

We use the `Hospital` database as our example to illustrate the usage of XSB-ODBC interface in this manual. We assume the basic knowledge of Microsoft ODBC interface and its ODBC administrator throughout the text. Please refer to “Inside WindowsTM 95” (or more recent documentation) for information on this topic.

2.2 Using the Interface

The XSB-ODBC module is a module and as such exports the predicates it supports. In order to use any predicate defined below, **it must be imported** from `odbc_call`. For example, before you can use the predicate to open a data source, you must include:

```
:- import odbc_open/3 from odbc_call.
```

2.2.1 Connecting to and Disconnecting from Data Sources

Assuming that the data source to be connected to is available, i.e. it has an entry in `ODBC.INI` file which can be checked by running Microsoft ODBC Administrator, it can be connected to in the following way:

```
| ?- odbc_open(data_source_name, username, passwd).
```

If the connection is successfully made, the predicate invocation will succeed. This step is necessary before anything can be done with the data sources since it gives XSB the opportunity to initialize system resources for the session.

This is an executable predicate, but you may want to put it as a query in a file that declares a database interface and will be loaded.

To close the current session use:

```
| ?- odbc_close.
```

and XSB will give all the resources it allocated for this session back to the system.

If you are connecting to only one data source at a time, the predicates above are sufficient. However, if you want to connect to multiple data sources at the same time, you must use extended versions of the predicates above. When connecting to multiple sources, you must give an atomic name to each source you want to connect to, and use that name whenever referring to that source. The names may be chosen arbitrarily but must be used consistently. The extended versions are:

```
| ?- odbc_open(data_source_name, username, passwd, connectionName).
```

and

```
| ?- odbc_close(connectionName).
```

A list of existing Data Source Names and descriptions can be obtained by backtracking through `odbc_data_sources/2`. For example:

```
| ?- odbc_data_sources(DSN,DSNDescr).
```

```
DSN = mycdf
```

```
DSNDescr = MySQL driver;
```

```
DSN = mywincdf
```

```
DSNDescr = TDS driver (Sybase/MS SQL);
```

2.2.2 Accessing Tables in Data Sources Using SQL

There are several ways that can be used to extract information from or modify a table in a data source. The most basic way is to use predicates that pass an SQL statement directly to the ODBC driver. The basic call is:

```
| ?- odbc_sql(BindVals,SQLStmt,ResultRow).
```

where `BindVals` is a list of (ground) values that correspond to the parameter indicators in the SQL statement (the '?'s); `SQLStmt` is an atom containing an SQL statement; and `ResultRow` is a returned list of values constituting a row from the result set returned by the SQL query. Thus for a select SQL statement, this call is nondeterministic, returning each retrieved row in turn.

The `BindVals` list should have a length corresponding to the number of parameters in the query, in particular being the empty list (`[]`) if `SQLStmt` contains no '?'s. If `SQLStmt` is not a select statement returning a result set, then `ResultRow` will be the empty list, and the call is deterministic. Thus this predicate can be used to do updates, DDL statements, indeed any SQL statement.

`SQLStmt` need not be an atom, but can be a (nested) list of atoms which flattens (and concatenates) to form an SQL statement.

`BindVals` is normally a list of values of primitive Prolog types: atoms, integers, or floats. The values are converted to the types of the corresponding database fields. However, complex Prolog values can also be stored in a database field. If a term of the form `term(VAL)` appears in the `BindVal` list, then `VAL` (a Prolog term) will be written in canonical form (as produced by `write_canonical`) to the corresponding database field (which must be `CHAR` or `BYTE`). If a term of the form `string(CODELIST)` appears in `BindVal`, then `CODELIST` must be a list of ascii-codes (as produced by `atom_codes`) and these codes will be converted to a `CHAR` or `BYTE` database type.

`ResultRow` for a select statement is normally a list of variables that will nondeterministically be bound to the values of the fields of the tuples returned by the execution of the select statement. The Prolog types of the values returned will be determined by the database types of the corresponding fields. A `CHAR` or `BYTE` database type will be returned as a Prolog atom; an `INTEGER` database field will be returned as a Prolog integer, and similarly for floats. However, the user can request that `CHAR` and `BYTE` database fields be returned as something other than an atom. If the term `string(VAR)` appears in `ResultRow`, then the corresponding database field must be `CHAR` or `BYTE`, and in this case, the variable `VAR` will be bound to the list of ascii-codes that make up the database field. This allows an XSB programmer to avoid adding an atom to the atom table unnecessarily. If the term `term(VAR)` appears in `ResultRow`, then the corresponding database field value is assumed to be a Prolog term in canonical form, i.e., can be read by `read_canonical/1`. The corresponding value will be converted into a Prolog term and bound to `VAR`. This allows a programmer to store complex Prolog terms in a database. Variables in such a term are local only to that term.

When connecting to multiple data sources, you should use the form:

```
| ?- odbc_sql(ConnectionName,BindVals,SQLStmt,ResultRow).
```

For example, we can define a predicate, `get_test_name_price`, which given a test ID, retrieves the name and price of that test from the test table in the hospital database:

```
get_test_name_price(Id,Nam,Pri) :-
    odbc_sql([Id], 'SELECT TName,Price FROM Test WHERE TId = ?', [Nam,Pri]).
```

The interface uses a cursor to retrieve this result and caches the cursor, so that if the same query is needed in the future, it does not need to be re-parsed, and re-optimized. Thus, if this predicate were to be called several times, the above form is more efficient than the following form, which must be parsed and optimized for each and every call:

```
get_test_name_price(Id,Nam,Pri) :-
    odbc_sql([], ['SELECT TName,Price FROM Test WHERE TId = ', Id, '''], [Nam,Pri]).
```

Note that to include a quote (') in an atom, it must be represented by using two quotes.

There is also a predicate:


```
| ?- odbc_sql_cnt(ConnectionName,BindVals,SQLStmt,Count).
```

This predicate is very similar to `odbc_sql/4` except that it can only be used for UPDATE, INSERT, and DELETE SQL statements. The first three arguments are just as in `odbc_sql/4`; the fourth must be a variable in which is returned the integer count of the number of rows affected by the SQL operation.

2.2.3 Cursor Management

The XSB-ODBC interface is limited to using 100 open cursors. When XSB systems use database accesses in a complicated manner, management of open cursors can be a problem due to the tuple-at-a-time access of databases from Prolog, and due to leakage of cursors through cuts and throws. Often, it is more efficient to call the database through set-at-a-time predicates such as `findall/3`, and then to backtrack through the returned information. For instance, the predicate `findall_odbc_sql/4` can be defined as:

```
findall_odbc_sql(ConnName,BindVals,SQLStmt,ResultRow):-
    findall(Res,odbc_sql(ConnName,BindVals,SQLStmt,Res),Results),
    member(ResultRow,Results).
```

As a convenience, therefore, the predicates `findall_odbc_sql/3` and `findall_odbc_sql/4` are defined in the ODBC interface.

2.2.4 Accessing Tables in Data Sources through the Relation Level

While all access to a database is possible using SQL as described above, the XSB-ODBC interface supports higher-level interaction for which the user need not know or write SQL statements; that is done as necessary by the interface. With the relation level interface, users can simply declare a predicate to access a table and the system generates the necessary underlying code, generating specialized code for each mode in which the predicate is called.

To declare a predicate to access a database table, a user must use the `odbc_import/2` interface predicate.

The syntax of `odbc_import/2` is as follows:

```
| ?- odbc_import('TableName'('FIELD1', 'FIELD2', ..., 'FIELDn'), 'PredicateName').
```

where `'TableName'` is the name of the database table to be accessed and `'PredicateName'` is the name of the XSB predicate through which access will be made. `'FIELD1', 'FIELD2', ... , 'FIELDn'` are the exact attribute names(case sensitive) as defined in the database table schema. The chosen columns define the view and the order of arguments for the database predicate `'PredicateName'`.

For example, to create a link to the `Test` table through the `'test'` predicate:

```
| ?- odbc_import('Test'('TId','TName','Length','Price'),test).
```

yes

When connecting to multiple data sources, you should use the form:

```
| ?- odbc_import(ConnectionName,
                'TableName'('FIELD1', 'FIELD2', ..., 'FIELDn'),
                'PredicateName').
```

2.2.5 Using the Relation Level Interface

Once the links between tables and predicates have been successfully established, information can then be extracted from these tables using the corresponding predicates. Continuing from the above example, now rows from the table `Test` can be obtained:

```
| ?- test(TId, TName, L, P).
```

```
TId = t001
TName = X-Ray
L = 5
P = 100
```

Backtracking can then be used to retrieve the next row of the table `Test`.

Records with particular field values may be selected in the same way as in Prolog; no mode specification for database predicates is required. For example:

```
| ?- test(TId, 'X-Ray', L, P).
```

will automatically generate the query:

```
SELECT rel1.TId, rel1.TName, rel1.Length, rel1.Price
FROM Test rel1
WHERE rel1.TName = ?
```

and

```
| ?- test('NULL'(_), 'X-Ray', L, P).
```

generates: (See Section [2.2.6](#))

```
SELECT NULL , rel1.TName, rel1.Length, rel1.Price
FROM Test rel1
WHERE rel1.TId IS NULL AND rel1.TName = ?
```

During the execution of this query the bind variable ? will be bound to the value 'X-Ray'.

Of course, the same considerations about cursors noted in Section 2.2.3 apply to the relation-level interface. Accordingly, the ODBC interface also defines the predicate `odbc_import/4` which allows the user to specify that rows are to be fetched through `findall/3`. For example, the call

```
odbc_import('Test'('TId','TName','Length','Price'),test,[findall(true)]).
```

will behave as described above *but* will make all database calls through `findall/3` and return rows by backtracking through a list rather than maintaining open cursors.

Also as a courtesy to Quintus Prolog users we have provided compatibility support for some PRODBI predicates which access tables at a relational level ².

```
| ?- odbc_attach(PredicateName, table(TableName)).
```

eg. invoke

```
| ?- odbc_attach(test2, table('Test')).
```

and then execute

```
| ?- test2(TId, TName, L, P).
```

to retrieve the rows.

2.2.6 Handling NULL values

The interface treats NULL's by introducing a single valued function 'NULL'/1 whose single value is a unique (Skolem) constant. For example a NULL value may be represented by

```
'NULL'(null123245)
```

Under this representation, two distinct NULL values will not unify. On the other hand, the search condition `IS NULL Field` can be represented in XSB as `Field = 'NULL'(_)`

Using this representation of NULL's the following protocol for queries and updates is established.

Queries

```
| ?- dept('NULL'(_),_,_).
```

Generates the query:

²This predicate is obsolescent and `odbc_import/{2,3,4}` should be used instead.

```
SELECT NULL , rel1.DNAME , rel1.LOC
FROM DEPT rel1
WHERE rel1.DEPTNO IS NULL;
```

Hence, 'NULL'(_) can be used to retrieve rows with NULL values at any field.

'NULL'/1 fails the predicate whenever it is used with a bound argument.

```
| ?- dept('NULL'(null12745),_,_). → fails always.
```

Query Results

When returning NULL's as field values, the interface returns NULL/1 function with a unique integer argument serving as a skolem constant.

Notice that the above guarantees the expected semantics for the join statements. In the following example, even if Deptno is NULL for some rows in emp or dept tables, the query still evaluates the join successfully.

```
| ?- emp(ENAME,_,_,_,Deptno),dept(Deptno,Dname,Loc) ..
```

Inserts

To insert rows with NULL values you can use Field = 'NULL'(_) or Field = 'NULL'(null12346). For example:

```
| ?- emp_ins('NULL'(_), ...). → inserts a NULL value for ENAME
```

```
| ?- emp_ins('NULL'('bound'), ...) → inserts a NULL value for ENAME.
```

Deletes

To delete rows with NULL values at any particular FIELD use Field = 'NULL'(_), 'NULL'/1 with a free argument. When 'NULL'/1's argument is bound it fails the delete predicate always. For example:

```
| ?- emp_del('NULL'(_), ...). → adds ENAME IS NULL to the generated SQL
statement
```

```
| ?- emp_del('NULL'('bound'), ...). → fails always
```

The reason for the above protocol is to preserve the semantics of deletes, when some free arguments of a delete predicate get bound by some preceding predicates. For example in the following clause, the semantics is preserved even if the Deptno field is NULL for some rows.

```
| ?- emp(_,_,_,_,Deptno), dept_del(Deptno).
```

2.2.7 The View Level Interface

The view level interface can be used to define XSB queries which include only imported database predicates (by using the relation level interface) described above and aggregate predicates (defined below). When these queries are invoked, they are translated into complex database queries, which are then executed taking advantage of the query processing ability of the DBMS.

One can use the view level interface through the predicate `odbc_query/2`:

```
| ?- odbc_query('QueryName'(ARG1, ..., ARGn), DatabaseGoal).
```

All arguments are standard XSB terms. `ARG1`, `ARG2`, ..., `ARGn` define the attributes to be retrieved from the database, while `DatabaseGoal` is an XSB goal (i.e. a possible body of a rule) that defines the selection restrictions and join conditions.

The compiler is a simple extension of [?] which generates SQL queries with bind variables and handles NULL values as described in Section 2.2.6. It allows negation, the expression of arithmetic functions, and higher-order constructs such as grouping, sorting, and aggregate functions.

Database goals are translated according to the following rules from [?]:

- Disjunctive goals translate to distinct SQL queries connected through the UNION operator.
- Goal conjunctions translate to joins.
- Negated goals translate to negated EXISTS subqueries.
- Variables with single occurrences in the body are not translated.
- Free variables translate to grouping attributes.
- Shared variables in goals translate to equi-join conditions.
- Constants translate to equality comparisons of an attribute and the constant value.
- Nulls are translated to IS NULL conditions.

For more examples and implementation details see [?].

In the following, we show the definition of a simple join view between the two database predicates *Room* and *Floor*.

Assuming the declarations:

```
| ?- odbc_import('Room'('RoomNo','CostPerDay','Capacity','FId'),room).

| ?- odbc_import('Floor'('FId','','FName'),floor).
```

use

```
| ?- odbc_query(query1(RoomNo,FName),
                (room(RoomNo,_,_,FId),floor(FId,_,FName))).
yes

| ?- query1(RoomNo,FloorName).
```

Prolog/SQL compiler generates the SQL statement:

```
SELECT rel1.RoomNo , rel2.FName FROM Room rel1 , Floor rel2
WHERE rel2.FId = rel1.FId;
```

Backtracking can then be used to retrieve the next row of the view.

```
| ?- query1('101','NULL'(_)).
```

generates the SQL statement:

```
SELECT rel1.RoomNo, NULL
FROM Room rel1 , Floor rel2
WHERE rel1.RoomId = ? AND rel2.FId = rel1.FId AND rel2.FName IS NULL;
```

The view interface also supports aggregate functions such as sum, avg, count, min and max. For example

```
| ?- odbc_import('Doctor'('DId', 'FId', 'DName','PhoneNo','ChargePerMin'),doctor).

yes
| ?- odbc_query(avgchargepermin(X),
                (X is avg(ChargePerMin, A1 ^ A2 ^ A3 ^ A4 ^
                        doctor(A1,A2, A3,A4,ChargePerMin)))).

yes
| ?- avgchargepermin(X).

SELECT AVG(rel1.ChargePerMin)
FROM doctor rel1;
```

```
X = 1.64
```

```
yes
```

A more complicated example is the following:

```

| ?- odbc_query(nonsense(A,B,C,D,E),
               (doctor(A, B, C, D, E),
                not floor('First Floor', B),
                not (A = 'd001'),
                E > avg(ChargePerMin, A1 ^ A2 ^ A3 ^ A4 ^
                      (doctor(A1, A2, A3, A4, ChargePerMin)))))).

| ?- nonsense(A,'4',C,D,E).

SELECT rel1.DId , rel1.FId , rel1.DName , rel1.PhoneNo , rel1.ChargePerMin
FROM doctor rel1
WHERE rel1.FId = ? AND NOT EXISTS
(SELECT *
FROM Floor rel2
WHERE rel2.FName = 'First Floor' and rel2.FId = rel1.FId
) AND rel1.DId <> 'd001' AND rel1.ChargePerMin >
(SELECT AVG(rel3.ChargePerMin)
FROM Doctor rel3
);

A = d004
C = Tom Wilson
D = 516-252-100
E = 2.5

```

All database queries defined by `odbc_query/{2,3}` can be queried with any mode.

Note that at each call to a database relation or rule, the communication takes place through bind variables. The corresponding restrictive SQL query is generated, and if this is the first call with that adornment, it is cached. A second call with same adornment would try to use the same database cursor if still available, without reparsing the respective SQL statement. Otherwise, it would find an unused cursor and retrieve the results. In this way efficient access methods for relations and database rules can be maintained throughout the session.

If connecting to multiple data sources, use the form:

```
:- odbc_query(connectionName,'QueryName'(ARG1, ..., ARGn), DatabaseGoal).
```

2.2.8 Insertions and Deletions of Rows through the Relational Level

Insertion and deletion operations can also be performed on an imported table. The two predicates to accomplish these operations are `odbc_insert/2` and `odbc_delete/2`. The syntax of `odbc_insert/2` is as follows: the first argument is the declared database predicate for insertions

and the second argument is some imported data source relation. The second argument can be declared with some of its arguments bound to constants. For example after `Room` is imported through `odbc_import`:

```
|?- odbc_import('Room'('RoomNo','CostPerDay','Capacity','FId'), room).
yes
```

Now we can do

```
| ?- odbc_insert(room_ins(A1,A2,A3),(room(A1,A2,A3,'3'))).

yes
| ?- room_ins('306','NULL'(_),2).

yes
```

This will insert the row: ('306',NULL, 2,'3') into the table `Room`. Note that any call to `room_ins/7` should have all its arguments bound.

See Section 2.2.6) for information about NULL value handling.

The first argument of `odbc_delete/2` predicate is the declared delete predicate and the second argument is the imported data source relation with the condition for requested deletes, if any. The condition is limited to simple comparisons. For example assuming `Room/3` has been imported as above:

```
| ?- odbc_delete(room_del(A), (room('306',A,B,C), A > 2)).

yes
```

After this declaration you can use:

```
| ?- room_del(3).
```

to generate the SQL statement:

```
DELETE From Room rel1
WHERE rel1.RoomNo = '306' AND rel1.CostPerDay = ? AND ? > 2
;
```

Note that you have to commit your inserts or deletes to tables to make them permanent. (See section 2.2.11).

These predicates also have the form in which an additional first argument indicates a connection, for use with multiple data sources.

Also, some ODBC drivers have been found that do not accept the form of SQL generated for deletes. In these cases, you must use the lower-level interface: `odbc_sql`.

2.2.9 Access to Data Dictionaries

The following utility predicates provide users with tools to access data dictionaries ³. A brief description of these predicates is as follows:

odbc_show_schema(accessible(Owner)) Shows the names of all accessible tables that are owned by Owner. (This list can be long!) If Owner is a variable, all tables will be shown, grouped by owner.

odbc_show_schema(user) Shows just those tables that belongs to user.

odbc_show_schema(tuples('Table')) Shows all rows of the database table named 'Table'.

odbc_show_schema(arity('Table')) The number of fields in the table 'Table'.

odbc_show_schema(columns('Table')) The field names of a table.

For retrieving above information use:

- `odbc_get_schema(accessible(Owner),List)`
- `odbc_get_schema(user,List)`
- `odbc_get_schema(arity('Table'),List)`
- `odbc_get_schema(columns('Table'),List)`

The results of above are returned in List as a list.

2.2.10 Other Database Operations

odbc_create_table('TableName','FIELDS') FIELDS is the field specification as in SQL.

```
eg. odbc_create_table('MyTable', 'Col1 NUMBER,
                               Col2 TEXT(50),
                               Col3 TEXT(13)').
```

odbc_create_index('TableName','IndexName', index(__,Fields)) Fields is the list of columns for which an index is requested. For example:

³Users of Quintus Prolog may note that these predicates are all PRODBI compatible.

```
odbc_create_index('Doctor', 'DocKey', index(_, 'DId')).
```

odbc_delete_table('TableName') To delete a table named 'TableName'

odbc_delete_view('ViewName') To delete a view named 'ViewName'

odbc_delete_index('IndexName') To delete an index named 'IndexName'

2.2.11 Transaction Management

Depending on how the transaction options are set in ODBC.INI for data sources, changes to the data source tables may not be committed (i.e., the changes become permanent) until the user explicitly issues a commit statement. Some ODBC drivers support autocommit, which, if on, means that every update operation is immediately committed upon execution. If autocommit is off, then an explicit commit (or rollback) must be done by the program to ensure the updates become permanent (or are ignored.).

The predicate `odbc_transaction/1` supports these operations.

odbc_transaction(autocommit(on)) Turns on autocommit, so that all update operations will be immediately committed on completion.

odbc_transaction(autocommit(off)) Turns off autocommit, so that all update operations will not be committed until explicitly done so by the program (using one of the following operations.)

odbc_transaction(commit) Commits all transactions up to this point. (Only has an effect if autocommit is off).

odbc_transaction(rollback) Rolls back all update operations done since the last commit point. (Only has an effect if autocommit is off).

2.2.12 Interface Flags

Users are given the option to monitor control aspects of the ODBC interface by setting ODBC flags via the predicates `set_odbc_flag/2` and `odbc_flag/2`.

The first aspect that can be controlled is whether to display SQL statements for SQL queries. This is done by the `show_query` flag. For example:

```
| ?- odbc_flag(show_query, Val).
```

```
Val = on
```

Indicates that SQL statements will now be displayed for all SQL queries, and is the default value for the ODBC interface. To turn it off execute the command `set_odbc_flag(show_query,on)`.

The second aspect that can be controlled is the action taken upon ODBC errors. Three possible actions may be useful in different contexts and with different drivers. First, the error may be ignored, so that a database call succeeds; second the error cause the predicate to fail, and third the error may cause an exception to be thrown to be handled by a catcher (or the default system error handler, see Volume 1).

- | ?- `odbc_flag(fail_on_error, ignore)` Ignores all ODBC errors, apart from writing a warning. In this case, it's the users' responsibility to check each of their actions and do error handling.
- | ?- `odbc_flag(fail_on_error, fail)` Interface fails whenever error occurs.
- | ?- `odbc_flag(fail_on_error, throw)` Throws an error-term of the form `error(odbc_error,Message,Backtrace)` in which `Message` is a textual description of the ODBC error, and `Backtrace` is a list of the continuations of the call. These continuations may be printed out by the error handler.

The default value of `fail_on_error` is `on`.

2.2.13 Datalog

Users can write recursive Datalog queries with exactly the same semantics as in XSB using imported database predicates or database rules. For example assuming `odbc_parent/2` is an imported database predicate, the following recursive query computes its transitive closure.

```
:- table(ancestor/2).
ancestor(X,Y) :- odbc_parent(X,Y).
ancestor(X,Z) :- ancestor(X,Y), odbc_parent(Y,Z).
```

This works with drivers that support multiple open cursors to the same connection at the same time. (Sadly, some don't.) In the case of drivers that don't support multiple open cursors, one can often replace each `odbc_import-ed` predicate call

```
...,predForTable(A,B,C),...
```

by

```
...,findall([A,B,C],predForTable(A,B,C),PredList),
    member([A,B,C],PredList)...
```

and get the desired effect.

2.3 Error messages

ERR - DB: Connection failed For some reason the attempt to connect to data source failed.

- Diagnosis: Try to see if the data source has been registered with Microsoft ODBC Administrator, the username and password are correct and MAXCURSORNUM is not set to a very large number.

ERR - DB: Parse error The SQL statement generated by the Interface or the first argument to `odbc_sql/1` or `odbc_sql_select/2` can not be parsed by the data source driver.

- Diagnosis: Check the SQL statement. If our interface generated the erroneous statement please contact us at `xsb-contact@cs.sunysb.edu`.

ERR - DB: No more cursors left Interface run out of non-active cursors either because of a leak or no more free cursors left.

- Diagnosis: System fails always with this error. `odbc_transaction(rollback)` or `odbc_transaction(commit)` should resolve this by freeing all cursors.

ERR - DB: FETCH failed Normally this error should not occur if the interface running properly.

- Diagnosis: Please contact us at `xsb-contact@cs.sunysb.edu`

2.4 Notes on specific ODBC drivers

MyODBC The ODBC driver for MySQL is called MyODBC, and it presents some particularities that should be noted.

First, MySQL, as of version 3.23.55, does not support strings of length greater than 255 characters. XSB's ODBC interface has been updated to allow the use of the BLOB datatype to encode larger strings.

More importantly, MyODBC implements `SQLDescribeCol` such that, by default, it returns actual lengths of columns in the result table, instead of the formal lengths in the tables. For example, suppose you have, in table A, a field f declared as "VARCHAR (200)". Now, you create a query of the form "SELECT f FROM A WHERE ...". If, in the result set, the largest size of f is 52, that's the length that `SQLDescribeCol` will return. This breaks XSB's caching of query-related data-structures. In order to prevent this behavior, you should configure your DSN setup so that you pass "Option=1" to MyODBC.

Chapter 3

The New XSB-Database Interface

By Saikat Mukherjee, Michael Kifer and Hui Wan

3.1 Introduction

The XSB-DB interface is a package that allows XSB users to access databases through various drivers. Using this interface, information in different DBMSs can be accessed by SQL queries. The interface defines Prolog predicates which makes it easy to connect to databases, query them, and disconnect from the databases. Central to the concept of a connection to a database is the notion of a *connection handle*. A connection handle describes a particular connection to a database. Similar to a connection handle is the notion of a query handle which describes a particular query statement. As a consequence of the handles, it is possible to open multiple database connections (to the same or different databases) and keep alive multiple queries (again from the same or different connections). The interface also supports dynamic loading of drivers. As a result, it is possible to query databases using different drivers concurrently ¹.

Currently, this package provides drivers for ODBC, a native MySQL driver, and a driver for the embedded MySQL server.

3.2 Configuring the Interface

Generally, each driver has to be configured separately, but if the database packages such as ODBC, MySQL, etc., are installed in standard places then the XSB configuration mechanism will do the job automatically.

Under Windows, first make sure that XSB is configured and built correctly for Windows, and that it runs. As part of that building process, the command

```
makexsb_wind
```

¹In Version 3.4, this package has not been ported to the multi-threaded engine.

must have been executed in the directory `XSB\build`. It will normally configure the ODBC driver without problems. For the MySQL driver one has to edit the file

```
packages\dbdrivers\mysql\cc\NMakefile.mak
```

to indicate where MySQL is installed. To build the embedded MySQL driver under Windows, the file

```
packages\dbdrivers\mysqlenbedded\cc\NMakefile.mak
```

might need to be edited. Then you should either rebuild XSB using the `makexsb_wind` command or by running

```
nmake /f NMakefile.mak
```

in the appropriate directories (`dbdrivers\mysql\cc` or `dbdrivers\mysqlenbedded\cc`). Note that you need a C++ compiler and `nmake` installed on your system for this to work.²

Under Unix, the `configure` script will build the drivers automatically if the `-with-dbdrivers` option is specified. If, however, ODBC and MySQL are not installed in their standard places, you will have to provide the following parameters to the `configure` script:

- `-with-odbc-libdir=LibDIR` – `LibDIR` is the directory where the library `libodbc.so` lives on your system.
- `-with-odbc-incdir=IncludeDIR` – `IncludeDIR` is the directory where the ODBC header files, such as `sql.h` live.
- `-with-mysql-libdir=MySQLlibdir` – `MySQLlibdir` is the directory where MySQL's shared libraries live on your system.
- `-with-mysql-incdir=MySQLincludeDir` – `MySQLincludeDir` is the directory where MySQL's header files live.

If you are also using the embedded MySQL server and want to take advantage of the corresponding XSB driver, you need to provide the following directories to tell XSB where the copy of MySQL that supports the embedded server is installed. This has to be done *only* if that copy is not in a standard place, like `/usr/lib/mysql`.

- `-with-mysqlembded-libdir=MySQLlibdir` – `MySQLlibdir` is the directory where MySQL's shared libraries live on your system. This copy of MySQL must be configured with support for the embedded server.
- `-with-mysqlembded-incdir=MySQLincludeDir` – `MySQLincludeDir` is the directory where MySQL's header files live.

² <http://www.microsoft.com/express/vc/>
<http://download.microsoft.com/download/vc15/Patch/1.52/W95/EN-US/Nmake15.exe>

Under Cygwin, the ODBC libraries come with the distribution; they are located in the directory `/cygdrive/c/cygwin/lib/w32api/` and are called `odbc32.a` and `odbc32.a`. (Check if your installation is complete and has these libraries!) Otherwise, the configuration of the interface under Cygwin is same as in unix (you do not need to provide any ODBC-specific parameters to the configure script under Cygwin).

If at the time of configuring XSB some database packages (*e.g.*, MySQL) are not installed on your system, you can install them later and configure the XSB interface to them then. For instance, to configure the ODBC interface separately, you can type

```
cd packages/dbdrivers/odbc
configure
```

Again, if ODBC is installed in a non-standard location, you might need to supply the options `-with-odbc-libdir` and `-with-odbc-incdir` to the configure script. Under Cygwin ODBC is always installed in a standard place, and `configure` needs no additional parameters.

Under Windows, separate configuration of the XSB-DB interfaces is also possible, but you need Visual Studio installed. For instance, to configure the MySQL interface, type

```
cd packages\dbdrivers\mysql\cc
nmake /f NMakefile.mak
```

As before, you might need to edit the `NMakefile.mak` script to tell the compiler where the required MySQL's libraries are. You also need the file `packages\dbdrivers\mysql\mysql_init.P` with the following content:

```
:- export mysql_info/2.
mysql_info(support, 'yes').
mysql_info(libdir, '').
mysql_info(ccflags, '').
mysql_info(ldflags, '').
```

Similarly, to configure the ODBC interface, do

```
cd packages\dbdrivers\odbc\cc
nmake /f NMakefile.mak
```

You will also need to create the file `packages\dbdrivers\odbc\odbc_init.P` with the following contents:

```
:- export odbc_info/2.
odbc_info(support, 'yes').
odbc_info(libdir, '').
odbc_info(ccflags, '').
odbc_info(ldflags, '').
```

3.3 Using the Interface

We use the `student` database as our example to illustrate the usage of the XSB-DB interface in this manual. The schema of the student database contains three columns viz. the student name, the student id, and the name of the advisor of the student.

The XSB-DB package has to be first loaded before using any of the predicates. This is done by the call:

```
| ?- [dbdrivers].
```

Next, the driver to be used for connecting to the database has to be loaded. Currently, the interface has support for a native MySQL driver (using the MySQL C API), and an ODBC driver. For example, to load the ODBC driver call:

```
| ?- load_driver(odbc).
```

Similarly, to load the mysql driver call:

```
| ?- load_driver(mysql).
```

or

```
| ?- load_driver(mysqlembedded).
```

3.3.1 Connecting to and Disconnecting from Databases

There are two predicates for connecting to databases, `db_connect/5` and `db_connect/6`. The `db_connect/5` predicate is for ODBC connections, while `db_connect/6` is for other (non-ODBC) database drivers.

```
| ?- db_connect(+Handle, +Driver, +DSN, +User, +Password).  
| ?- db_connect(+Handle, +Driver, +Server, +Database, +User, +Password).
```

The `db_connect/5` predicate assumes that an entry for a data source name (DSN) exists in the `odbc.ini` file. The `Handle` is the connection handle name used for the connection. The `Driver` is the driver being used for the connection. The `User` and `Password` are the user name and password being used for the connection. The user is responsible for giving the name to the handle. To connect to the data source `mydb` using the user name `xsb` and password `xsb` with the `odbc` driver, the call is as follows:

```
| ?- db_connect(ha, odbc, mydb, xsb, xsb).
```


where `ha` is the user-chosen handle name (a Prolog atom) for the connection.

The `db_connect/6` predicate is used for drivers other than ODBC. The arguments `Handle`, `Driver`, `User`, and `Password` are the same as for `db_connect/5`. The `Server` and `Database` arguments specify the server and database to connect to. For example, for a connection to a database called `test` located on the server `wolfe` with the user name `xsb`, the password `foo`, and using the `mysql` driver, the call is:

```
| ?- db_connect(ha, mysql, wolfe, test, xsb, foo).
```

where `ha` is the handle name the user chose for the connection.

If the connection is successfully made, the predicate invocation will succeed. This step is necessary before anything can be done with the data sources since it gives XSB the opportunity to initialize system resources for the session.

To close a database connection use:

```
| ?- db_disconnect(Handle).
```

where `handle` is the connection handle name. For example, to close the connection to above `mysql` database call:

```
| ?- db_disconnect(ha).
```

and XSB will give all the resources it allocated for this session back to the system.

3.3.2 Querying Databases

The interface supports two types of querying. In direct querying, the query statement is not prepared while in prepared querying the query statement is prepared before being executed. The results from both types of querying are retrieved tuple at a time. Direct querying is done by the predicate:

```
| ?- db_query(ConnectionHandle, QueryHandle, SQLQueryList, ReturnList).
```

`ConnectionHandle` is the name of the handle used for the database connection. `QueryHandle` is the name of the *query handle* for this particular query. For prepared queries, the query handle is used both in order to execute the query and to close it and free up space. For direct querying, the query handle is used only for closing query statements (see below). The `SQLQueryList` is a list of terms which is used to build the SQL query. The terms in this list are ground atoms. `ReturnList` is a list of variables each of which correspond to a return value in the query. It is upto the user to specify the correct number of return variables corresponding to the query. Also, as in the case of a connection handle, the user is responsible for giving the name to the query handle. For example, a query on the student database to select all the students for a given advisor is accomplished by the call:

```
| ?- X = adv,
    db_query(ha, qa, ['select T.name from student T where T.advisor = ', X], [P]),
    fail.
```

where `ha` and `qa` are respectively the connection handle and query handle name the user chose.

Observe that the query list is composed of the SQL string and a ground value for the advisor. The return list is made of one variable corresponding to the student name. The failure drive loop retrieves all the tuples.

Preparing a query is done by the call to the predicate:

```
| ?- db_prepare(ConnectionHandle, QueryHandle, SQLQueryList).
```

As before, `ConnectionHandle` and `QueryHandle` specify the handles for the connection and the query. The `SQLQueryList` is a list of terms which build up the query string. The placeholder ‘?’ is used for values which have to be bound during the execution of the statement. For example, to prepare a query for selecting the advisor name for a student name using our student database:

```
| ?- db_prepare(ha, qa, ['select T.advisor from student T where T.name = ?']).
```

A prepared statement is executed using the predicate:

```
| ?- db_prepare_execute(QueryHandle, BindList, ReturnList).
```

The `BindList` contains the ground values corresponding to the ‘?’ in the prepared statement. The `ReturnList` is a list of variables for each argument in a tuple of the result set.

For direct querying, the query handle is closed automatically when all the tuples in the result set have been retrieved. In order to explicitly close a query handle, and free all the resources associated with the handle, a call is made to the predicate:

```
| ?- db_statement_close(QueryHandle).
```

where `QueryHandle` is the query handle for the statement to be closed.

The interface is also able to transparently handle Prolog terms. Users can both save and retrieve terms without any special processing.

3.4 Error Handling

Each predicate in the XSB-DB interface throws an exception with the functor

```
xsb_error(database(Number), Message)
```

where `Number` is a string with the error number and `Message` is a string with a slightly detailed error message. It is upto the user to catch this exception and proceed with error handling. This is done by the throw-catch error handling mechanism in XSB. For example, in order to catch the error which will be thrown when the user attempts to close a database connection for a handle (`ha`) which does not exist:

```
| ?- catch(db_disconnect(ha),
        xsb_error(database(Number), Message), handler(Number, Message)).
```

It is the user's responsibility to define the handler predicate which can be as simple as printing out the error number and message or may involve more complicated processing.

A list of error numbers and messages that are thrown by the XSB-DB interface is given below:

- **XSB_DBI_001: XSB_DBI ERROR: Driver already registered**
This error is thrown when the user tries to load a driver, using the `load_driver` predicate, which has already been loaded previously.
- **XSB_DBI_002: XSB_DBI ERROR: Driver does not exist**
This error is thrown when the user tries to connect to a database, using `db_connect`, with a driver which has not been loaded.
- **XSB_DBI_003: XSB_DBI ERROR: Function does not exist in this driver**
This error is thrown when the user tries to use a function support for which does not exist in the corresponding driver. For example, this error is generated if the user tries to use `db_prepare` for a connection established with the `mysql` driver.
- **XSB_DBI_004: XSB_DBI ERROR: No such connection handle**
This error is thrown when the user tries to use a connection handle which has not been created.
- **XSB_DBI_005: XSB_DBI ERROR: No such query handle**
This error is thrown when the user tries to use a query handle which has not been created.
- **XSB_DBI_006: XSB_DBI ERROR: Connection handle already exists**
This error is thrown when the user tries to create a connection handle in `db_connect` using a name which already exists as a connection handle.
- **XSB_DBI_007: XSB_DBI ERROR: Query handle already exists**
This error is thrown when the user tries to create a query handle, in `db_query` or `db_prepare`, using a name which already exists as a query handle for a different query.
- **XSB_DBI_008: XSB_DBI ERROR: Not all parameters supplied**
This error is thrown when the user tries to execute a prepared statement, using `db_prepare_execute`, without supplying values for all the parameters in the statement.
- **XSB_DBI_009: XSB_DBI ERROR: Unbound variable in parameter list**
This error is thrown when the user tries to execute a prepared statement, using `db_prepare_execute`, without binding all the parameters of the statement.

- **XSB_DBI_010: XSB_DBI ERROR: Same query handle used for different queries**
This error is thrown when the user issues a prepare statement (`db_prepare`) using a query handle that has been in use by another prepared statement and which has not been closed. Query handles must be closed before reuse.
- **XSB_DBI_011: XSB_DBI ERROR: Number of requested columns exceeds the number of columns in the query**
This error is thrown when the user `db_query` specifies more items to be returned in the last argument than the number of items in the `SELECT` statement in the corresponding query.
- **XSB_DBI_012: XSB_DBI ERROR: Number of requested columns is less than the number of columns in the query**
This error is thrown when the user `db_query` specifies fewer items to be returned in the last argument than the number of items in the `SELECT` statement in the corresponding query.
- **XSB_DBI_013: XSB_DBI ERROR: Invalid return list in query**
Something else is wrong with the return list of the query.
- **XSB_DBI_014: XSB_DBI ERROR: Too many open connections**
There is a limit (200) on the number of open connections.
- **XSB_DBI_015: XSB_DBI ERROR: Too many registered drivers**
There is a limit (100) on the number of database drivers that can be registered at the same time.
- **XSB_DBI_016: XSB_DBI ERROR: Too many active queries**
There is a limit (2000) on the number of queries that can remain open at any given time.

3.5 Notes on specific drivers

ODBC Driver

The ODBC driver has been tested in Linux using the `unixodbc` driver manager. It currently supports the following functionality: (a) connecting to a database using a DSN, (b) direct querying of the database, (c) using prepared statements to query the database, (d) closing a statement handle, and (d) disconnecting from the database. The ODBC driver has also been tested under Windows and Cygwin.

MySQL Driver

The MySQL driver provides access to the native MySQL C API. Currently, it has support for the following functionality: (a) connecting to a database using `db_connect`, (b) direct querying of the database, (c) using prepared statements to query the database, (d) closing a statement handle, and (e) disconnecting from the database.

The MySQL driver has been tested under Linux and Windows.

Driver for the Embedded MySQL Server

This driver provides access to the Embedded MySQL Server Library `libmysqld`. Currently, it has support for the following functionality: (a) connecting to a database `db_connect`, (b) direct querying of the database, (c) using prepared statements to query the database, (d) closing a statement handle, and (e) disconnecting from the database.

The MySQL driver for Embedded MySQL Server has been tested under Linux.

In order to use this driver, you will need:

- MySQL with Embedded Server installed on your machine. If you don't have a precompiled binary distribution of MySQL, which was configured with `libmysqld` support (the embedded server library), you will need to build MySQL from sources and configure it with the `-with-embedded-server` option.
- append to `/etc/my.cnf` (or `/etc/mysql/my.cnf` – whichever is used on your machine) or `~/my.cnf`:

```
[mysqlembedded_driver_SERVER]
language = /usr/share/mysql/english
datadir = .....
```

You will probably need to replace `/usr/share/mysql/english` with a directory appropriate for your MySQL installation.

You might also need to set the `datadir` option to specify the directory where the databases managed by the embedded server are to be kept. This has to be done if there is a possibility of running the embedded MySQL server alongside the regular MySQL server. In that case, the `datadir` directory of the embedded server must be different from the `datadir` directory of the regular server (which is likely to be specified using the `datadir` option in `/etc/my.cnf` or `/etc/mysql/my.cnf`). This is because specifying the same directory might lead to a corruption of your databases. See <http://dev.mysql.com/doc/refman/5.1/en/multiple-servers.html> for further details on running multiple servers.

Please note that loading the embedded MySQL driver increases the memory footprint of XSB. This additional memory is released automatically when XSB exits. If you need to release the memory before exiting XSB, you can call `driverMySQLEmbedded_lib_end` after disconnecting from MySQL. Note that once `driverMySQLEmbedded_lib_end` is called, no further connections to MySQL are allowed from the currently running session of XSB (or else XSB will exit abnormally).

Chapter 4

Introduction to XSB Packages

An XSB package is a piece of software that extends XSB functionality but is not critical to programming in XSB. Around a dozen packages are distributed with XSB, ranging from simple meta-interpreters to complex software systems. Some packages provide interfaces from XSB to other software systems, such as Perl, SModels or Web interfaces (as in the `libwww` package). Others, such as the CHR and Flora packages, extend XSB to different programming paradigms.

Each package is distributed in the `$XSB_DIR/packages` subdirectory, and has two parts: an initialization file, and a subdirectory in which package source code files and executables are kept. For example, the `xsbdoc` package has files `xsbdoc.P`, `xsbdoc.xwam`, and a subdirectory, `xsbdoc`. If a user doesn't want to retain `xsbdoc` (or any other package) he or she may simply remove the initialization files and the associated subdirectory without affecting the core parts of the XSB system.

Several of the packages are documented in this manual in the various chapters that follow. However, many of the packages contain their own manuals. For these packages, we provide only a summary of their functionality in Chapter 14.

Chapter 5

Wildcard Matching

By Michael Kifer

XSB has an efficient interface to POSIX wildcard matching functions. To take advantage of this feature, you must build XSB using a C compiler that supports POSIX 2.0 (for wildcard matching). This includes GCC and probably most other compilers. This also works under Windows, provided you install Cygnus' CygWin and use GCC to compile ¹.

The `wildmatch` package provides the following functionality:

1. Telling whether a wildcard, like the ones used in Unix shells, match against a given string. Wildcards supported are of the kind available in `tcsh` or `bash`. Alternating characters (*e.g.*, “[`abc`]” or “[`^abc`]”) are supported.
2. Finding the list of all file names in a given directory that match a given wildcard. This facility generalizes `directory/2` (in module `directory`), and it is much more efficient.
3. String conversion to lower and upper case.

To use this package, you need to type:

```
| ?- [wildmatch].
```

If you are planning to use it in an XSB program, you need this directive:

```
:- import glob_directory/4, wildmatch/3, convert_string/3 from wildmatch.
```

The calling sequence for `glob_directory/4` is:

```
glob_directory(+Wildcard, +Directory, ?MarkDirs, -FileList)
```

¹This package has not yet been ported to the multi-threaded engine.

The parameter `Wildcard` can be either a Prolog atom or a Prolog string. `Directory` is also an atom or a string; it specifies the directory to be globbed. `MarkDirs` indicates whether directory names should be decorated with a trailing slash: if `MarkDirs` is bound, then directories will be so decorated. If `MarkDirs` is an unbound variable, then trailing slashes will not be added.

`FileList` gets the list of files in `Directory` that match `Wildcard`. If `Directory` is bound to an atom, then `FileList` gets bound to a list of atoms; if `Directory` is a Prolog string, then `FileList` will be bound to a list of strings as well.

This predicate succeeds if at least one match is found. If no matches are found or if `Directory` does not exist or cannot be read, then the predicate fails.

The calling sequence for `wildmatch/3` is as follows:

```
wildmatch(+Wildcard, +String, ?IgnoreCase)
```

`Wildcard` is the same as before. `String` represents the string to be matched against `Wildcard`. Like `Wildcard`, `String` can be an atom or a string. `IgnoreCase` indicates whether case of letters should be ignored during matching. Namely, if this argument is bound to a non-variable, then the case of letters is ignored. Otherwise, if `IgnoreCase` is a variable, then the case of letters is preserved.

This predicate succeeds when `Wildcard` matches `String` and fails otherwise.

The calling sequence for `convert_string/3` is as follows:

```
convert_string(+InputString, +OutputString, +ConversionFlag)
```

The input string must be an atom or a character list. The output string must be unbound. Its type will “atom” if so was the input and it will be a character list if so was the input string. The conversion flag must be the atom `tolower` or `toupper`.

This predicate always succeeds, unless there was an error, such as wrong type argument passed as a parameter.

Chapter 6

pcre: Pattern Matching and Substitution Using PCRE

By Mandar Pathak

6.1 Introduction

This package employs the PCRE library to enable XSB perform pattern matching and string substitution based on Perl regular expressions.

6.2 Pattern matching

The `pcre` package provides two ways of doing pattern matching: first-match mode and bulk-match mode. The syntax of the `pcre:match/4` predicate is:

```
?- pcre:match(+Pattern, +Subject, -MatchList, +Mode).
```

To find only the first match, the `Mode` parameter must be set to the atom `one`. To find all matches, the `Mode` parameter is set to the atom `bulk`. The result of the matching is returned as a list of the form:

```
match(Match,Prematch,Postmatch,[Subpattern1, Subpattern2,...])
```

The `Pattern` and the `Subject` arguments of `pcre:match` must be XSB atoms. If there is a match in the subject, then the result is returned as a list of the form shown above. *Match* refers to the substring which matched the entire pattern. *Prematch* contains part of the subject-string that precedes the matched substring. *Postmatch* contains part of the subject following the matched substring. The list of subpatterns (the 4-th argument of the `match` data structure) corresponds to the substrings which matched the parenthesized expressions in the given pattern. For example:

```
?- pcre:match('(\d{5}-\d{4})\ [A-Z]{2}',
'Hello12345-6789 NYwalk', X, 'one').
X = [match(12345-6789 NY,Hello,walk,[12345-6789])]
```

In this example, the match was found for substring ‘12345-6789 NY’. The prematch is ‘Hello’ and the postmatch is ‘walk’. The substring ‘12345-6789’ matched the parenthesized expression $(\d{5}-\d{4})$ and hence it is returned as part of the subpatterns list. Consider another example:

```
?- pcre:match('[a-z]+@[a-z]+\.(com|net|edu)',
'a@b.com@c.net@d.edu', X, 'bulk').
X = [match(a@b.com,,@c.net@d.edu,[com]),
     match(com@c.net,a@b.,@d.edu,[net]),
     match(om@c.net,a@b.c,@d.edu,[net]),
     match(m@c.net,a@b.co,@d.edu,[net]),
     match(net@d.edu,a@b.com@c.,,[edu]),
     match(et@d.edu,a@b.com@c.n.,,[edu]),
     match(t@d.edu,a@b.com@c.ne.,,[edu])]
```

This example uses the bulk match mode of the `pcre_match/4` to find all possible matches which resemble a very basic email address. In case there is no prematch or postmatch to a matched substring, an empty string is returned.

In general, there can be any number of parenthesized subpatterns in a given pattern and the subpattern match-list in the 4-th argument of the `match` data structure can have 0, 1, 2, or more elements.

6.3 String Substitution

The `pcre` package also provides a way to perform string substitution via the `pcre:substitute/4` predicate. It has the following syntax:

```
?- pcre:substitute(+Pattern, +Subject, +Substitution, -Result).
```

Pattern is the regular expression against which *Subject* is matched. Each match found is then replaced by the *Substitution*, and the result is returned in the variable *Result*. Here, *Pattern*, *Subject* and *Substitution* have to be XSB atoms whereas *Result* must be an unbound variable. The following example illustrates the use of this predicate:

```
?- pcre:substitute('is','This is a Mississippi issue', 'was', X).
X = Thwas was a Mwasswassippi wassue
```

Note that the predicate `pcre:substitute/4` always works in a bulk mode. If one needs to substitute only *one* occurrence of a pattern, this is easy to do using the `pcre:match/4` predicate. For instance, if one wants to replace the third occurrence of “is” in the above string, we could issue the query

```
?- pcre:match('is','This is a Mississippi issue',X,bulk).
```

take the third element in the returned list, which is

```
match(is,'This is a M','sissippi issue',[])
```

and then concatenate the 2-nd argument with “was” and with the 3-d argument of that `match` data structure.

More examples of the use of the `pcre` package can be found in

```
$XSBDIR/examples/pcretest.P
```

6.4 Installation and configuration

XSB’s `pcre` package requires that the PCRE library is installed. For Windows, the PCRE library files are included with the installation. For Linux and Mac, the `libpcre` and `libpcre-dev` packages must be installed using the distribution’s package manager.

6.4.1 Configuring for Linux, Mac, and other Unices

If a particular Linux distribution does not include these libraries they must be downloaded and built manually. Please visit

```
http://www.pcre.org/
```

to download the latest distribution and follow the instructions given with the package.

To configure `pcre` on Linux, Mac, or on some other Unix variant, switch to the `XSB/build` directory and type:

```
cd ../packages/pcre
./configure
./makexsb
```

6.4.2 Configuring for Windows

Configuring `pcre` on Windows requires creating the DLL for Windows. To create the DLL, open the Visual C++ command prompt, switch to the root XSB directory, and type:

```
cd packages\pcre\cc
nmake /f NMakefile.mak
```

This builds the DLL required by XSB's `pcre` package on Windows. To ensure that the build went ahead smoothly, open the directory

```
{XSB_DIR}\config\x86-pc-windows\bin
```

and verify that the file `pcre4pl.dll` exists there.

Once the package has been configured, it must be loaded before it can be used:

```
?- [pcre].
```

Chapter 7

curl: The XSB Internet Access Package

By Aneesh Ali

7.1 Introduction

The `curl` package is an interface to the `libcurl` library, which provides access to most of the standard Web protocols. The supported protocols include FTP, FTPS, HTTP, HTTPS, SCP, SFTP, TFTP, TELNET, DICT, LDAP, LDAPS, FILE, IMAP, SMTP, POP3 and RTSP. Libcurl supports SSL certificates, HTTP POST, HTTP PUT, FTP uploading, HTTP form based upload, proxies, cookies, user+password authentication (Basic, Digest, NTLM, Negotiate, Kerberos4), file transfer resume, http proxy tunneling etc.

The `curl` package accepts input in the form of URLs and Prolog atoms. To load the `curl` package, the user should type

```
?- [curl].
```

The `curl` package is integrated with file I/O of XSB in a transparent fashion and for many purposes Web pages can be treated just as yet another kind of a file. We first explain how Web pages can be accessed using the standard file I/O feature and then describe other predicates, which provide a lower-level interface.

7.2 Integration with File I/O

The `curl` package is integrated with XSB File I/O so that a web page can be opened as any other file. Once a Web page is opened, it can be read or written just like the a normal file.

7.2.1 Opening a Web Document

Web documents are opened by the usual predicates **see/1**, **open/3**, **open/4**.

see(*url*(+*Url*))

see(*url*(+*Url*, *Options*))

open(*url*(+*Url*), +*Mode*, -*Stream*)

open(*url*(+*Url*), +*Mode*, -*Stream*, +*Options*)

Url is an atom that specifies a URL. *Stream* is the file stream of the open file. *Mode* can be

read to create an input stream or **write**, to create an output stream. For reading, the contents of the Web page are cached in a temporary file. For writing, a temporary empty file is created. This file is posted to the corresponding URL at closing.

The *Options* parameter is a list that controls loading. Members of that list can be of the following form:

redirect(*Bool*)

Specifies the redirection option. The supported values are true and false. If true, any number of redirects is allowed. If false, redirections are ignored. The default is true.

secure(*CrtName*)

Specifies the secure connections (https) option. *CrtName* is the name of the file holding one or more certificates to verify the peer with.

auth(*UserName*, *Password*)

Sets the username and password basic authentication.

timeout(*Seconds*)

Sets the maximum time in seconds that is allowed for the transfer operation.

user__agent(*Agent*)

Sets the User-Agent: header in the http request sent to the remote server.

7.2.2 Closing a Web Document

Web documents opened by the predicates **see/1**, **open/3**, and **open/4** above must be closed by the predicates **close/2** or **close/3**. The data written to the stream is first posted to the URL. If that succeeds, the stream is closed. ??? And if it does not succeed????

close(+*Stream*, +*Source*)

close(+*Stream*, +*Source*, +*Options*)

Source can be of the form **url**(*url*). *Stream* is a file stream. *Options* is a list of options supported normally for close.

7.3 Low Level Predicates

This section describes additional predicates provided by the `curl` packages, which extend the functionality provided by the file I/O integration.

7.3.1 Loading web documents

Web documents are loaded by the predicate `load__page/5`, which has many options. The parameters of this predicate are described below.

load__page(*+Source*, *+Options*, *-Properties*, *-Content*, *-Warn*)

Source can be of the form `url(url)` or an atom *url* (check!!!). The document is returned in *Content*. *Warn* is bound to a (possibly empty) list of warnings generated during the process.

Properties is bound to a list of properties of the document. They include *Directory name*, *File name*, *File suffix*, *Page size*, and *Page time*. The `load__page/5` predicate caches a copy of the Web page that it fetched from the Web in a local file, which is specified by the above properties *Directory name*, *File name*, and *File suffix*. The remaining two parameters indicate the size and the last modification time of the fetched Web page. The directory and the file name The *Options* parameter is the same as in the URL opening predicates.

7.3.2 Retrieve the properties of a web document

The properties of a web document are loaded by the predicates `url__properties/3` and `url__properties/2`.

url__properties(*+Url*, *+Options*, *-Properties*)

The *Options* and *Properties* are same as in `load__page/5`.

url__properties(*+Url*, *-Properties*)

What are the default options???

7.3.3 Encode Url

Sometimes it is necessary to convert a URL string into something that can be used, for example, as a file name. This is done by the following predicate.

url__properties(*+Source*, *-Properties*)

(???? `url__properties` for encoding??? Explain the difference with `url__properties`)

Source has the form `url(url)` or an atom *url*, where *url* is an atom. (check!!!) *Properties* is bound to a list of properties of the URL. They include *Directory Name*, *File Name* and, *File Suffix* of the URL.

7.3.4 Obtaining the Redirection URL

If the originally specified URL was redirected, the URL of the page that was actually fetched by `load_page/5` can be found with the help of the following predicate:

```
get_redir_url(+Source, -UrlNew)  
    Source can be of the form url(url), file(filename) or a string.
```

7.4 Installation and configuration

The `curl` package of XSB requires that the `libcurl` package is installed. For Windows, the `libcurl` library files are included with the installation. For Linux and Mac, the `libcurl` and `libcurl-dev` packages need to be installed using the distribution's package manager. In some Linux distributions, `libcurl-dev` might be called `libcurl-gnutls-dev` or `libcurl-openssl-dev`. In addition, the release number might be attached. For instance, `libcurl4` and `libcurl4-openssl-dev`.

If a particular Linux distribution does not include the above packages and for other Unix variants, the `libcurl` package must be downloaded and built manually. See

<http://curl.haxx.se/download.html>

To configure `curl` on Linux, Mac, or on some other Unix variant, switch to the `XSB/build` directory and type

```
cd XSB/packages/curl  
./configure  
./makexsb
```


Chapter 8

sgml and xpath: SGML/XML/HTML Parsers and XPath

By Rohan Shirwaikar

8.1 Introduction

This suite of packages consists of the `sgml` package, which can parse XML, HTML, XHTML, and even SGML documents and the `xpath` package, which supports XPath queries on XML documents. The `sgml` package is an adaptation of a similar package in SWI Prolog and a port of SWI's codebase with some minor changes. The `xpath` package provides an interface to the popular `libxml2` library, which supports XPath and XML parsing, and is used in Mozilla based browsers. At present, the XML parsing capabilities of `libxml2` are not utilized explicitly in XSB, but such support might be provided in the future. The `sgml` package does not rely on `libxml2` ¹.

Installation and configuration. The `sgml` package does not require any installation steps under Unix-based systems or under Cygwin. Under native Windows, if you downloaded XSB from CVS, you need to compile the package as follows:

```
cd XSB\packages\sgml\cc
nmake /f NMakefile.mak
```

You need MS Visual Studio for that. If you downloaded a prebuilt version of XSB, then the `sgml` package should have already been compiled for you and no installation is required.

The details of the `xpath` package and the corresponding configuration instructions appear in Section 8.4.

¹This package has not yet been tested for thread-safety

8.2 Overview of the SGML Parser

The `sgml` package accepts input in the form of files, URLs and Prolog atoms. To load the `sgml` parser, the user should type

```
?- [sgml].
```

at the prompt. If `test.html` is a file with the following contents

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 3.2//EN">
```

```
<html>
```

```
<head>
```

```
<title>Demo</title>
```

```
</head>
```

```
<body>
```

```
<h1 align=center>This is a demo</h1>
```

```
<p>Paragraphs in HTML need not be closed.
```

```
<p>This is called 'omitted-tag' handling.
```

```
</body>
```

```
</html>
```

then the following call

```
?- load_html_structure(file('test.html'), Term, Warn).
```

will parse the document and bind `Term` to the following Prolog term:

```
[ element(html,
    [],
    [ element(head,
        [],
        [ element(title,
            [],
            [ 'Demo'
            ])
        ],
        element(body,
            [],
            [ '\n',
              element(h1,
                [ align = center
```

```

        ],
        [ 'This is a demo'
        ]),
    '\n\n',
    element(p,
        [],
        [ 'Paragraphs in HTML need not be closed.\n'
        ]),
    element(p,
        [],
        [ 'This is called 'omitted-tag\' handling.'
        ])
    ])
])
].

```

The XML document is converted into a list of Prolog terms of the form `element(Name,Attributes,Content)`. Each term corresponds to an XML element. *Name* represents the name of the element. *Attributes* is a list of attribute-value pairs of the element. *Content* is a list of child-elements and CDATA. For instance,

```
<aaa>foo<bbb>foo1</bbb></aaa>
```

will be parsed as

```
element(aaa,[],[foo, element(bbb,[],[foo1])])
```

Entities (e.g. `<`;) are returned as part of CDATA, unless they cannot be represented. See `load_sgml_structure/3` for details.

8.3 Predicate Reference

8.3.1 Loading Structured Documents

SGML, HTML, and XML documents are parsed by the predicate `load_structure/4`, which has many options. For convenience, a number of commonly used shorthands are provided to parse SGML, XML, HTML, and XHTML documents respectively.

```
load_sgml_structure(+Source, -Content, -Warn)
```

```
load_xml_structure(+Source, -Content, -Warn)
```

```
load_html_structure(+Source, -Content, -Warn)
```

```
load_xhtml_structure(+Source, -Content, -Warn)
```

The parameters of these predicates have the same meaning as those in `load__structure/4`, and are described below.

The above predicates (in fact, just `load_xml_structure/3` and `load_html_structure/3`) are the most commonly used predicates of the `sgml` package. The other predicates described in this section are needed only for advanced uses of the package.

load__structure(*+Source*, *-Content*, *+Options*, *-Warn*)

Source can have one of the following forms: `url(url)`, `file(file name)`, `string('document as a Prolog atom')`. The parsed document is returned in *Content*. *Warn* is bound to a (possibly empty) list of warnings generated during the parsing process. *Options* is a list of parameters that control parsing, which are described later.

The list *Content* can have the following members:

A Prolog atom

Atoms are used to represent character strings, i.e., CDATA.

element(*Name*, *Attributes*, *Content*)

Name is the name of the element tag. Since SGML is case-insensitive, all element names are returned as lowercase atoms.

Attributes is a list of pairs the form *Name=Value*, where *Name* is the name of an attribute and *Value* is its value. Values of type CDATA are represented as atoms. The values of multi-valued attributes (NAMES, etc.) are represented as a lists of atoms. Handling of the attributes of types NUMBER and NUMBERS depends on the setting of the `number(+NumberMode)` option of `set__sgml_parser/2` or `load__structure/3` (see later). By default the values of such attributes are represented as atoms, but the `number(...)` option can also specify that these values must be converted to Prolog integers.

Content is a list that represents the content for the element.

entity(*Code*)

If a character entity (e.g., `Α`) is encountered that cannot be represented in the Prolog character set, this term is returned. It represents the code of the encountered character (e.g., `entity(913)`).

entity(*Name*)

This is a special case of `entity(Code)`, intended to handle special symbols by their name rather than character code. If an entity refers to a character entity holding a single character, but this character cannot be represented in the Prolog character set, this term is returned. For example, if the contents of an element is `Α < Β` then it will be represented as follows:

```
[ entity('Alpha'), ' < ', entity('Beta') ]
```

Note that entity names are case sensitive in both SGML and XML.

sdata(*Text*)

If an entity with declared content-type SDATA is encountered, this term is used. The data of the entity instantiates *Text*.

ndata(*Text*)

If an entity with declared content-type `NDATA` is encountered, this term is used. The data instantiates *Text*.

pi(*Text*)

If a processing instruction is encountered (`<?...?>`), *Text* holds the text of the processing instruction. Please note that the `<?xml ...?>` instruction is ignored and is not treated as a processing instruction.

The *Options* parameter is a list that controls parsing. Members of that list can be of the following form:

dtd(*?DTD*)

Reference to a DTD object. If specified, the `<!DOCTYPE ...>` declaration supplied with the document is ignored and the document is parsed and validated against the provided DTD. If the DTD argument is a variable, then the variable *DTD* gets bound to the DTD object created out of the DTD supplied with the document.

dialect(*+Dialect*)

Specify the parsing dialect. The supported dialects are `sgml` (default), `xml` and `xmlns`.

space(*+SpaceMode*)

Sets the space handling mode for the initial environment. This mode is inherited by the other environments, which can override the inherited value using the XML reserved attribute `xml:space`. See Section 8.3.2 for details.

number(*+NumberMode*)

Determines how attributes of type `NUMBER` and `NUMBERS` are handled. If `token` is specified (the default) they are passed as an atom. If `integer` is specified the parser attempts to convert the value to an integer. If conversion is successful, the attribute is represented as a Prolog integer. Otherwise the value is represented as an atom. Note that SGML defines a numeric attribute to be a sequence of digits. The `-` (minus) sign is not allowed and `1` is different from `01`. For this reason the default is to handle numeric attributes as tokens. If conversion to integer is enabled, negative values are silently accepted and the minus sign is ignored.

defaults(*+Bool*)

Determines how default and fixed attributes from the DTD are used. By default, defaults are included in the output if they do not appear in the source. If `false`, only the attributes occurring in the source are emitted.

file(*+Name*)

Sets the name of the input file for error reporting. This is useful if the input is a stream that is not coming from a file. In this case, errors and warnings will not have the file name in them, and this option allows one to force inclusion of a file name in such messages.

line(*+Line*)

Sets the starting line-number for reporting errors. For instance, if `line(10)` is specified and an error is found at line *X* then the error message will say that the error occurred at line *X*+10. This option is used when the input stream does not start with the first line of a file.

max_errors(+Max)

Sets the maximum number of errors. The default is 50. If this number is reached, the following exception is raised:

```
error(limit_exceeded(max_errors, Max), _)
```

8.3.2 Handling of White Spaces

Four modes for handling white-spaces are provided. The initial mode can be switched using the `space(SpaceMode)` option to `load_structure/3` or `set_sgml_parser/2`. In XML mode, the mode is further controlled by the `xml:space` attribute, which may be specified both in the DTD and in the document. The defined modes are:

space(sgml)

Newlines at the start and end of an element are removed. This is the default mode for the SGML dialect.

space(preserve)

White space is passed literally to the application. This mode leaves all white space handling to the application. This is the default mode for the XML dialect.

space(default)

In addition to `sgml` space-mode, all consecutive whitespace is reduced to a single space-character.

space(remove)

In addition to `default`, all leading and trailing white-space is removed from CDATA objects. If, as a result, the CDATA becomes empty, nothing is passed to the application. This mode is especially handy for processing data-oriented documents, such as RDF. It is not suitable for normal text documents. Consider the HTML fragment below. When processed in this mode, the spaces surrounding the three elements in the example below are lost. This mode is not part of any standard: XML 1.0 allows only `default` and `preserve`.

Consider adjacent `bold` `and` `<it>italic</it>` words.

The parsed term will be `['Consider adjacent',element(b,[],[bold]),element(ul,[],[and]),element(it,[],[italics]),words]`.

8.3.3 XML documents

The parser can operate in two modes: the `sgml` mode and the `xml` mode, as defined by the `dialect(Dialect)` option. HTML is a special case of the SGML mode with a particular DTD. Regardless of this option, if the first line of the document reads as below, the parser is switched automatically to the XML mode.

```
<?xml ... ?>
```

Switching to XML mode implies:

- *XML empty elements*
The construct `<element attribute ... attribute/>` is recognized as an empty element.
- *Predefined entities*
The following entities are predefined: `<`; (`<`), `>`; (`>`), `&`; (`&`), `'`; (`'`) and `"`; (`"`).
- *Case sensitivity*
In XML mode, names of tags and attributes are case-sensitive, except for the DTD reserved names (i.e. `ELEMENT`, *etc.*).
- *Character classes*
In XML mode, underscore (`_`) and colon (`:`) are allowed in names.
- *White-space handling*
White space mode is set to **preserve**. In addition, the XML reserved attribute **xml:space** is honored; it may appear both in the document and the DTD. The **remove** extension (see `space(remove)` earlier) is allowed as a value of the **xml:space** attribute. For example, the DTD statement below ensures that the **pre** element preserves space, regardless of the default processing mode.

```
<!ATTLIST pre xml:space nmtoken #fixed preserve>
```

XML Namespaces

Using the dialect `xmlns`, the parser will recognize XML namespace prefixes. In this case, the names of elements are returned as a term of the format

URL:LocalName

If an identifier has no namespace prefix and there is no default namespace, it is returned as a simple atom. If an identifier has a namespace prefix but this prefix is undeclared, the namespace prefix rather than the related URL is returned.

Attributes declaring namespaces (`xmlns:ns=url`) are represented in the translation as regular attributes.

8.3.4 DTD-Handling

The DTD (**D**ocument **T**ype **D**efinition) are internally represented as objects that can be created, freed, defined, and inspected. Like the parser itself, it is filled by opening it as a Prolog output stream and sending data to it. This section summarizes the predicates for handling the DTD.

new_dtd(*+DocType*, *-DTD*, *-Warn*)

Creates an empty DTD for the named *DocType*. The returned DTD-reference is an opaque term that can be used in the other predicates of this package. *Warn* is the list of warnings generated.

free_dtd(*+DTD*, *-Warn*)

Deallocate all resources associated to the DTD. Further use of *DTD* is invalid. *Warn* is the list of warnings generated.

open_dtd(*+DTD*, *+Options*, *-Warn*)

This opens and loads a DTD from a specified location (given in the *Options* parameter (see next)). *DTD* represents the created DTD object after the source is loaded. *Options* is a list options. Currently the only option supported is *source(location)*, where *location* can be of one of these forms:

```
url(url)
file(fileName)
string('document as a Prolog atom').
```

dtd(*+DocType*, *-DTD*, *-Warn*)

Certain DTDs are part of the system and have known doctypes. Currently, 'HTML' and 'XHTML' are the only recognized built-in doctypes. Such a DTD can be used for parsing simply by specifying the doctype. Thus, the **dtd/3** predicate takes the doctype name, finds the DTD associated with the given doctype, and creates a dtd object for it. *Warn* is the list of warnings generated.

dtd(*+DocType*, *-DTD*, *+DtdFile* *-Warn*)

The predicate parses the DTD present at the location *DtdFile* and creates the corresponding DTD object. *DtdFile* can have one of the following forms: **url**(*url*), **file**(*fileName*), **string**('document as a Prolog atom').

8.3.5 Low-level Parsing Primitives

The following primitives are used only for more complex types of parsing, which might not be covered by the **load_structure/4** predicate.

new_sgml_parser(*-Parser*, *+Options*, *-Warn*)

Creates a new parser. *Warn* is the list of warnings generated. A parser can be used one or multiple times for parsing documents or parts thereof. It may be bound to a DTD or the DTD may be left implicit. In this case the DTD is created from the document prologue or (if it is not in the prologue) parsing is performed without a DTD. The *Options* list can contain the following parameters:

dtd(*?DTD*)

If *DTD* is bound to a DTD object, this DTD is used for parsing the document and

the document's prologue is ignored. If *DTD* is a variable, the variable gets bound to a created DTD. This DTD may be created from the document prologue or build implicitly from the document's content.

free_sgml_parser(*+Parser*, *-Warn*)

Destroy all resources related to the parser. This does not destroy the DTD if the parser was created using the `dtd(DTD)` option. *Warn* is the list of warnings generated during parsing (can be empty).

set_sgml_parser(*+Parser*, *+Option*, *-Warn*)

Sets attributes to the parser. *Warn* is the list of warnings generated. *Options* is a list that can contain the following members:

file(*File*)

Sets the file for reporting errors and warnings. Sets the `linenumber` to 1.

line(*Line*)

Sets the starting line for error reporting. Useful if the stream is not at the start of the (file) object for generating proper line-numbers. This option has the same meaning as in the `load_structure/4` predicate.

charpos(*Offset*)

Sets the starting character location. See also the `file(File)` option. Used when the stream does not start from the beginning of a document.

dialect(*Dialect*)

Set the markup dialect. Known dialects:

sgml

The default dialect. This implies markup is case-insensitive and standard SGML abbreviation is allowed (abbreviated attributes and omitted tags).

xml

This dialect is selected automatically if the processing instruction `<?xml ...>` is encountered.

xmlns

Process file as XML file with namespace support.

qualify_attributes(*Boolean*)

Specifies how to handle unqualified attributes (i.e., without an explicit namespace) in XML namespace (`xmlns`) dialect. By default, such attributes are not qualified with namespace prefixes. If `true`, such attributes are qualified with the namespace of the element they appear in.

space(*SpaceMode*)

Define the initial handling of white-space in `PCDATA`. This attribute is described in Section 8.3.2.

number(*NumberMode*)

If `token` is specified (the default), attributes of type number are represented as a Prolog atom. If `integer` is specified, such attributes are translated into Prolog integers. If the

conversion fails (e.g., due to an overflow) a warning is issued and the value is represented as an atom.

doctype(*Element*)

Defines the top-level element of the document. If a `<!DOCTYPE ...>` declaration has been parsed, this declaration is used. If there is no `DOCTYPE` declaration then the parser can be instructed to use the element given in `doctype(_)` as the top level element. This feature is useful when parsing part of a document (see the `parse` option to `sgml_parse/3`).

sgml_parse(*+Parser*, *+Options*, *-Warn*)

Parse an XML file. The parser can operate in two input and two output modes. Output is a structured term as described with `load_structure/4`.

Warn is the list of warnings generated. A full description of *Options* is given below.

document(*+Term*)

A variable that will be unified with a list describing the content of the document (see `load_structure/4`).

source(*+Source*)

Source can have one of the following forms: `url(url)`, `file(fileName)`, `string('document as a Prolog atom')`. This option *must* be given.

content_length(*+Characters*)

Stop parsing after the given number of *Characters*. This option is useful for parsing input embedded in *envelopes*, such as HTTP envelopes.

parse(*Unit*)

Defines how much of the input is parsed. This option is used to parse only parts of a file.

file

Default. Parse everything upto the end of the input.

element

The parser stops after reading the first element. Using `source(Stream)`, this implies reading is stopped as soon as the element is complete, and another call may be issued on the same stream to read the next element.

declaration

This may be used to stop the parser after reading the first declaration. This is useful if we want to parse only the `doctype` declaration.

max_errors(*+MaxErrors*)

Sets the maximum number of errors. If this number is exceeded, further writes to the stream will yield an I/O error exception. Printing of errors is suppressed after reaching this value. The default is 100.

syntax_errors(*+ErrorMode*)

Defines how syntax errors are handled.

quiet

Suppress all messages.

print

Default. Print messages.

8.3.6 External Entities

While processing an SGML document the document may refer to external data. This occurs in three places: external parameter entities, normal external entities and the DOCTYPE declaration. The current version of this tool deals rather primitively with external data. External entities can only be loaded from a file.

Two types of lines are recognized by this package:

```
DOCTYPE doctype file
```

```
PUBLIC "Id " file
```

The parser loads the entity from the file specified as *file*. The file can be local or a URL.

8.3.7 Exceptions

Exceptions are generated by the parser in two cases. The first case is when the user specifies wrong input. For example when specifying

```
load_structure( string('<m></m>'), Document, [line(xyz)], Warn)
```

The string *xyz* is not in the domain of *line*. Hence in this case a domain error exception will be thrown.

Exceptions are generated when XML being parsed is not well formed. For example if the input XML contains

```
'<m></m1>'
```

exceptions will be thrown.

In both cases the format of the exception is

```
error( sgml( error term), error message)
warning( sgml( warning term), warning message)
```

where *error term* or *warning term* can be of the form

- *pointer to the parser instance,*
- *line at which error occurred,*

- *error code*.
- *functor(argument)*, where *functor* and *argument* depend on the type of exception raised. For example,

```

resource-error(no-memory) — if memory is unavailable
permission-error(file-name) — no permission to read a file
A system-error(description) -- internal system error
type-error(expected,actual) — data type error
domain-error(functor,offending-value) — the offending value is not in the domain
of the functor. For instance, in load_structure( string('<m></m>'), Document,
[line(xyz)], Warn), xyz is not in the domain of line.
existence-error(resource) — resource does not exist
limit-exceeded(limit,maxval) — value exceeds the limit.

```

8.3.8 Unsupported features

The current parser is rather limited. While it is able to deal with many serious documents, it omits several less-used features of SGML and XML. Known missing SGML features include

- *NOTATION on entities*
Though notation is parsed, notation attributes on external entity declarations are not represented in the output.
- *NOTATION attributes*
SGML notations may have attributes, declared using `<!ATTLIST #NOT name attrib>`. Those data attributes are provided when you declare an external CDATA, NDATA, or SDATA entity. XML does not support external CDATA, NDATA, or SDATA entities, nor any of the other uses to which data attributes are put in SGML.
- *SGML declaration*
The 'SGML declaration' is fixed, though most of the parameters are handled through indications in the implementation.
- *The RANK feature*
It is regarded as obsolete.
- *The LINK feature*
It is regarded as too complicated.
- *The CONCUR feature*
Concurrent markup allows a document to be tagged according to more than one DTD at the same time. It is not supported.
- *The Catalog files*
Catalog files are not supported.

In the XML mode, the parser recognizes SGML constructs that are not allowed in XML. Also various extensions of XML over SGML are not yet realized. In particular, XInclude is not implemented.

8.3.9 Summary of Predicates

dtd/2	Find or build a DTD for a document type
free_dtd/1	Free a DTD object
free_sgml_parser/1	Destroy a parser
load_dtd/2	Read DTD information from a file
load_structure/4	Parse XML/SGML/HTML data into Prolog term
load_sgml_structure/3	Parse SGML file into Prolog term
load_html_structure/3	Parse HTML file into Prolog term
load_xml_structure/3	Parse XML file into Prolog term
load_xhtml_structure/3	Parse XHTML file into Prolog term
new_dtd/2	Create a DTD object
new_sgml_parser/2	Create a new parser
open_dtd/3	Open a DTD object as an output stream
set_sgml_parser/2	Set parser options (dialect, source, <i>etc.</i>)
sgml_parse/2	Parse the input
xml_name/1	Test atom for valid XML name
xml_quote_attribute/2	Quote text for use as an attribute
xml_quote_cdata/2	Quote text for use as PCDATA

8.4 XPath support

XPath is a query language for addressing parts of an XML document. In XSB, this support is provided by the `xpath` package. To use this package the `libxml2` XML parsing library must be installed on the machine. It comes with most Linux distributions, since it is part of the Gnome desktop, or one can download it from <http://xmlsoft.org/>. It is available for Linux, Solaris, Windows, and MacOS. Note that both the library itself and the `.h` files of that library must be installed. In some Linux distributions, the `.h` files might reside in a separate package from the package that contains the actual library. For instance, the library (`libxml2.so`) might be in the package called `libxml2` (which is usually installed by default), while the `.h` files might be in the package `libxml2-dev` (which is usually *not* in default installations).

On Unix-based systems (and MacOS), the package might need to be configured at the time XSB is configured using XSB's `configure` script found in the XSB's `build` directory. Normally, if `libxml2` is installed by a Linux package manager, nothing special is required: the package will be configured by default. If the library is in a non-standard place, then the configure option `-with-xpath-dir=directory-of-libxml2` must be given. It must specify the directory where `lib/*libxml2.so` (or `libxml2.dylib` in Mac) and `include/libxml2` can be found.

Examples: If `libxml2` is in a default location, then XSB can be configured simply like this:

```
./configure
```

Otherwise, use

```
./configure --with-xpath-dir=/usr/local
```

if, for example, `libxml2.so` is in `/usr/local/lib/i386-linux-gnu/libxml2.so` and the included `.h` files are in `/usr/local/include/libxml2/*`.

On Windows and under Cygwin, the `libxml2` library is already included in the XSB distribution and does not need to be downloaded. If you are using a prebuilt XSB distribution for Windows, then you do not need to do anything—the package has already been built for you.

For Cygwin, you only need to run the `./configure` script without any options. This needs to be done regardless of whether you downloaded XSB from CVS or a released prebuilt version.

If you downloaded XSB from CVS and want to use it under native Windows (not Cygwin), then you would need to compile the XPath package, and you need Microsoft's Visual Studio. To compile the package one should do the following:

```
cd packages\xpath\cc
nmake /f NMakefile.mak
```

The following section assumes that the reader is familiar with the syntax of XPath and its capabilities. To load the `xpath` package, type

```
:-[xpath].
```

The program needs to include the following directive:

```
:- import parse_xpath/4 from xpath.
```

XPath query evaluation is done by using the `parse_xpath` predicate.

parse_xpath(*+Source*, *+XPathQuery*, *-Output*, *+NamespacePrefixList*)

Source is a term of the format `url(url)`, `file(filename)` or `string('XML-document-as-a-string')`. It specifies that the input XML document is contained in a file, can be fetched from a URL, or is given directly as a Prolog atom.

XPathQuery is a standard XPath query which is to be evaluated on the XML document in *Source*.

Output gets bound to the output term. It represents the XML element returned after the XPath query is evaluated on the XML document in *Source*. The output term is of the form `string('XML-document')`. It can then be parsed using the `sgml` package described earlier.

NamespacePrefixList is a space separated list of pairs of the form *prefix = namespace*. This specifies the namespace prefixes that are used in the XPath query.

For example if the xpath expression is `'/x:html/x:head/x:meta'` where `x` is a prefix that stands for `'http://www.w3.org/1999/xhtml'`, then `x` would have to be defined as follows:

```
?- parse_xpath(url('http://w3.org'), '/x:html/x:head/x:meta', 04,  
               'x=http://www.w3.org/1999/xhtml').
```

In the above, the xpath query is `'/x:html/x:head/x:meta'` and the prefix has been defined as `'x=http://www.w3.org/1999/xhtml'`.

Chapter 9

rdf: The XSB RDF Parser

By Aneesh Ali

9.1 Introduction

RDF is a W3C standard for representing meta-data about documents on the Web as well as exchanging frame-based data (e.g. ontologies). RDF has a formal data model defined in terms of *triples*. In addition, a *graph* model is defined for visualization and an XML serialization for exchange. This chapter describes the API provided by the XSB RDF parsing package. The package and its documentation are adaptations from SWI Prolog.

9.2 High-level API

The RDF translator is built in Prolog on top of the **sgml2pl** package, which provides XML parsing. The transformation is realized in two passes. It is designed to operate in various environments and therefore provides interfaces at various levels. First we describe the top level, which parses RDF-XML file into a list of triples. These triples are *not* asserted into the Prolog database because it is not necessarily the final format the user wishes to use and it is not clear how the user might want to deal with multiple RDF documents. Some options are using global URI's in one pool, in Prolog modules, or using an additional argument.

load_rdf(*+File*, *-Triples*)

Same as `load_rdf(+File, -Triples, [])`.

load_rdf(*+File*, *-Triples*, *+Options*)

Read the RDF-XML file *File* and return a list of *Triples*. *Options* is a list of additional processing options. Currently defined options are:

base_uri(*BaseURI*)

If provided, local identifiers and identifier-references are globalized using this URI. If

omitted, local identifiers are not tagged.

blank_nodes(*Mode*)

If *Mode* is **share** (default), blank-node properties (i.e. complex properties without identifier) are reused if they result in exactly the same triple-set. Two descriptions are shared if their intermediate description is the same. This means they should produce the same set of triples in the same order. The value **noshare** creates a new resource for each blank node.

expand_foreach(*Boolean*)

If *Boolean* is **true**, expand **rdf:aboutEach** into a set of triples. By default the parser generates **rdf(each(Container), Predicate, Subject)**.

lang(*Lang*)

Define the initial language (i.e. pretend there is an **xml:lang** declaration in an enclosing element).

ignore_lang(*Bool*)

If **true**, **xml:lang** declarations in the document are ignored. This is mostly for compatibility with older versions of this library that did not support language identifiers.

convert_typed_literal(*:ConvertPred*)

If the parser finds a literal with the **rdf:datatype=Type** attribute, call *ConvertPred(+Type, +Content, -Literal)*. *Content* is the XML element contents returned by the XML parser (a list). The predicate must unify *Literal* with a Prolog representation of *Content* according to *Type* or throw an exception if the conversion cannot be made.

This option serves two purposes. First of all it can be used to ignore type declarations for backward compatibility of this library. Second it can be used to convert typed literals to a meaningful Prolog representation (e.g., convert '42' to the Prolog integer 42 if the type is **xsd:int** or a related type).

namespaces(*-List*)

Unify *List* with a list of *NS=URL* for each encountered **xmlns:NS=URL** declaration found in the source.

entity(*+Name, +Value*)

Override entity declaration in file. As it is common practice to declare namespaces using entities in RDF/XML, this option allows changing the namespace without changing the file. Multiple such options are allowed.

The *Triples* list is a list of the form **rdf(Subject, Predicate, Object)** triples. *Subject* is either a plain resource (an atom), or one of the terms **each(URI)** or **prefix(URI)** with the usual meaning. *Predicate* is either a plain atom for explicitly non-qualified names or a term *NameSpace:Name*. If *NameSpace* is the defined RDF name space it is returned as the atom **rdf**. *Object* is a URI, a *Predicate* or a term of the form **literal(Value)** for literal values. *Value* is either a plain atom or a parsed XML term (list of atoms and elements).

9.2.1 RDF Object representation

The *Object* (3rd) part of a triple can have several different types. If the object is a resource it is returned as either a plain atom or a term *NameSpace:Name*. If it is a literal it is returned as

`literal(Value)`, where *Value* can have one of the form below.

- An atom
If the literal *Value* is a plain atom is a literal value not subject to a datatype or `xml:lang` qualifier.
- `lang(LanguageID, Atom)`
If the literal is subject to an `xml:lang` qualifier *LanguageID* specifies the language and *Atom* the actual text.
- A list
If the literal is an XML literal as created by `parseType="Literal"`, the raw output of the XML parser for the content of the element is returned. This content is a list of `element(Name, Attributes, Content)` and atoms for CDATA parts as described with the `sgml` package.
- `type(Type, StringValue)`
If the literal has an `rdf:datatype=Type` a term of this format is returned.

9.2.2 Name spaces

RDF name spaces are identified using URIs. Unfortunately various URI's are in common use to refer to RDF. The RDF parser therefore defines the `rdf_name_space/1` predicate as `multifile`, which can be extended by the user. For example, to parse Netscape OpenDirectory (<http://www.mozilla.org/rdf/doc/inference.html>) given in the `structure.rdf` file (<http://rdf.dmoz.org/rdf/structure.rdf.u8.gz>), the following declarations are used:

```
:- multifile
    rdf_parser:rdf_name_space/1.

rdf_parser:rdf_name_space('http://www.w3.org/TR/RDF/').
rdf_parser:rdf_name_space('http://directory.mozilla.org/rdf').
rdf_parser:rdf_name_space('http://dmoz.org/rdf').
```

The above statements will then extend the initial definition of this predicate provided by the parser:

```
rdf_name_space('http://www.w3.org/1999/02/22-rdf-syntax-ns#').
rdf_name_space('http://www.w3.org/TR/REC-rdf-syntax').
```

9.2.3 Low-level access

The predicates `load_rdf/2` and `load_rdf/3` described earlier are not always sufficient. For example, they cannot deal with documents where the RDF statement is embedded in an XML document. It also cannot deal with really large documents (e.g. the Netscape OpenDirectory project, currently about 90 MBytes), without requiring huge amounts of memory.

For really large documents, the **sgml2pl** parser can be instructed to handle the content of a specific element (i.e. `<rdf:RDF>`) element-by-element. The parsing primitives defined in this section can be used to process these one-by-one.

xml_to_rdf(*+XML*, *+BaseURI*, *-Triples*)

Process an XML term produced by **sgml**'s `load_structure/4` using the `dialect(xmlns)` output option. *XML* is either a complete `<rdf:RDF>` element, a list of RDF-objects (container or description), or a single description of container.

9.3 Testing the RDF translator

A test-suite and a driver program are provided by `rdf_test.P` in the `XSB/examples/rdf` directory. To run these tests, load this file into Prolog and execute `test_all`. The test files found in the directory `examples/rdf/suite` are then converted into triples. The expected output is in `examples/rdf/expectedoutput`. One can also run the tests selectively, using the following predicates:

suite(*+N*)

Run test *N* using the file `suite/tN.rdf` and display its RDF representation and the triples.

test_file(*+File*)

Process *File* and display its RDF representation and the triples.

Chapter 10

Constraint Packages

Constraint packages are an important part of modern logic programming, but approaches to constraints differ both in their semantics and in their implementation. At a semantic level, *Constraint Logic Programming* associates constraints with logical variables, and attempts to determine solutions that are inconsistent with or entailed by those constraints. At an implementational level, the constraints can either be manipulated by accessing attributed variables or by adding *constraint handling rules* to a program. The former approach of attributed variables can be much more efficient than constraint handling rules (which are themselves implemented through attributed variables) but are much more difficult to use than constraint handling rules. These variable-based approaches differ from that of *Answer Set Programming* in which a constraint problem is formulated as a set of rules, which are consistent if a stable model can be constructed for them.

XSB supports all of these approaches. Two packages based on attributed variables are presented in this chapter: CLP(R) and the `bounds` package, which provides a simple library for handling finite domains. XSB's CHR package is described in Chapter 11, and XSB's Answer Set Programming Package, XASP is described in Chapter 12.

Before describing the individual packages, we note that these packages can be freely used with variant tabling, the mechanisms for which handle attributed variables. However in Version 3.4, calling a predicate P that is tabled using call subsumption will raise an error if the call to P contains any constrained variables (attributed variables).

10.1 clpr: The CPL(R) package

The CLP(R) library supports solutions of linear equations and inequalities over the real numbers and the lazy treatment of nonlinear equations¹. In displaying sets of equations and disequations, the library removes redundancies, performs projections, and provides for linear optimization. The goal of the XSB port is to provide the same CLP(R) functionality as in other platforms, but also to allow constraints to be used by tabled predicates. This section provides a general introduction

¹The CLP(R) package is based on the clpqr package included in SWI Prolog version 5.6.49. This package was originally written by Christian Holzbaur and ported to SWI by Leslie De Koninck. Terrance Swift ported the package to XSB and wrote this XSB manual section.

to the CLP(R) functionality available in XSB, for further information on the API described in Section 10.1.1 see <http://www.ai.univie.ac.at/clpqr>, or the Sicstus Prolog manual (the CLP(R) library should behave similarly on XSB and Sicstus at the level of this API).

The `clpr` package may be loaded by the command `[clpr]`. Loading the package imports exported predicates from the various files in the `clpr` package into `usermod` (see Volume 1, Section 3.3) so that they may be used in the interpreter. Modules that use the exported predicates need to explicitly import them from the files in which they are defined (e.g. `bv`, as shown below).

XSB's tabling engine supports the use of attributed variables (Section 1.2), which in turn have been used to port real constraints to XSB under the CLP(R) library of Christian Holzbauer [?]. Constraint equations are represented using the Prolog syntax for evaluable functions (Volume 1, Section 6.2.1). Formally:

<i>ConstraintSet</i> ->	<i>C</i> <i>C</i> , <i>C</i>	
<i>C</i> ->	<i>Expr</i> ::= <i>Expr</i>	equation
	<i>Expr</i> = <i>Expr</i>	equation
	<i>Expr</i> < <i>Expr</i>	strict inequation
	<i>Expr</i> > <i>Expr</i>	strict inequation
	<i>Expr</i> =< <i>Expr</i>	nonstrict inequation
	<i>Expr</i> >= <i>Expr</i>	nonstrict inequation
	<i>Expr</i> /= <i>Expr</i>	disequation
<i>Expr</i> ->	<i>variable</i>	Prolog variable
	<i>number</i>	floating point number
	+ <i>Expr</i>	
	- <i>Expr</i>	
	<i>Expr</i> + <i>Expr</i>	
	<i>Expr</i> - <i>Expr</i>	
	<i>Expr</i> * <i>Expr</i>	
	<i>Expr</i> / <i>Expr</i>	
	abs(<i>Expr</i>)	
	sin(<i>Expr</i>)	
	cos(<i>Expr</i>)	
	tan(<i>Expr</i>)	
	pow(<i>Expr</i> , <i>Expr</i>)	raise to the power
	exp(<i>Expr</i> , <i>Expr</i>)	raise to the power
	min(<i>Expr</i> , <i>Expr</i>)	minimum of two expressions
	max(<i>Expr</i> , <i>Expr</i>)	maximum of two expressions
	#(<i>Expr</i>)	symbolic numerical constants

10.1.1 The CLP(R) API

From the command line, it is usually easiest to load the `clpr` package and call the predicates below directly from `usermod` (the module implicitly used by the command line). However, when calling

```

:- import {}/1 from clpr.

root(N, R) :-
  root(N, 1, R).
  root(0, S, R) :- !, S=R.
  root(N, S, R) :-
N1 is N-1,
{ S1 = S/2 + 1/S },
root(N1, S1, R).

```

Figure 10.1: Example of a file with a CLP(R) predicate

any of these predicates from compiled code, they must be explicitly imported from their modules (e.g. `{}` must be explicitly imported from `clpr`). Figure 10.1.1 shows an example of how this is done. ‘

```
{+Constraints}                                     module: clpr
```

When the CLP(R) package is loaded, inclusion of equations in braces (`{}`) adds **Constraints** to the constraint store where they are checked for satisfiability.

Example:

```

| ?- [clpr].
[clpr loaded]
[itf loaded]
[dump loaded]
[bv_r loaded]
[nf_r loaded]

yes

| ?- {X = Y+1, Y = 3*X}.

X = -0.5000
Y = -1.5000;

yes

```

Error Cases

- **Constraints** is not instantiated
 - `instantiation_error`
- **Constraints** is not an equation, an inequation or a disequation
 - `domain_error('constraint relation',Rel)`
- **Constraints** contains an expression `Expr` that is not a numeric expression
 - `domain_error('numeric expression',Expr)`

entailed(+Constraint) module: clpr
 Succeeds if **Constraint** is logically implied by the current constraint store. **entailed/1** does not change the constraint store.

Example:

```
| ?- {A =< 4},entailed(A =\= 5).
{ A =< 4.0000 }
```

yes

Error Cases

- **Constraints** is not instantiated
 - **instantiation_error**
- **Constraints** is not an equation, an inequation or a disequation
 - **domain_error('constraint relation',Rel)**

inf(+Expr,-Val) clpr
sup(+Expr,-Val) clpr
minimize(Expr) clpr
maximize(Expr) module: clpr

These four related predicates provide various mechanisms to compute the maximum and minimum of expressions over variables in a constraint store. In the case where the expression is not bounded from above over the reals **sup/2** and **maximize/1** will fail; similarly if the expression is not bounded from below **inf/2** and **minimize/1** will fail.

Examples:

```
| ?- {X = 2*Y,Y >= 7},inf(X,F).
{ X >= 14.0000 }
{ Y = 0.5000 * X }
```

```
X = _h8841
Y = _h9506
F = 14.0000
```

```
| ?- {X = 2*Y,Y >= 7},minimize(X).
X = 14.0000
Y = 7.0000
```

```
| ?- {X = 2*Y,Y =< 7},maximize(X-2).
```

```
X = 14.0000
Y = 7.0000
```

```
| ?- {X = 2*Y,Y =< 7},sup(X-2,Z).
{ X =< 14.0000 }
{ Y = 0.5000 * X }
```

```

X = _h8975
Y = _h9640
Z = 12.0000

```

```

yes
| ?- {X = 2*Y,Y =< 7},maximize(X-2).

```

```

X = 14.0000
Y = 7.0000

```

```

yes

```

```

inf(+Expr,-Val, +Vector, -Vertex)                                clpr
sup(+Expr,-Val, +Vector, -Vertex)                                module: clpr

```

These predicates work like `inf/2` and `sup/2` with the following addition. `Vector` is a list of Variables, and for each variable V in `Vector`, the value of V at the extremal point `Val` is returned in corresponding position in the list `Vertex`.

Example:

```

| ?= { 2*X+Y =< 16, X+2*Y =< 11,X+3*Y =< 15, Z = 30*X+50*Y},
      sup(Z, Sup, [X,Y], Vertex).
{ X + 3.0000 * Y =< 15.0000 }
{ X + 0.5000 * Y =< 8.0000 }
{ X + 2.0000 * Y =< 11.0000 }
{ Z = 30.0000 * X + 50.0000 * Y }

```

```

X = _h816
Y = _h869
Z = _h2588
Sup = 310.0000
Vertex = [7.0000,2.0000]

```

```

bb_inf(+IntegerList,+Expr,-Inf,-Vertex, +Eps)                    module: clpr

```

Works like `inf/2` in `Expr` but assumes that all the variables in `IntegerList` have integral values. `Eps` is a positive number between 0 and 0.5 that specifies how close an element of `IntegerList` must be to an integer to be considered integral – i.e. for such an X , $\text{abs}(\text{round}(X) - X) < \text{Eps}$. Upon success, `Vertex` is instantiated to the integral values of all variables in `IntegerList`. `bb_inf/5` works properly for non-strict inequalities only.

Example:

```

| ?- {X > Y + Z,Y > 1, Z > 1},bb_inf([Y,Z],X,Inf,Vertex,0).
{ Z > 1.0000 }
{ Y > 1.0000 }
{ X - Y - Z > 0.0000 }

```

```

X = _h14286
Y = _h10914
Z = _h13553

```



```

Inf = 4.0000
Vertex = [2.0000,2.0000]

```

```
yes
```

Error Cases

- IntegerList is not instantiated
 - instantiation_error

```
bb_inf(+IntegerList,+Expr,-Inf)                                module: clpr
Works like bb_inf/5, but with the neighborhood, Eps, set to 0.001.
```

Example

```

|?- {X >= Y+Z, Y > 1, Z > 1}, bb_inf([Y,Z],X,Inf)
{ Z > 1.0000 }
{ Y > 1.0000 }
{ X - Y - Z >= 0.0000 }

```

```

X = _h14289
Y = _h10913
Z = _h13556
Inf = 4.

```

```
yes
```

```
dump(+Variables,+NewVars,-CodedVars)                          module: dump
For a list of variables Variables and a list of variable names NewVars, returns in CodedVars
the constraints on the variables, without affecting the constraint store.
```

Example:

```

| ?- {X > Y+1, Y > 2},
    dump([X,Y], [x,y], CS).
{ Y > 2.0000 }
{ X - Y > 1.0000 }

```

```

X = _h17748
Y = _h17139
CS = [y > 2.0000,x - y > 1.0000];

```

Error Cases

- Variables is not instantiated to a list of variables
 - instantiation_error

```
projecting_assert(+Clause)                                    module: dump
In XSB, when a subgoal is tabled, the tabling system automatically determines the relevant
projected constraints for an answer and copies them into and out of a table. However,
```

when a clause with constrained variables is asserted, this predicate must be used rather than `assert/1` in order to project the relevant constraints. This predicate works with either standard or trie-indexed dynamic code.

Example:

```
| ?- {X > 3},projecting_assert(q(X)).
{ X > 3.0000 }

X = _h396

yes
| ?- listing(q/1).
q(A) :-
    clpr : {A > 3.0000}.

yes
| ?- q(X),entailed(X > 2).
{ X > 3.0000 }

X = _h358

yes
| ?- q(X),entailed(X > 4).

no
```

10.2 The bounds Package

Version 3.4 of XSB does not support a full-fledged CLP(FD) package. However it does support a simplified package that maintains an upper and lower bound for logical variables. `bounds` can thus be used for simple constraint problems in the style of finite domains, as long as these problems that do not rely on too heavily on propagation of information about constraint domains ²

Perhaps the simplest way to explain the functionality of `bounds` is by example. The query

```
|?- X in 1..2,X #> 1.
```

first indicates via `X in 1..2` that the lower bound of `X` is 1 and the higher bound 2, and then constraints `X`, which is not yet bound, to be greater than 1. Applying this latter constraint to `X` forces the lower bound to equal the upper bound, instantiating `X`, so that the answer to this query is `X = 2`.

Next, consider the slightly more complex query

```
|?- X in 1..3,Y in 1..3,Z in 1..3,all_different([X,Y,Z]),X = 1, Y = 2.
```

²The `bounds` package was written by Tom Schrijvers, and ported to XSB from SWI Prolog version 5.6.49 by Terrance Swift, who also wrote this manual section.

`all_different/3` constraints `X`, `Y` and `Z` each to be different, whatever their values may be. Accordingly, this constraint together with the bound restrictions, implies that instantiating `X` and `Y` also causes the instantiation of `Z`. In a similar manner, the query

```
|?- X in 1..3,Y in 1..3,Z in 1..3,sum([X,Y,Z],#=:9),
```

constrains the sum of the three variables to equal 9 – and in this case assigns them a concrete value due to their domain restrictions.

In many constraint problems, it does not suffice to know whether a set of constraints is satisfiable; rather, concrete values may be needed that satisfy all constraints. One way to produce such values is through the predicate `labelling/2`

```
|?- X in 1..5,Y in 1..5,X #< Y,labeling([max(X)], [X,Y]))
```

In this query, it is specified that `X` and `Y` are both to be instantiated not just by any element of their domains, but by a value that assigns `X` to be the maximal element consistent with the constraints. Accordingly `X` is instantiated to 4 and `Y` to 5.

Because constraints in `bounds` are based on attributed variables which are handled by XSB's variant tabling mechanisms, constrained variables can be freely used with variant tabling as the following fragment shows:

```
table_test(X):- X in 2..3,p(X).
```

```
:- table p/1.
```

```
p(X):- X in 1..2.
```

```
?- table_test(Y).
```

```
Y = 2
```

For a more elaborate example, we turn to the *SEND MORE MONEY* example, in which the problem is to assign numbers to each of the letters *S,E,N,D,M,O,R,Y* so that the number *SEND* plus the number *MORE* equals the number *MONEY*. Borrowing a solution from the SWI manual [?], the `bounds` package solves this problem as:

```
send([S,E,N,D], [M,O,R,E], [M,O,N,E,Y]) :-
    Digits = [S,E,N,D,M,O,R,Y],
    Carries = [C1,C2,C3,C4],
    Digits in 0..9,
    Carries in 0..1,
    M #= C4,
    0 + 10 * C4 #= M + S + C3,
    N + 10 * C3 #= 0 + E + C2,
    E + 10 * C2 #= R + N + C1,
    Y + 10 * C1 #= E + D,
    M #>= 1,
    S #>= 1,
    all_different(Digits),
    label(Digits).
```

In many cases, it may be useful to test whether a given constraint is true or false. This can be done by unifying a variable with the truth value of a given constraint – i.e. by *reifying* the constraint. As an example, the query

```
|?- X in 1..10, Y in 1..10, Z in 0..1, X #< Y, X #= Y #<=> Z, label([Z]).
```

sets the bounded variable `Z` to the truth value of `X #= Y`, or 0³.

A reader familiar with the finite domain library of Sicstus [?] will have noticed that the syntax of **bounds** is consistent with that library. It is important to note however, **bounds** maintains only the upper and lower bounds of a variables as its attributes, (along, of course with constraints on those variables) rather than an explicit vector of permissible values. As a result, **bounds** may not be suitable for large or complex constraint problems.

10.2.1 The bounds API

Note that **bounds** does not perform error checking, but instead relies on the error checking of lower-level comparison and arithmetic operators.

in(-Variable, +Bound) bounds

Adds the constraint **Bound** to **Variable**, where **Bound** should be of the form **Low..High**, with **Low** and **High** instantiated to integers. This constraint ensures that any value of **Variable** must be greater than or equal to **Low** and less than or equal to **High**. Unlike some finite-domain constraint systems, it does *not* materialize a vector of currently allowable values for **Variable**.

Variables that have not had their domains explicitly constrained are considered to be in the range `min_integer..max_integer`.

#>(Expr1, Expr2) bounds

#<(Expr1, Expr2) bounds

#>=(Expr1, Expr2) bounds

#<=(Expr1, Expr2) bounds

#=(Expr1, Expr2) bounds

#=(Expr1, Expr2) bounds

Ensures that a given relation holds between **Expr1** and **Expr2**. Within these constraints, expressions may contain the functions `+/2`, `-/2`, `*/2`, `+/2`, `+/2`, `+/2`, `mod/2`, and `abs/1` in addition to integers and variables.

#<=>(Const1, Const2) bounds

#=>(Const1, Const2) bounds

#<=(Const1, Const2) bounds

Constrains the truth-value of **Const1** to have the specified logical relation (“iff”, “only-if” or “if”) to **Const2**, where **Const1** and **Const2** have one of the six relational operators above.

³The current version of the **bounds** package does not always seem to propagate entailment into the values of reified variables.

all_different(+VarList) bounds
 VarList must be a list of variables: constrains all variables in VarList to have different values.

sum(VarList,Op,?Value) bounds
 VarList must be a list of variables and Value an integer or variable: constrains the sum of all variables in VarList to have the relation Op to Value (see preceding example).

labeling(+Opts,+VarList) bounds
 This predicate succeeds if it can assign a value to each variable in VarList such that no constraint is violated. Note that assigning a value to each constrained variable is equivalent to deriving a solution that satisfies all constraints on the variables, which may be intractible depending on the constraints. Opts allows some control over how value assignment is performed in deriving the solution.

- **leftmost** Assigns values to variables in the order in which they occur. For example the query:

```
|?- X in 1..4,Y in 1..3,X #< Y,labeling([leftmost],[X,Y]),writeln([X,Y]),fail.
[1,2]
[1,3]
[2,3]
```

no

instantiates X and Y to all values that satisfy their constraints, and does so by considering each value in the domain of X, checking whether it violates any constraints, then considering each value of Y and checking whether it violates any constraints.

- **ff** This “first-fail” strategy assigns values to variables based on the size of their domains, from smallest to largest. By adopting this strategy, it is possible to perform a smaller search for a satisfiable solution because the most constrained variables may be considered first (though the bounds of the variable are checked rather than a vector of allowable values).
- **min** and **max** This strategy labels variables in the order of their minimal lower bound or maximal upper bound.
- **min(Expr)** and **max(Expr)** This strategy labels the variables so that their assignment causes Expr to have a minimal or maximal value. Consider for example how these strategies would affect the labelling of the preceding query:

```
|?- X in 1..4,Y in 1..3,X #< Y,labeling([min(Y)],[X,Y]),writeln([X,Y]),fail.
[1,2]
```

no

```
|?- X in 1..4,Y in 1..3,X #< Y,labeling([max(X)],[X,Y]),writeln([X,Y]),fail.
[2,3]
```

no

label(+VarList) bounds
 Shorthand for labeling([leftmost],+VarList).

indomain(?Var) bounds

Unifies **Var** with an element of its domain, and upon successive backtracking, with all other elements of its domain.

serialized(+BeginList,+Durations) bounds

serialized/2 can be useful for scheduling problems. As input it takes a list of variables or integers representing the beginnings of temporal events, along with a list of non-negative integers indicating the duration of each event in **BeginList**. The effect of this predicate is to constrain each of the events in **BeginList** to have a start time such that their durations do not overlap. As an example, consider the query

```
|?- X in 1..10, Y in 1..10, serialized([X,Y],[8,1]),label([X,Y]),writeln((X,Y)),fail.
```

In this query event **X** is taken to have duration of 8 units, while event **Y** is taken to have duration of 1 unit. Executing this query will instantiate **X** and **Y** to many different values, such as (1,9), (1,10), and (2,10) where **X** is less than **Y**, but also (10,1), (10,2) and many others where **Y** is less than **X**. Refining the query as

```
X in 1..10, Y in 1..10, serialized([X,Y],[8,1]),X #< Y,label([X,Y]),writeln((X,Y)),fail.
```

removes all solutions where **Y** is less than **X**.

lex_chain(+List) bounds

lex_chain/1 takes as input a list of lists of variables and integers, and enforces the constraint that each element in a given list is less than or equal to the elements in all succeeding lists. As an example, consider the query

```
|?- X in 1..3,Y in 1..3,lex_chain([[X],[2],[Y]]),label([X,Y]),writeln([X,Y]),fail.  
[1,2]  
[1,3]  
[2,2]  
[2,3]
```

lex_chain/1 ensures that **X** is less than or equal to 2 which is less than or equal to **Y**.

Chapter 11

Constraint Handling Rules

11.1 Introduction

Constraint Handling Rules (CHR) is a committed-choice bottom-up language embedded in XSB. It is designed for writing constraint solvers and is particularly useful for providing application-specific constraints. It has been used in many kinds of applications, like scheduling, model checking, abduction, type checking among many others.

CHR has previously been implemented in other Prolog systems (SICStus, Eclipse, Yap, hProlog), Haskell and Java. The XSB CHR system is based on the hProlog CHR system.

In this documentation we restrict ourselves to giving a short overview of CHR in general and mainly focus on XSB-specific elements. For a more thorough review of CHR we refer the reader to [?]. More background on CHR can be found at [?].

In Section 11.2 we present the syntax of CHR in XSB and explain informally its operational semantics. Next, Section 11.3 deals with practical issues of writing and compiling XSB programs containing CHR. Section 11.4 provides a few useful predicates to inspect the constraint store and Section 11.5 illustrates CHR with two example programs. How to combine CHR with tabled predicates is covered in Section 11.6. Finally, Section 11.7 concludes with a few practical guidelines for using CHR.

11.2 Syntax and Semantics

11.2.1 Syntax

The syntax of CHR rules in XSB is the following:

```
rules --> rule, rules.  
rules --> [].  
  
rule --> name, actual_rule, pragma, [atom('.'.)].
```

```

name --> xsb_atom, [atom('@')].
name --> [].

actual_rule --> simplification_rule.
actual_rule --> propagation_rule.
actual_rule --> simpagation_rule.

simplification_rule --> constraints, [atom('<=>')], guard, body.
propagation_rule --> constraints, [atom('==>')], guard, body.
simpagation_rule --> constraints, [atom('\')], constraints, [atom('<=>')],
                    guard, body.

constraints --> constraint, constraint_id.
constraints --> constraint, [atom(',')], constraints.

constraint --> xsb_compound_term.

constraint_id --> [].
constraint_id --> [atom('#')], xsb_variable.

guard --> [].
guard --> xsb_goal, [atom('|')].

body --> xsb_goal.

pragma --> [].
pragma --> [atom('pragma')], actual_pragmas.

actual_pragmas --> actual_pragma.
actual_pragmas --> actual_pragma, [atom(',')], actual_pragmas.

actual_pragma --> [atom('passive(')], xsb_variable, [atom(')')].

```

Additional syntax-related terminology:

- **head:** the constraints in an `actual_rule` before the arrow (either `<=>` or `==>`)

11.2.2 Semantics

In this subsection the operational semantics of CHR in XSB are presented informally. They do not differ essentially from other CHR systems.

When a constraint is called, it is considered an active constraint and the system will try to apply the rules to it. Rules are tried and executed sequentially in the order they are written.

A rule is conceptually tried for an active constraint in the following way. The active constraint is matched with a constraint in the head of the rule. If more constraints appear in the head they are looked for among the suspended constraints, which are called passive constraints in this context. If the necessary passive constraints can be found and all match with the head of the rule and the guard of the rule succeeds, then the rule is committed and the body of the rule executed. If not all the necessary passive constraint can be found, the matching fails or the guard fails, then the body is not executed and the process of trying and executing simply continues with the following rules. If for a rule, there are multiple constraints in the head, the active constraint will try the rule sequentially multiple times, each time trying to match with another constraint.

This process ends either when the active constraint disappears, i.e. it is removed by some rule, or after the last rule has been processed. In the latter case the active constraint becomes suspended.

A suspended constraint is eligible as a passive constraint for an active constraint. The other way it may interact again with the rules, is when a variable appearing in the constraint becomes bound to either a non-variable or another variable involved in one or more constraints. In that case the constraint is triggered, i.e. it becomes an active constraint and all the rules are tried.

Rule Types There are three different kinds of rules, each with their specific semantics:

- **simplification:**

The simplification rule removes the constraints in its head and calls its body.

- **propagation:**

The propagation rule calls its body exactly once for the constraints in its head.

- **simpagation:**

The simpagation rule removes the constraints in its head after the \setminus and then calls its body. It is an optimization of simplification rules of the form:

$$constraints_1, constraints_2 \leq => constraints_1, body$$

Namely, in the simpagation form:

$$constraints_1 \setminus constraints_2 \leq => body$$

The $constraints_1$ constraints are not called in the body.

Rule Names Naming a rule is optional and has no semantical meaning. It only functions as documentation for the programmer.

Pragmas The semantics of the pragmas are:

- **passive/1**: the constraint in the head of a rule with the identifier specified by the **passive/1** pragma can only act as a passive constraint in that rule.

Additional pragmas may be released in the future.

11.3 CHR in XSB Programs

11.3.1 Embedding in XSB Programs

Since `chr` is an XSB package, it must be explicitly loaded before being used.

```
?- [chr].
```

CHR rules are written in a `tt` `.chr` file. They should be preceded by a declaration of the constraints used:

```
:- constraints ConstraintSpec1, ConstraintSpec2, ...
```

where each `ConstraintSpec` is a functor description of the form `name/arity` pair. Ordinary code may be freely written between the CHR rules.

The CHR constraints defined in a particular `.chr` file are associated with a CHR module. The CHR module name can be any atom. The default module is `user`. A different module name can be declared as follows:

```
:- chr_module(modulename).
```

One should never load different files with the same CHR module name.

11.3.2 Compilation

Files containing CHR rules are required to have a `.chr` extension, and their compilation has two steps. First the `.chr` file is preprocessed into a `.P` file containing XSB code. This `.P` file can then be loaded in the XSB emulator and used normally.

```
load_chr(File)                                     chr_pp
    load_chr/1 takes as input a file name whose extension is either .chr or that has no extension.
    It preprocesses File if the times of the CHR rule file is newer than that of the corresponding
    Prolog file, and then consults the Prolog file.
```

```
preprocess(File,PFile)                             chr_pp
    preprocess/2 takes as input a file name whose extension is either .chr or that has no
    extension. It preprocesses File if the times of the CHR rule file is newer than that of the
    corresponding Prolog file, but does not consult the Prolog file.
```

11.4 Useful Predicates

The `chr` module contains several useful predicates that allow inspecting and printing the content of the constraint store.

<code>show_store(+Mod)</code>	<code>chr</code>
Prints all suspended constraints of module <code>Mod</code> to the standard output.	
<code>suspended_chr_constraints(+Mod,-List)</code>	<code>chr</code>
Returns the list of all suspended CHR constraints of the given module.	

11.5 Examples

Here are two example constraint solvers written in CHR.

- The program below defines a solver with one constraint, `leq/2`, which is a less-than-or-equal constraint.

```
:- chr_module(leq).

:- export cycle/3.

:- import length/2 from basics.

:- constraints leq/2.
reflexivity @ leq(X,X) <=> true.
antisymmetry @ leq(X,Y), leq(Y,X) <=> X = Y.
idempotence @ leq(X,Y) \ leq(X,Y) <=> true.
transitivity @ leq(X,Y), leq(Y,Z) ==> leq(X,Z).

cycle(X,Y,Z):-
    leq(X,Y),
    leq(Y,Z),
    leq(Z,X).
```

- The program below implements a simple finite domain constraint solver.

```
:- chr_module(dom).

:- import member/2 from basics.

:- constraints dom/2.
```

```

dom(X,[]) <=> fail.
dom(X,[Y]) <=> X = Y.
dom(X,L1), dom(X,L2) <=> intersection(L1,L2,L3), dom(X,L3).

intersection([],_,[]).
intersection([H|T],L2,[H|L3]) :-
    member(H,L2), !,
    intersection(T,L2,L3).
intersection([_|T],L2,L3) :-
    intersection(T,L2,L3).

```

These and more examples can be found in the `examples/chr/` folder accompanying this XSB release.

11.6 CHR and Tabling

The advantage of CHR in XSB over other Prolog systems, is that CHR can be combined with tabling. Hence part of the constraint solving can be performed once and reused many times. This has already shown to be useful for applications of model checking with constraints.

However the use of CHR constraints is slightly more complicated for tabled predicates. This section covers how exactly to write a tabled predicate that has one or more arguments that also appear as arguments in suspended constraints. In the current release the CHR-related parts of the tabled predicates have to be written by hand. In a future release this may be substituted by an automatic transformation.

11.6.1 General Issues and Principles

The general issue is how call constraints should be passed in to the tabled predicate and how answer constraints are passed out of the predicate. Additionally, in some cases care has to be taken not to generate infinite programs.

The recommended approach is to write the desired tabled predicate as if no additional code is required to integrate it with CHR. Next transform the tabled predicate to take into account the combination of tabling and CHR. Currently this transformation step has to be done by hand. In the future we hope to replace this hand coding with programmer declarations that guide automated transformations.

Hence we depart from an ordinary tabled predicate, say `p/1`:

```

:- table p/1.

p(X) :-
    ... /* original body of p/1 */.

```

In the following we will present several transformations or extensions of this code to achieve a particular behavior. At least the transformation discussed in subsection 11.6.2 should be applied to obtain a working integration of CHR and tabling. Further extensions are optional.

11.6.2 Call Abstraction

Currently only one type of call abstraction is supported: full constraint abstraction, i.e. all constraints on variables in the call should be removed. The technique to accomplish this is to replace all variables in the call that have constraints on them with fresh variables. After the call, the original variables should be unified with the new ones.

In addition, the call environment constraint store should be replaced with an empty constraint store before the call and on return the answer store should be merged back into the call environment constraint store.

The previously mentioned tabled predicate `p/1` should be transformed to:

```
:- import merge_answer_store/1,
        get_chr_store/1,
        set_chr_store/1,
        get_chr_answer_store/2
   from chr.

:- table tabled_p/2.

p(X) :-
    tabled_p(X1,AnswerStore),
    merge_answer_store(AnswerStore),
    X1 = X.

tabled_p(X,AnswerStore) :-
    get_chr_store(CallStore),
    set_chr_store(_EmptyStore)
    orig_p(X),
    get_chr_answer_store(chrmod,AnswerStore),
    set_chr_store(CallStore).

orig_p(X) :-
    ... /* original body of p/1 */.
```

This example shows how to table the CHR constraints of a single CHR module `chrmod`. If multiple CHR modules are involved, one should add similar arguments for the other modules.

11.6.3 Answer Projection

To get rid of irrelevant constraints, most notably on local variables, the answer constraint store should in some cases be projected on the variables in the call. This is particularly important for

programs where otherwise an infinite number of answers with ever growing answer constraint stores could be generated.

The current technique of projection is to provide an additional `project/1` constraint to the CHR solver definition. The argument of this constraint is the list of variables to project on. Appropriate CHR rules should be written to describe the interaction of this `project/1` constraint with other constraints in the store. An additional rule should take care of removing the `project/1` constraint after all such interaction.

The `project/1` constraint should be posed before returning from the tabled predicate.

If this approach is not satisfactory or powerful enough to implement the desired projection operation, you should resort to manipulating the underlying constraint store representation. Contact the maintainer of XSB's CHR system for assistance.

Example Take for example a predicate `p/1` with a less than or equal constraint `leq/2` on variables and integers. The predicate `p/1` has local variables, but when `p` returns we are not interested in any constraints involving local variables. Hence we project on the argument of `p/1` with a project constraint as follows:

```
:- import memberchk/2 from lists.

:- import merge_answer_store/1,
    get_chr_store/1,
    set_chr_store/1,
    get_chr_answer_store/2
    from chr.

:- table tabled_p/2.

:- constraints leq/2, project/1.

... /* other CHR rules */
project(L) \ leq(X,Y) <=>
    ( var(X), \+ memberchk(X,L)
      ; var(Y), \+ memberchk(Y,L)
    ) | true.

project(_) <=> true.

p(X) :-
    tabled_p(X1,AnswerStore),
    merge_answer_store(AnswerStore),
    X1 = X.

tabled_p(X,AnswerStore) :-
    get_chr_store(CallStore),
    set_chr_store(_EmptyStore)
    orig_p(X),
    project([X]),
```

```

    get_chr_answer_store(chrmod,AnswerStore),
    set_chr_store(CallStore).

orig_p(X) :-
    ... /* original body of p/1 */.

```

The example in the following subsection shows projection in a full application.

11.6.4 Answer Combination

Sometimes it is desirable to combine different answers to a tabled predicate into one single answer or a subset of answers. Especially when otherwise there would be an infinite number of answers. If the answers are expressed as constraints on some arguments and the logic of combining is encoded as CHR rules, answers can be combined by merging the respective answer constraint stores.

Another case where this is useful is when optimization is desired. If the answer to a predicate represents a valid solution, but an optimal solution is desired, the answer should be represented as constraints on arguments. By combining the answer constraints, only the most constrained, or optimal, answer is kept.

Example An example of a program that combines answers for both termination and optimisation is the shortest path program below:

```

:- chr_module(path).

:- import length/2 from lists.

:- import merge_chr_answer_store/1,
    get_chr_store/1,
    set_chr_store/1,
    get_chr_answer_store/2
    from chr.

breg_retskel(A,B,C,D) :- '_$builtin'(154).

:- constraints geq/2, plus/3, project/1.

geq(X,N) \ geq(X,M) <=> number(N), number(M), N =< M | true.

reflexivity @ geq(X,X) <=> true.
antisymmetry @ geq(X,Y), geq(Y,X) <=> X = Y.
idempotence @ geq(X,Y) \ geq(X,Y) <=> true.
transitivity @ geq(X,Y), geq(Y,Z) ==> var(Y) | geq(X,Z).

plus(A,B,C) <=> number(A), number(B) | C is A + B.
plus(A,B,C), geq(A,A1) ==> plus(A1,B,C1), geq(C,C1).
plus(A,B,C), geq(B,B1) ==> plus(A,B1,C1), geq(C,C1).

```

```

project(X) \ plus(_,_,_) # ID <=> true pragma passive(ID).
project(X) \ geq(Y,Z) # ID <=> (Y \== X ; var(Z) )| true pragma passive(ID).
project(_) <=> true.

path(X,Y,C) :-
  tabled_path(X,Y,C1,AS),
  merge_chr_answer_store(AS),
  C = C1.

:- table tabled_path/4.

tabled_path(X,Y,C,AS) :-
  '$_$savecp'(Breg),
  breg_retskel(Breg,4,Skel,Cs),
  copy_term(p(X,Y,C,AS,Skel),p(OldX,OldY,OldC,OldAS,OldSkel)),
  get_chr_store(GS),
  set_chr_store(_GS1),
  orig_path(X,Y,C),
  project(C),
  ( get_returns(Cs,OldSkel,Leaf),
    OldX == X, OldY == Y ->
      merge_chr_answer_store(OldAS),
      C = OldC,
      get_chr_answer_store(path,MergedAS),
      sort(MergedAS,AS),
      ( AS = OldAs ->
        fail
      );
      delete_return(Cs,Leaf)
    ),
  ;
  get_chr_answer_store(path,UnsortedAS),
  sort(UnsortedAS,AS)
),
  set_chr_store(GS).

orig_path(X,Y,C) :- edge(X,Y,C1), geq(C,C1).
orig_path(X,Y,C) :- path(X,Z,C2), edge(Z,Y,C1), plus(C1,C2,C0), geq(C,C0).

edge(a,b,1).
edge(b,a,1).
edge(b,c,1).
edge(a,c,3).
edge(c,a,1).

```

The predicate `orig_path/3` specifies a possible path between two nodes in a graph. In `tabled_path/4` multiple possible paths are combined together into a single path with the shortest distance. Hence the tabling of the predicate will reject new answers that have a worse distance and will replace the old answer when a better answer is found. The final answer gives the optimal solution, the shortest path. It is also necessary for termination to keep only the best answer. When cycles appear in the

graph, paths with longer and longer distance could otherwise be put in the table, contributing to the generation of even longer paths. Failing for worse answers avoids this infinite build-up.

The predicate also includes a projection to remove constraints on local variables and only retain the bounds on the distance.

The sorting canonicalizes the answer stores, so that they can be compared.

11.6.5 Overview of Tabling-related Predicates

<code>merge_answer_store(+AnswerStore)</code>	chr
Merges the given CHR answer store into the current global CHR constraint store.	
<code>get_chr_store(-ConstraintStore)</code>	chr
Returns the current global CHR constraint store.	
<code>set_chr_store(?ConstraintStore)</code>	chr
Set the current global CHR constraint store. If the argument is a fresh variable, the current global CHR constraint store is set to be an empty store.	
<code>get_chr_answer_store(+Mod, -AnswerStore)</code>	chr
Returns the part of the current global CHR constraint store of constraints in the specified CHR module, in the format of an answer store usable as a return argument of a tabled predicate.	

11.7 Guidelines

In this section we cover several guidelines on how to use CHR to write constraint solvers and how to do so efficiently.

- **Set semantics:** The CHR system allows the presence of identical constraints, i.e. multiple constraints with the same functor, arity and arguments. For most constraint solvers, this is not desirable: it affects efficiency and possibly termination. Hence appropriate simpagation rules should be added of the form:

$$constraint \backslash constraint \leq => true$$

- **Multi-headed rules:** Multi-headed rules are executed more efficiently when the constraints share one or more variables.

11.8 CHRd

An alternate implementation of CHR can be found in the CHRd package. The main objective of the CHRd package is to optimize processing of constraints in the environment where termination is guaranteed by the tabling engine, (and where termination benefits provided by the existing solver

are not critical). CHRd takes advantage of XSB's tabling to simplify CHR's underlying storage structures and solvers. Specifically, we entirely eliminate the thread-global constraint store in favor of a distributed one, realized as a collection of sets of constraints entirely associated with program variables. This decision limits the applicability of CHRd to a restricted class of CHR programs, referred to as direct-indexed CHR, in which all constraints in the head of a rule are connected by shared variables. Most CHR programs are direct-indexed, and other programs may be easily converted to fall into this class. Another advance of CHRd is its set-based semantics which removes the need to maintain the propagation history, thus allowing further simplicity in the representation of the constraints. The CHRd package itself is described in [?], and both the semantics of CHRd and the class of direct-indexed CHR are formally defined in [?].

Chapter 12

XASP: Answer Set Programming with XSB and Smodels

By Luis Castro, Terrance Swift, David S. Warren ¹

The term *Answer Set Programming (ASP)* describes a paradigm in which logic programs are interpreted using the (extended) stable model semantics. While the stable model semantics is quite elegant, it has radical differences from traditional program semantics based on Prolog. First, stable model semantics applies only to ground programs; second stable model semantics is not goal-oriented – determining whether a stable model is true in a program involves examining each clause in a program, regardless of whether the goal would depend on the clause in a traditional evaluation ².

Despite (or perhaps because of) these differences, ASP has proven to be a useful paradigm for solving a variety of combinatorial programs. Indeed, determining a stable model for a logic program can be seen as an extension of the NP-complete problem of propositional satisfiability, so that satisfiability problems that can be naturally represented as logic programs can be solved using ASP.

The current generation of ASP systems are very efficient for determining whether a program has a stable model (analogous to whether the program, taken as a set of propositional axioms, is satisfiable). However, ASP systems have somewhat primitive file-based interfaces. XSB is a natural complement to ASP systems. Its basis in Prolog provides a procedural counterpart for ASP, as described in Chapter 5 of Volume 1 of this manual; and XSB's computation of the Well-founded semantics has a well-defined relationship to stable model semantics. Furthermore, deductive-database-like capabilities of XSB allow it to be an efficient and flexible grounder for many ASP problems.

The XASP package provides various mechanisms that allow tight linkage of XSB programs to the Smodels [?] stable model generator. The main interface is based on a store of clauses that can

¹ Thanks to Barry Evans for helping resuscitate the XASP installation procedure, and to Gonalo Lopes for the installation procedure on Windows.

²In Version 3.4, the Smodels API has not been tested with the multi-threaded engine, and Smodels itself is not thread-safe.

be incrementally asserted or deleted by an XSB program. Clauses in this store can make use of all of the cardinality and weight constraint syntax supported by Smodels, in addition to default negation. When the user decides that the clauses in a store are a complete representation of a program whose stable model should be generated, the clauses are copied into Smodels buffers. Using the Smodels API, the generator is invoked, and information about any stable models generated are returned. This use of XASP is roughly analogous to building up a constraint store in CLP, and periodically evaluating that store, but integration with the store is less transparent in XASP than in CLP. In XASP, clauses must be explicitly added to a store and evaluated; furthermore clauses are not removed from the store upon backtracking, unlike constraints in CLP.

The XNMR interpreter provides a second, somewhat more implicit use of XASP. In the XNMR interface a query Q is evaluated as is any other query in XSB. However, conditional answers produced for Q and for its subgoals, upon user request, can be considered as clauses and sent to Smodels for evaluation. In backtracking through answers for Q , the user backtracks not only through answer substitutions for variables of Q , but also through the stable models produced for the various bindings.

12.1 Installing the Interface

Installing the Smodels interface of XASP sometimes can be tricky for two reasons. First, XSB must dynamically load the Smodels library, and dynamic loading introduces platform dependencies. Second since Smodels is written in C++ and XSB is written in C, the load must ensure that names are properly resolved and that C++ libraries are loaded, steps that may be addressed differently by different compilers³. However, by following the steps outlined below in the section for Unix or Windows, XASP should be running in a matter of minutes.

12.1.1 Installing the Interface under Unix

In order to use the Smodels interface, several steps must be performed.

1. *Creating a library for Smodels.* Smodels itself must be compiled as a library. Unlike previous versions of XSB, which required a special configuration step for Smodels, Version 3.4 requires no special configuration, since XSB includes source code for Smodels 2.33 as a subdirectory of the `$XSBDIR/packages/xasp` directory (denoted `$XASPDIR`). We suggest making Smodels out of this directory⁴. Thus, to make the Smodels library

- (a) Change directory to `$XASPDIR/smodels`
- (b) On systems other than OS X, type

```
make lib
```

³XSB's compiler can automatically call foreign compilers to compile modules written in C, but in Version 3.4 of XSB C++ modules must be compiled with external commands, such as the `make` command shown below.

⁴Although distributed with XSB, Smodels is distributed under the GNU General Public License, a license that is slightly stricter than the license XSB uses. Users distributing applications based on XASP should be aware of any restrictions imposed by GNU General Public License.

on OS X, type ⁵

```
make -f Makefile.osx lib
```

If the compilation step ran successfully, there should be a file `libsmodels.so` (or `libsmodels.dylib` on MacOS X or `libsmodels.dll` on Windows...) in `$XASPDIR/smodels/.libs`

(c) Change directory back to `$XASPDIR`

2. *Compiling the XASP files* Next, platform-specific compilation of XASP files needs to be performed. This can be done by consulting `prologMake.P` and executing the goal

```
?- make.
```

It is important to note that under Version 3.4, code compiled by the single threaded engine will only be executable by the single threaded engine, and code compiled by the multi-threaded engine will only be executable by the multi-threaded engine.

3. *Checking the Installation* To see if the installation is working properly, `cd` to the subdirectory `tests` and type:

```
sh testsuite.sh <$XSBDIR>
```

If the test suite succeeded it will print out a message along the lines of

```
PASSED testsuite for /Users/terranceswift/XSBNEW/XSB/config/powerpc-apple-darwin7.5.1/bin/xsb
```

12.1.2 Installing XASP under Windows using Cygwin

To install XASP under Windows, you must use Version 3.4 of XSB or later and Version 2.31 or later of Smodels ⁶. You should also have a recent version of Cygwin (e.g. 1.5.20 or later) with all the relevant development packages installed, such as `devel`, `make`, `automake`, `patchtools`, and possibly `x11` (for `makedepend`) Without an appropriate Cygwin build environment many of these steps will simply fail, sometimes with quite cryptic error messages.

1. *Patch and Compile Smodels* First, uncompress `smodels-2.31.tar.gz` in some directory, (for presentation purposes we use `/cygdrive/c/smodels-2.31` — that is, `c:\smodels-2.31`). After that, you must apply the patch provided with this package. This patch enables the creation of a DLL from Smodels. Below is a sample session (system output omitted) with the required commands:

```
$ cd /cygdrive/c/smodels-2.31
$ cat $XSB/packages/xasp/patch-smodels-2.31 | patch -p1
$ make lib
```

⁵A special makefile is needed for OS X since the GNU libtool is called glibtool on this platform.

⁶This section was written by Goncalo Lopes.

After that, you should have a file called `smodels.dll` in the current directory, as well as a file called `smodels.a`. You should make the former "visible" to Windows. Two alternatives are either (a) change the `PATH` environment variable to contain `c:\smodels-2.31`, or (b) copy `smodels.dll` to some other directory in your `PATH` (such as `c:\windows`, for instance). One simple way to do this is to copy `smodels.dll` to `$XSB/config/i686-pc-cygwin/bin`, *after* the configure XSB step (step 2), since that directory has to be in your path in order to make XSB fully functional.

2. *Configure XSB*. In order to properly configure XSB, you must tell it where the Smodels sources and library (the `smodels.a` file) are. In addition, you must compile XSB such that it doesn't use the Cygwin DLL (using the `-mno-cygwin` option for gcc). The following is a sample command:

```
$ cd $XSB/build
$ ./configure --enable-no-cygwin --with-smodels="/cygdrive/c/smodels-2.31"
```

You can optionally include the extended Cygwin w32 API using the configuration option `--with-includes=<PATH_TO_API>`, (this allows XSB's build procedure to find `makedepend` for instance), but you'll probably do fine with just the standard Cygwin apps.

There are some compiler variables which may not be automatically set by the configure script in `xsb_config.h`, namely the configuration names and some activation flags. To correct this, do the following:

- (a) cd to `$XSB/config/i686-pc-cygwin`
- (b) open the file `xsb_config.h` and add the following lines:

```
#define CONFIGURATION "i686-pc-cygwin"
#define FULL_CONFIG_NAME "i686-pc-cygwin"
#define SLG_GC
```

(Still more flags may be needed depending on Cygwin configuration)

After applying these changes, cd back to the `$XSB/build` directory and compile XSB:

```
$ ./makexsb
```

Now you should have in `$XSB/config/i686-pc-cygwin/bin` directory both a `xsb.exe` and a `xsb.dll`.

3. *Compiling XASP*. First, go to the XASP directory and execute the `makelinks.sh` script in order to make the headers and libraries in Smodels be accessible to XSB, i.e.:

```
$ cd $XSB/packages/xasp
$ sh makelinks.sh /cygdrive/c/smodels-2.31
```

Now you must copy the `smoMakefile` from the `config` directory to the `xasp` directory and run both its directives:

```
$ cp $XSB/config/i686-pc-cygwin/smoMakefile .
$ make -f smoMakefile module
$ make -f smoMakefile all
```

At this point, you can consult `xnmr` as you can with any other package, or `xsb` with the `xnmr` command line parameter, like this: (don't forget to add XSB bin directory to the `$PATH` environment variable)

```
$ xsb xnmr
```

Lots of error messages will probably appear because of some runtime load compiler, but if everything goes well you can ignore all of them since your `xasppkg` will be correctly loaded and everything will be functioning smoothly from there on out.

12.2 The Smodels Interface

The Smodels interface contains two levels: the *cooked* level and the *raw* level. The cooked level interns rules in an XSB *clause store*, and translates general weight constraint rules [?] into a *normal form* that the Smodels engine can evaluate. When the programmer has determined that enough clauses have been added to the store to form a semantically complete sub-program, the program is *committed*. This means that information in the clauses is copied to Smodels and interned using Smodels data structures so that stable models of the clauses can be computed and examined. By convention, the cooked interface ensures that the atom `true` is present in all stable models, and the atom `false` is false in all stable models. The raw level models closely the Smodels API, and demands, among other things, that each atom in a stable sub-program has been translated into a unique integer. The raw level also does not provide translation of arbitrary weight constraint rules into the normal form required by the Smodels engine. As a result, the raw level is significantly more difficult to directly use than the cooked level. While we make public the APIs for both the raw and cooked level, we provide support only for users of the cooked interface.

As mentioned above Smodels extends normal programs to allow weight constraints, which can be useful for combinatorial problems. However, the syntax used by Smodels for weight constraints does not follow ISO Prolog syntax so that the XSB syntax for weight constraints differs in some respects from that of Smodels. Our syntax is defined as follows, where A is a Prolog atom, N a non-negative integer, and I an arbitrary integer.

- $GeneralLiteral ::= WeightConstraint \mid Literal$
- $WeightConstraint ::= weightConst(Bound, WeightList, Bound)$
- $WeightList ::= List\ of\ WeightLiterals$
- $WeightLiteral ::= Literal \mid weight(Literal, N)$
- $Literal ::= A \mid not(A)$

- $Bound ::= I \mid \text{undef}$

Thus an example of a weight constraint might be:

- `weightConst(1,[weight(a,1),weight(not(b),1)],2)`

We note that if a user does not wish to put an upper or lower bound on a weight constraint, she may simply set the bound to `undef` or to an integer less than 0.

The intuitive semantics of a weight constraint `weightConst(Lower,WeightList,Upper)`, in which `List` is a list of *WeightLiterals* that it is true in a model M whenever the sum of the weights of the literals in the constraint that are true in M is between the lower `Lower` and `Upper`. Any literal in a *WeightList* that does not have a weight explicitly attached to it is taken to have a weight of 1.

In a typical session, a user will initialize the Smodels interface, add rules to the clause store until it contains a semantically meaningful sub-problem. He can then specify a compute statement if needed, commit the rules, and compute and examine stable models via backtracking. If desired, the user can then re-initialize the interface, and add rules to or retract rules from the clause store until another semantically meaningful sub-program is defined; and then commit, compute and examine another stable model ⁷.

The process of adding information to a store and periodically evaluating it is vaguely reminiscent of the Constraint Logic Programming (CLP) paradigm, but there are important differences. In CLP, constraints are part of the object language of a Prolog program: constraints are added to or projected out of a constraint store upon forward execution, removed upon backwards execution, and iteratively checked. When using this interface, on the other hand, an XSB program essentially acts as a compiler for the clause store, which is treated as a target language. Clauses must be explicitly added or removed from the store, and stable model computation cannot occur incrementally – it must wait until all clauses have been added to the store. We note in passing that the `xnmr` module provides an elegant but specialized alternative. `xnmr` integrates stable models into the object language of XSB, by computing "relevant" stable models from the the residual answers produced by query evaluation. It does not however, support the weighted constraint rules, compute statements and so on that this module supports.

Neither the raw nor the cooked interface currently supports explicit negation.

Examples of use of the various interfaces can be found in the subdirectory `intf_examples`

`smcInit`

Initializes the XSB clause store and the Smodels API. This predicate must be executed before building up a clause store for the first time. The corresponding raw predicate, `smrInit(Num)`, initializes the Smodels API assuming that it will require at most `Num` atoms.

`smcReInit`

Reinitializes the Smodels API, but does *not* affect the XSB clause store. This predicate is

⁷Currently, only normal rules can be retracted.

provided so that a user can reuse rules in a clause store in the context of more than one sub-program.

`smcAddRule(+Head,+Body)`

Interns a ground rule into the XSB clause store. **Head** must be a *GeneralLiteral* as defined at the beginning of this section, and **Body** must be a list of *GeneralLiterals*. Upon interning, the rule is translated into a normal form, if necessary, and atoms are translated to unique integers. The corresponding raw predicates, `smrAddBasicRule/3`, `smrAddChoiceRule/3`, `smrAddConstraintRule/4`, and `smrAddWeightRule/3` can be used to add raw predicates immediately into the SModels API.

`smcRetractRule(+Head,+Body)`

Retracts a ground (basic) rule from the XSB clause store. Currently, this predicate cannot retract rules with weight constraints: **Head** must be a *Literal* as defined at the beginning of this section, and **Body** must be a list of *GeneralLiterals*.

`smcSetCompute(+List)`

Requires that **List** be a list of literals – i.e. atoms or the default negation of atoms). This predicate ensures that each literal in **List** is present in the stable models returned by Smodels. By convention the cooked interface ensures that **true** is present and **false** absent in all stable models. After translating a literal it calls the raw interface predicates `smrSetPosCompute/1` and `smrSetNegCompute/1`

`smcCommitProgram`

This predicate translates all of the clauses from the XSB clause store into the data structures of the Smodels API. It then signals to the API that all clauses have been added, and initializes the Smodels computation. The corresponding raw predicate, `smrCommitProgram`, performs only the last two of these features.

`smComputeModel`

This predicate calls Smodels to compute a stable model, and succeeds if a stable model can be computed. Upon backtracking, the predicate will continue to succeed until all stable models for a given program cache have been computed. `smComputeModel/0` is used by both the raw and the cooked levels.

`smcExamineModel(+List,-Atoms)`

`smcExamineModel(+List,-Atoms)` filters the literals in **List** to determine which are true in the most recently computed stable model. These true literals are returned in the list **Atoms**. `smrExamineModel(+N,-Atoms)` provides the corresponding raw interface in which integers from 0 to N, true in the most recently computed stable model, are input and output.

`smEnd`

Reclaims all resources consumed by Smodels and the various APIs. This predicate is used by both the cooked and the raw interfaces.

`print_cache`

This predicate can be used to examine the XSB clause store, and may be useful for debugging.

12.2.1 Using the Smodels Interface with Multiple Threads

If XASP has been compiled under the multi-threaded engine, the Smodels interface will be fully thread-safe: this means that Smodels and all interface predicates described in this section can be used concurrently by different threads. In multi-threaded XASP, each XSB thread can initialize and query its own instance of Smodels, and build up its own private clause store at both the cooked and raw levels (shared clause stores are not yet available). Figure 12.1 provides a simple example of how this can be done. For each thread that will generate stable models, a message queue is created that will be used to communicate back results. Two threads are then created and these threads concurrently add rules to their private clause stores, call Smodels, and send the results back to the calling thread using the appropriate message queue. Of course the example here is just one of many possible: answers could be returned using different configurations of message queues, through shared tables, through shared asserted code, and so on.

12.3 The `xnmr_int` Interface

. This module provides the interface from the `xnmr` module to Smodels. It does not use the `sm_int` interface, but rather directly calls the Smodels C interface, and can be thought of as a special-purpose alternative to `sm_int`.

`init_smodels(+Query)`

Initializes smodels with the residual program produced by evaluating `Query`. `Query` must be a call to a tabled predicate that is currently completely evaluated (and should have a delay list)

`atom_handle(?Atom,?AtomHandle)`

The *handle* of an atom is set by `init_smodels/1` to be an integer uniquely identifying each atoms in the residual program (and thus each atom in the Herbrand base of the program for which the stable models are to be derived). The initial query given to `init_smodels` has the atom-handle of 1.

`in_all_stable_models(+AtomHandle,+Neg)`

`in_all_stable_models/2` returns true if `Neg` is 0 and the atom numbered `AtomHandle` returns true in all stable models (of the residual program set by the previous call to `init_smodels/1`). If `Neg` is nonzero, then it is true if the atom is in NO stable model.

`pstable_model(+Query,-Model,+Flag)`

returns nondeterministically a list of atoms true in the partial stable model total on the atoms relevant to instances of `Query`, if `Flag` is 0. If `Flag` is 1, it only returns models in which the instance of `Query` is true.

`a_stable_model`

This predicate invokes Smodels to find a (new) stable model (of the program set by the previous invocation of `init_smodels/1`.) It will compute all stable models through backtracking. If there are no (more) stable models, it fails. Atoms true in a stable model can be examined by `in_current_stable_model/1`.

```

:- ensure_loaded(xasp).
:- import smcInit/0, smcAddRule/2, smcCommitProgram/0 smcSetCompute/1,
      smComputeModel/0, smcExamineModel/1, smEnd/0 from sm_int.
:- import thread_create/1 from thread.
:- import thread_get_message/2, thread_send_message/2, message_queue_create/1 from mutex_xsb.

test:-
    message_queue_create(Queue1),
    message_queue_create(Queue2),
    thread_create(test1(Queue1)),
    thread_create(test2(Queue2)),
    read_models(Queue1),
    read_models(Queue2).

test1(Queue) :-
    smcInit,
    smcAddRule(a1,[]),
    smcAddRule(b1,[]),
    smcAddRule(d1,[a1,not(c1)]),
    smcAddRule(c1,[b1,not(d1)]),
    smcCommitProgram,
    write('All Solutions: '),nl,
    (   smComputeModel,
        smcExamineModel(Model),
        thread_send_message(Queue,solution(program1,Model)),
        fail
    );
    thread_send_message(Queue,no_more_solutions),
    smEnd ).

test2(Queue) :-
    smcInit,
    smcAddRule(a2,[]),
    smcAddRule(b2,[]),
    smcAddRule(d2,[a2,not(c2)]),
    smcAddRule(c2,[b2,not(d2)]),
    smcCommitProgram,
    write('All Solutions: '),nl,
    (   smComputeModel,
        smcExamineModel(Model),
        thread_send_message(Queue,solution(program2,Model)),
        fail
    );
    thread_send_message(Queue,no_more_solutions),
    smEnd ).

read_models(Queue):-
    repeat,
    thread_get_message(Queue,Message),
    (Message = no_more_solutions ->
        true
    ; writeln(Message),
        fail ).

```

Figure 12.1: Using the Smodels Interface with Multi-Threading

`in_current_stable_model(?AtomHandle)`

This predicate is true of handles of atoms true in the current stable model (set by an invocation of `a_stable_model/0`.)

`current_stable_model(-AtomList)`

returns the list of atoms true in the current stable model.

`print_current_stable_model`

prints the current stable model to the stream to which answers are sent (i.e `stdfbk`)

Chapter 13

PITA: Probabilistic Inference with Tabling and Answer subsumption

By Fabrizio Riguzzi

“Probabilistic Inference with Tabling and Answer subsumption” (PITA) [?] is a package for uncertain reasoning. In particular, it allows various forms of Probabilistic Logic Programming and Possibilistic Logic Programming. It accepts the language of Logic Programs with Annotated Disjunctions (LPADs) [?, ?] and CP-logic programs [?, ?].

An example of LPAD/CP-logic program is

$$\begin{aligned} (heads(Coin) : 0.5) \vee (tails(Coin) : 0.5) &\leftarrow toss(Coin), \neg biased(Coin). \\ (heads(Coin) : 0.6) \vee (tails(Coin) : 0.4) &\leftarrow toss(Coin), biased(Coin). \\ (fair(Coin) : 0.9) \vee (biased(Coin) : 0.1) & \\ &toss(Coin). \end{aligned}$$

The first clause states that if we toss a coin that is not biased it has equal probability of landing heads and tails. The second states that if the coin is biased it has a slightly higher probability of landing heads. The third states that the coin is fair with probability 0.9 and biased with probability 0.1 and the last clause states that we toss a coin with certainty.

PITA computes the probability of queries by transforming the input program into a normal logic program and then calling a modified version of the query on the transformed programs.

13.0.1 Installation

PITA uses **GLib 2.0** and **CUDD**. GLib is a standard GNU package so it is easy to install it using the package management software of your Linux distribution.

To install CUDD, follow the instructions at <http://vlsi.colorado.edu/~fabio/CUDD/> to get the package (or get directly from <ftp://vlsi.colorado.edu/pub/cudd-2.4.2.tar.gz>), for

example `cudd-2.4.2.tar.gz`. After decompressing, you will have a directory `cudd-2.4.2` with various subdirectories. Compile CUDD following the included instructions.

To install PITA with XSB, run `XSB configure` in the `build` directory with option `-with-pita=DIR` where `DIR` is the folder where CUDD is.

Syntax

Disjunction in the head is represented with a semicolon and atoms in the head are separated from probabilities by a colon. For the rest, the usual syntax of Prolog is used. For example, the CP-logic clause

$$h_1 : p_1 \vee \dots \vee h_n : p_n \leftarrow b_1, \dots, b_m, \neg c_1, \dots, \neg c_l$$

is represented by

```
h1:p1 ; ... ; hn:pn :- b1,...,bm,\+ c1,...,\+ cl
```

No parentheses are necessary. The `pi` are numeric expressions. It is up to the user to ensure that the numeric expressions are legal, i.e. that they sum up to less than one.

If the clause has an empty body, it can be represented like this

```
h1:p1 ; ... ; hn:pn.
```

If the clause has a single head with probability 1, the annotation can be omitted and the clause takes the form of a normal prolog clause, i.e.

```
h1:- b1,...,bm,\+ c1,...,\+ cl.
```

stands for

```
h1:1 :- b1,...,bm,\+ c1,...,\+ cl.
```

The body of clauses can contain a number of built-in predicates including:

```
is/2 >/2 </2 >=/2 <=/2 ==/2 ==\=/2 true/0 false/0
=/2 ==/2 \=/2 \==/2 length/2 member/2
```

The coin example above thus is represented as (see file `coin.cpl` in subdirectory `examples`)

```
heads(Coin):1/2 ; tails(Coin):1/2:-
    toss(Coin),\+biased(Coin).
heads(Coin):0.6 ; tails(Coin):0.4:-
    toss(Coin),biased(Coin).
fair(Coin):0.9 ; biased(Coin):0.1.
toss(coin).
```

Subdirectory `examples` contains other example programs.

13.0.2 Use

Probabilistic Logic Programming

First write your program in a file with extension `.cpl`. If you want to use inference on LPADs load PITA in XSB with

```
:- [pita].
```

load you program, say `coin.cpl`, with

```
:- load(coin).
```

and compute the probability of query atom `heads(coin)` by

```
:- prob(heads(coin),P).
```

`load(file)` reads `file.cpl`, translates it into a normal program, writes the result in `file.P` and loads `file.P`.

PITA offers also the predicate `parse(infile,outfile)` which translates the LPAD in `infile` into a normal program and writes it to `outfile`.

Moreover, you can use `prob(goal,P,CPUTime,WallTime)` that returns the probability of goal `P` together with the CPU and wall time used.

In case the modeling assumptions of PRISM hold, i.e.:

- the probability of a conjunction (A, B) is computed as the product of the probabilities of A and B (independence assumption),
- the probability of a disjunction $(A; B)$ is computed as the sum of the probabilities of A and B (exclusiveness assumption),

you can perform faster inference with an optimized version of PITA in package `pitaindexc.P`. It accepts the same commands of `pita.P`. `pitaindexc.P` simulates PRISM and does not need CUDD and GLib.

If you want to compute the Viterbi path and probability of a query (the Viterbi path is the explanation with the highest probability) as with the predicate `viterbif/3` of PRISM, you can use package `pitavitind.P`.

The package `pitacount.P` can be used to count the explanations for a query, provided that the independence assumption holds. To count the number of explanations for a query use

```
:- oount(heads(coin),C).
```

`pitacount.P` does not need CUDD and GLib.

Possibilistic Logic Programming

PITA can be used also for answering queries to possibilistic logic program [?], a form of logic programming based on possibilistic logic [?]. The package `pitaposs.P` provides possibilistic inference. You have to write the possibilistic program as an LPAD in which the rules have a single head whose annotation is the lower bound on the necessity of the clauses. To compute the highest lower bound on the necessity of a query use

```
:- poss(heads(coin),P).
```

`pitaposs.P` does not need CUDD and GLib.

Chapter 14

Other XSB Packages

Many of the XSB packages are maintained somewhat independently of XSB and have their own manuals. For these packages: *Flora2*, *XMC*, *xsbdoc* and *Cold Dead Fish* we provide summaries; full information can be obtained in the packages themselves. In addition, we provide full documentation here for two of the smaller packages, *slx* and *GAP*.

14.1 Programming with FLORA-2

\mathcal{F} LORA-2 is a sophisticated object-oriented knowledge base language and application development platform. It is implemented as a set of run-time libraries and a compiler that translates a unified language of F-logic [?], HiLog [?], and Transaction Logic [?, ?] into tabled Prolog code.

Applications of \mathcal{F} LORA-2 include intelligent agents, Semantic Web, ontology management, integration of information, and others.

The programming language supported by \mathcal{F} LORA-2 is a dialect of F-logic with numerous extensions, which include a natural way to do meta-programming in the style of HiLog and logical updates in the style of Transaction Logic. \mathcal{F} LORA-2 was designed with extensibility and flexibility in mind, and it provides strong support for modular software design through its unique feature of dynamic modules. Other extensions, such as the versatile syntax of FLORID path expressions, are borrowed from FLORID, a C++-based F-logic system developed at Freiburg University.¹ Extensions aside, the syntax of \mathcal{F} LORA-2 differs in many important ways from FLORID, from the original version of F-logic, as described in [?], and from an earlier implementation of \mathcal{F} LORA. These syntactic changes were needed in order to bring the syntax of \mathcal{F} LORA-2 closer to that of Prolog and make it possible to include simple Prolog programs into \mathcal{F} LORA-2 programs without choking the compiler. Other syntactic deviations from the original F-logic syntax are a direct consequence of the added support for HiLog, which obviates the need for the “@” sign in method invocations (this sign is now used to denote calls to \mathcal{F} LORA-2 modules).

\mathcal{F} LORA-2 is distributed in two ways. First, it is part of the official distribution of XSB and thus is installed together with XSB. Second, a more up-to-date version of the system is available

¹ See <http://www.informatik.uni-freiburg.de/~dbis/florid/> for more details.

on *FLORA-2*'s Web site at

`http://flora.sourceforge.net`

These two versions can be installed at the same time and used independently (*e.g.*, if you want to keep abreast with the development of *FLORA-2* or if a newer version was released in-between the releases of XSB). The installation instructions are somewhat different in these two cases. Here we only describe the process of configuring the version *FLORA-2* included with XSB.

Installing *FLORA-2* under UNIX. To configure a version of *FLORA-2* that was downloaded as part of the distribution of XSB, simply configure XSB as usual:

```
cd XSB/build
configure
makexsb
```

and then run

```
makexsb packages
```

If you downloaded XSB from its CVS repository earlier and are updating your copy using the `cvs update` command, then it might be a good idea to also do the following:

```
cd packages/flora2
makeflora clean
makeflora
```

Installing *FLORA-2* in Windows. First, you need Microsoft's `nmake`. Then use the following commands to configure *FLORA-2* (assuming that XSB is already installed and configured):

```
cd flora2
makeflora clean
makeflora path-to-prolog-executable
```

Also make sure that the `packages` directory contains a shortcut called `flora2.P` to the file `packages\flora2\flora2.P`.

Running *FLORA-2*. *FLORA-2* is fully integrated into the underlying XSB engine, including its module system. In particular, *FLORA-2* modules can invoke predicates defined in other Prolog modules, and Prolog modules can query the objects defined in *FLORA-2* modules.

Due to certain problems with XSB, *FLORA-2* runs best when XSB is configured with *local* scheduling, which is the default XSB configuration. However, with this type of scheduling, many Prolog intuitions that relate to the operational semantics do not work. Thus, the programmer

must think “more declaratively” and, in particular, to not rely on the order in which answers are returned.

The easiest way to get a feel of the system is to start *FLORA-2* shell and begin to enter queries interactively. The simplest way to do this is to use the shell script

```
.../flora2/runflora
```

where “...” is the directory where *FLORA-2* is downloaded. For instance, to invoke the version supplied with XSB, you would type something like

```
~/XSB/packages/flora2/runflora
```

At this point, *FLORA-2* takes over and F-logic syntax becomes the norm. To get back to the Prolog command loop, type **Control-D** (Unix) or **Control-Z** (Windows), or

```
| ?- _end.
```

If you are using *FLORA-2* shell frequently, it pays to define an alias, say (in Bash):

```
alias runflora='~/XSB/packages/flora2/runflora'
```

FLORA-2 can then be invoked directly from the shell prompt by typing **runflora**. It is even possible to tell *FLORA-2* to execute commands on start-up. For instance,

```
foo> runflora -e "_help."
```

will cause the system to execute the help command right after the initialization. Then the usual *FLORA-2* shell prompt is displayed.

FLORA-2 comes with a number of demo programs that live in

```
.../flora2/demos/
```

The demos can be run issuing the command “**_demo(demo-filename).**” at the *FLORA-2* prompt, *e.g.*,

```
flora2 ?- _demo(flogic_basics).
```

There is no need to change to the demo directory, as **flDemo** knows where to find these programs.

14.2 Summary of xmc: Model-checking with XSB

No documentation yet available.

the Ciao [?] system's *lpdoc* which has been adapted to generate a reference manual automatically from one or more XSB source files. The target format of the documentation can be Postscript, HTML, PDF, or nicely formatted ASCII text. *xsbdoc* can be used to automatically generate a description of full applications, library modules, README files, etc. A fundamental advantage of using *xsbdoc* to document programs is that it is much easier to maintain a true correspondence between the program and its documentation, and to identify precisely to what version of the program a given printed manual corresponds. Naturally, the *xsbdoc* manual is generated by *xsbdoc* itself.

The quality of the documentation generated can be greatly enhanced by including within the program text:

- *assertions* (indicating types, modes, etc. ...) for the predicates in the program, via the directive `pred/1`; and
- *machine-readable comments* (in the “literate programming” style).

The assertions and comments included in the source file need to be written using the forthcoming XSB *assertion language*, which supports most of the features of Ciao's assertion language within a simple and (hopefully) intuitive syntax.

xsbdoc is distributed under the *GNU general public license*.

Unlike *lpdoc*, *xsbdoc* does not use Makefiles, and instead maintains information about how to generate a document within Prolog *format files*. As a result, *xsbdoc* can in principle be run in any environment that supports the underlying software, such as XSB, L^AT_EX, dvips and so on. It has been tested on Linux and Windows running with Cygwin.

14.3 slx: Extended Logic Programs under the Well-Founded Semantics

As explained in the section *Using Tabling in XSB*, XSB can compute normal logic programs according to the well-founded semantics. In fact, XSB can also compute *Extended Logic Programs*, which contain an operator for explicit negation (written using the symbol `-`) in addition to the negation-by-failure of the well-founded semantics (`\+` or `not`). Extended logic programs can be extremely useful when reasoning about actions, for model-based diagnosis, and for many other uses [?]. The library, *slx* provides a means to compile programs so that they can be executed by XSB according to the *well-founded semantics with explicit negation* [?]. Briefly, WFSX is an extension of the well-founded semantics to include explicit negation and which is based on the *coherence principle* in which an atom is taken to be default false if it is proven to be explicitly false, intuitively:

$$-p \Rightarrow \text{not } p.$$

This section is not intended to be a primer on extended logic programming or on WFSX semantics, but we do provide a few sample programs to indicate the action of WFSX. Consider the program

```
s:- not t.
```

```
t:- r.
t.
```

```
r:- not r.
```

If the clause `-t` were not present, the atoms `r`, `t`, `s` would all be undefined in WFSX just as they would be in the well-founded semantics. However, when the clause `t` is included, `t` becomes true in the well-founded model, while `s` becomes false. Next, consider the program

```
s:- not t.
```

```
t:- r.
-t.
```

```
r:- not r.
```

In this program, the explicitly false truth value for `t` obtained by the rule `-t` overrides the undefined truth value for `t` obtained by the rule `t:- r`. The WFSX model for this program will assign the truth value of `t` as false, and that of `s` as true. If the above program were contained in the file `test.P`, an XSB session using `test.P` might look like the following:

```
> xsb

| ?- [slx].
[slx loaded]

yes
| ?- slx_compile('test.P').
[Compiling ./tmptest]
[tmptest compiled, cpu time used: 0.1280 seconds]
[tmptest loaded]

| ?- s.

yes
| ?- t.

no
| ?- naf t.

yes
| ?- r.

no
| ?- naf r.

no
| ?- und r.
```

yes

In the above program, the query `?- t.` did not succeed, because `t` is false in WFSX: accordingly the query `na f t` did succeed, because it is true that `t` is false via negation-as-failure, in addition to `t` being false via explicit negation. Note that after being processed by the SLX preprocessor, `r` is undefined but does not succeed, although `und r` will succeed.

We note in passing that programs under WFSX can be paraconsistent. For instance in the program.

```
p:- q.

q:- not q.
   -q.
```

both `p` and `q` will be true *and* false in the WFSX model. Accordingly, under SLX preprocessing, both `p` and `na f p` will succeed.

```
slx_compile(+File)                                module: slx
Preprocesses and loads the extended logic program named File. Default negation in File must be represented using the operator not rather than using tnot or \+. If L is an objective literal (e.g. of the form A or  $-A$  where A is an atom), a query ?- L will succeed if L is true in the WFSX model, na f L will succeed if L is false in the WFSX model, and und L will succeed if L is undefined in the WFSX model.
```

14.4 gapza: Generalized Annotated Programs

Generalized Annotated Programs (GAPs) [?] offer a powerful computational framework for handling paraconsistency and quantitative information within logic programs. The tabling of XSB is well-suited to implementing GAPs, and the gap library provides a meta-interpreter that has proven robust and efficient enough for a commercial application in data mining. The current meta-interpreter is limited to range-restricted programs.

A description of GAPs along with full documentation for this meta-interpreter is provided in [?] (currently also available at <http://www.cs.sunysb.edu/~tswift>). Currently, the interface to the GAP library is through the following call.

```
meta(?Annotated_atom)                             module: gap
If Annotated_atom is of the form Atom:[Lattice_type,Annotation] the meta-interpreter computes bindings for Atom and Annotation by evaluating the program according to the definitions provided for Lattice_type.
```